

# Implementing Load Balancing with Instance Groups

---



**Vitthal Srinivasan**

CO-FOUNDER, LOONYCORN

[www.loonycorn.com](http://www.loonycorn.com)

# Overview

**Introducing instance templates**

**Managed instance groups and unmanaged instance groups**

**Architectural overview of the HTTP(S) load balancer and its components**

**Cross-regional load balancers with unmanaged and managed instance groups**

**Autoscaling with managed instance groups**

# Instance Groups

---

# Instance Groups

A group of machines which can be created and managed together to avoid individually controlling each instance in the group

# Two Kinds of Instance Groups



**Managed**

**Unmanaged**

# Two Kinds of Instance Groups



Managed



Unmanaged

## Unmanaged

Groups of dissimilar instances that you can add and remove from the group

Do not offer autoscaling, rolling updates or instance templates

Not recommended, used only when you need to apply **load balancing to pre-existing** configurations

# Two Kinds of Instance Groups



**Managed**

**Unmanaged**





Managed

Uses an **instance template** to create a group of **identical** instances

Changes to the instance group changes all instances in the group


# Managed Instance Group

Group of identical GCE VM instances, created from the same instance template that are managed by the platform


# Instance Template

A specification of machine type, boot disk (or container image), zone, labels and other instance properties that can be used to instantiate either individual VM instances or a Managed Instance Group

# Instance Template



# Attractions of MIGs



## **Autoscaling**

**Associate autoscaling policy with  
MIG**




## **Autohealing**

**Associate health check and  
autohealing policy with MIG**

# Autoscaling and Autohealing

---

# Attractions of MIGs



## Autoscaling


Associate autoscaling policy with  
MIG



## Autohealing

Associate health check and  
autohealing policy with MIG

# Attractions of MIGs



## Autoscaling

Associate autoscaling policy with  
MIG

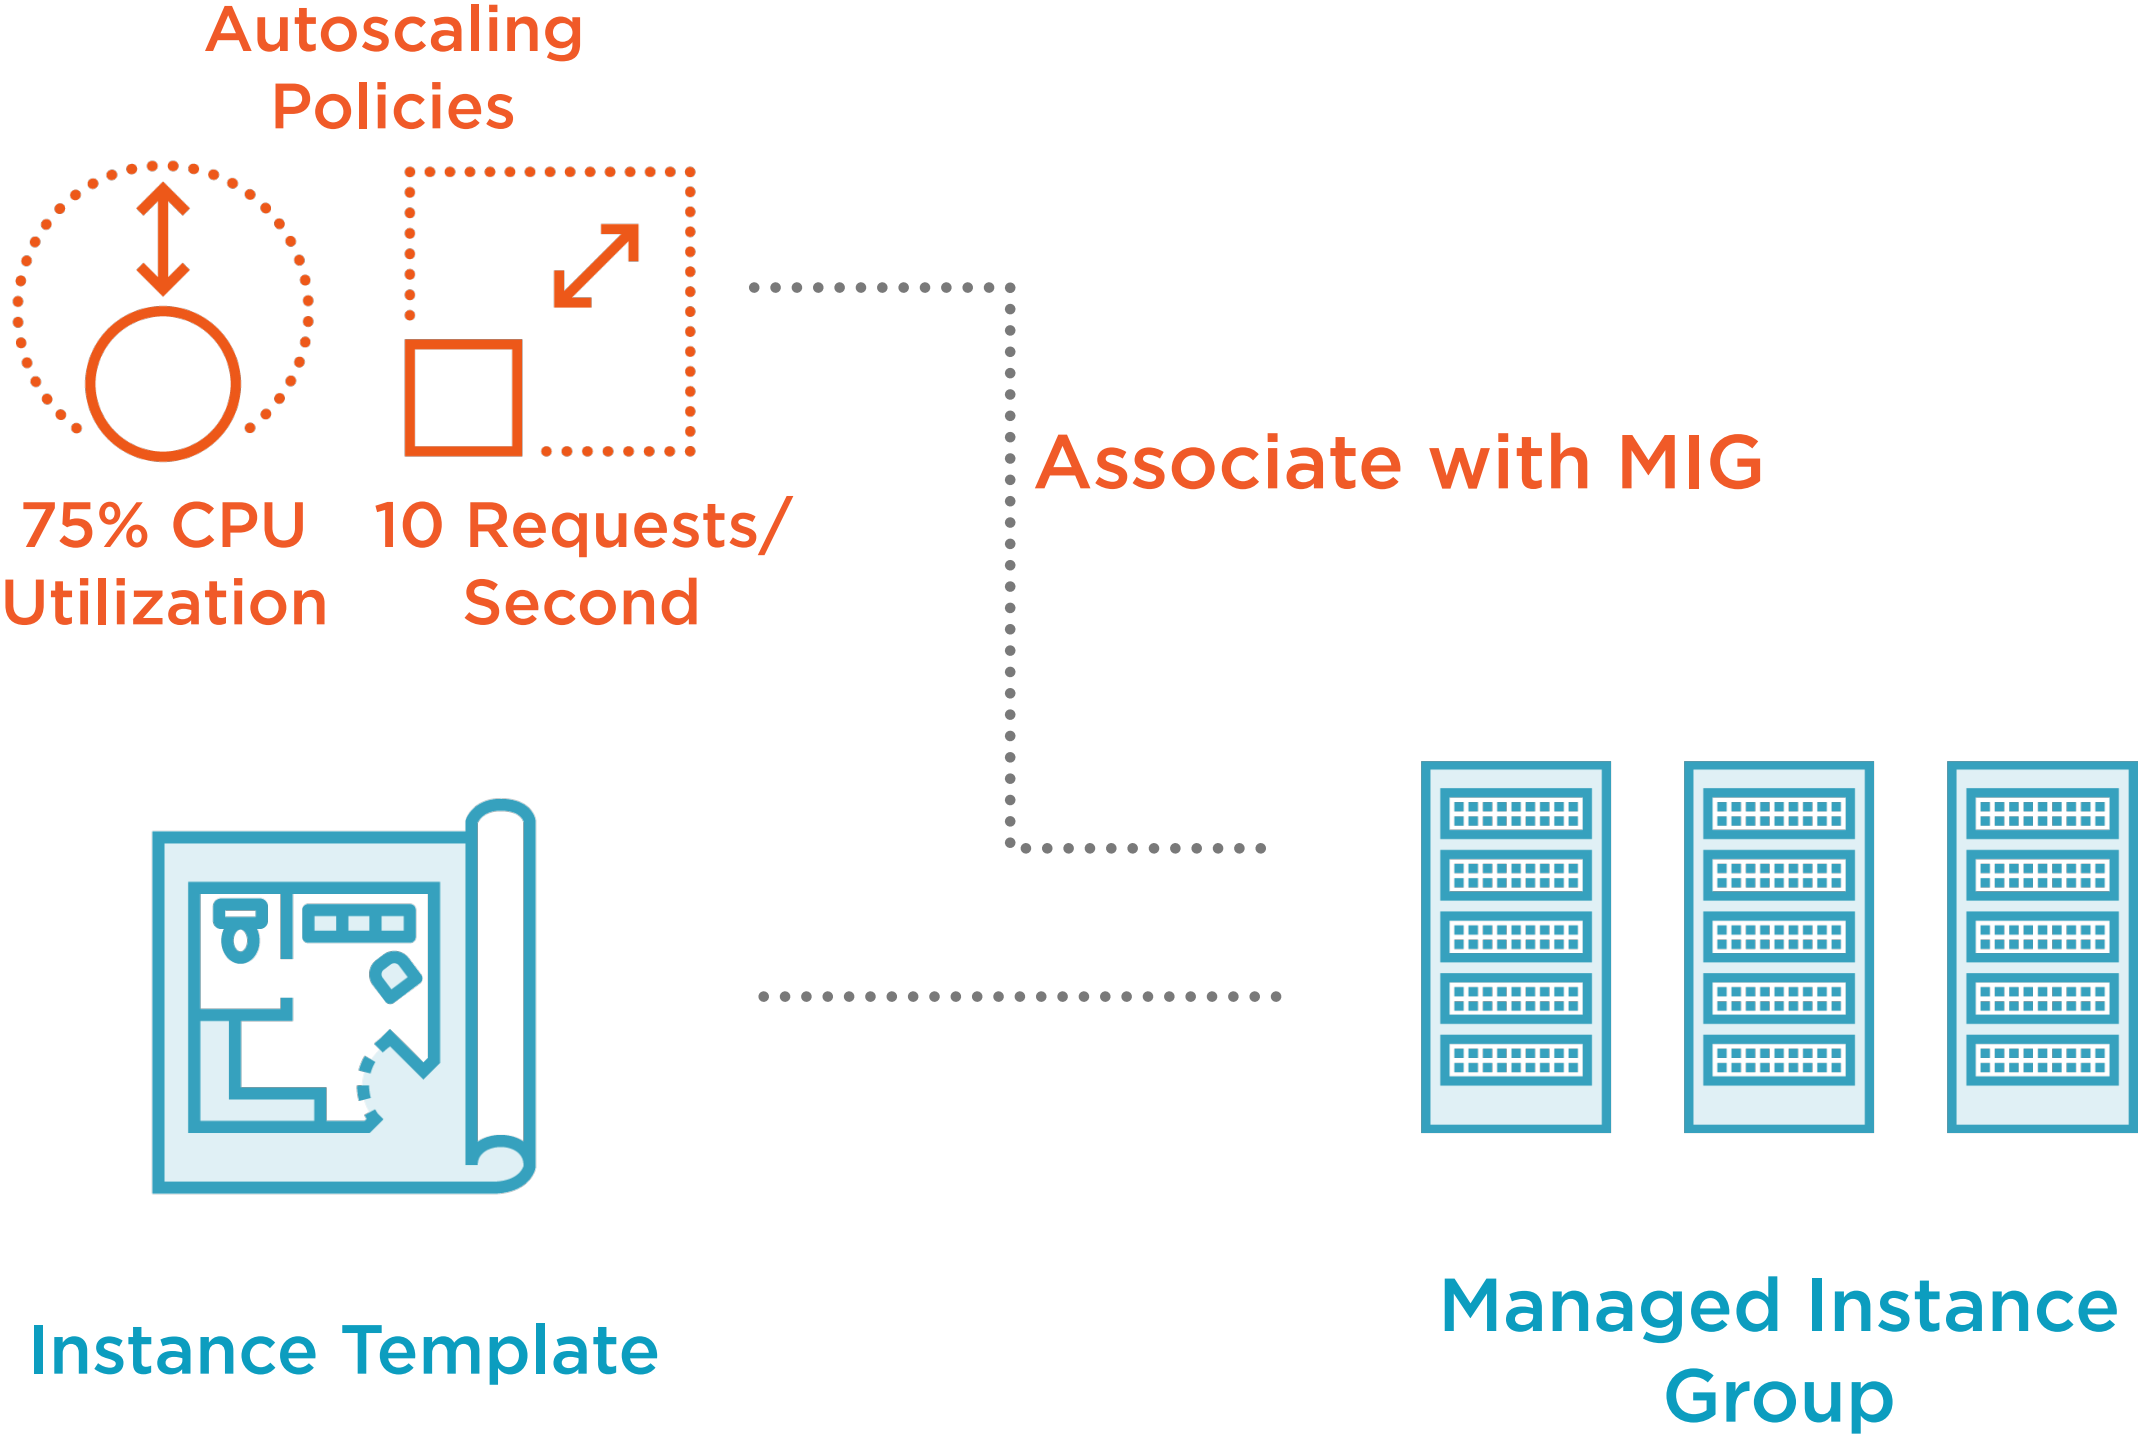


## Autohealing

Associate health check and  
autohealing policy with MIG




# Autoscaling Policies




# Autoscaling Policies

## Autoscaling Policies



Instance Template


Service monitors each VM instance vs. each policy



Managed Instance Group

# Autoscaling Policies


## Autoscaling Policies



Check, on average whether policy is being satisfied




Instance Template



Managed Instance Group


# Autoscaling Policies

## Autoscaling Policies



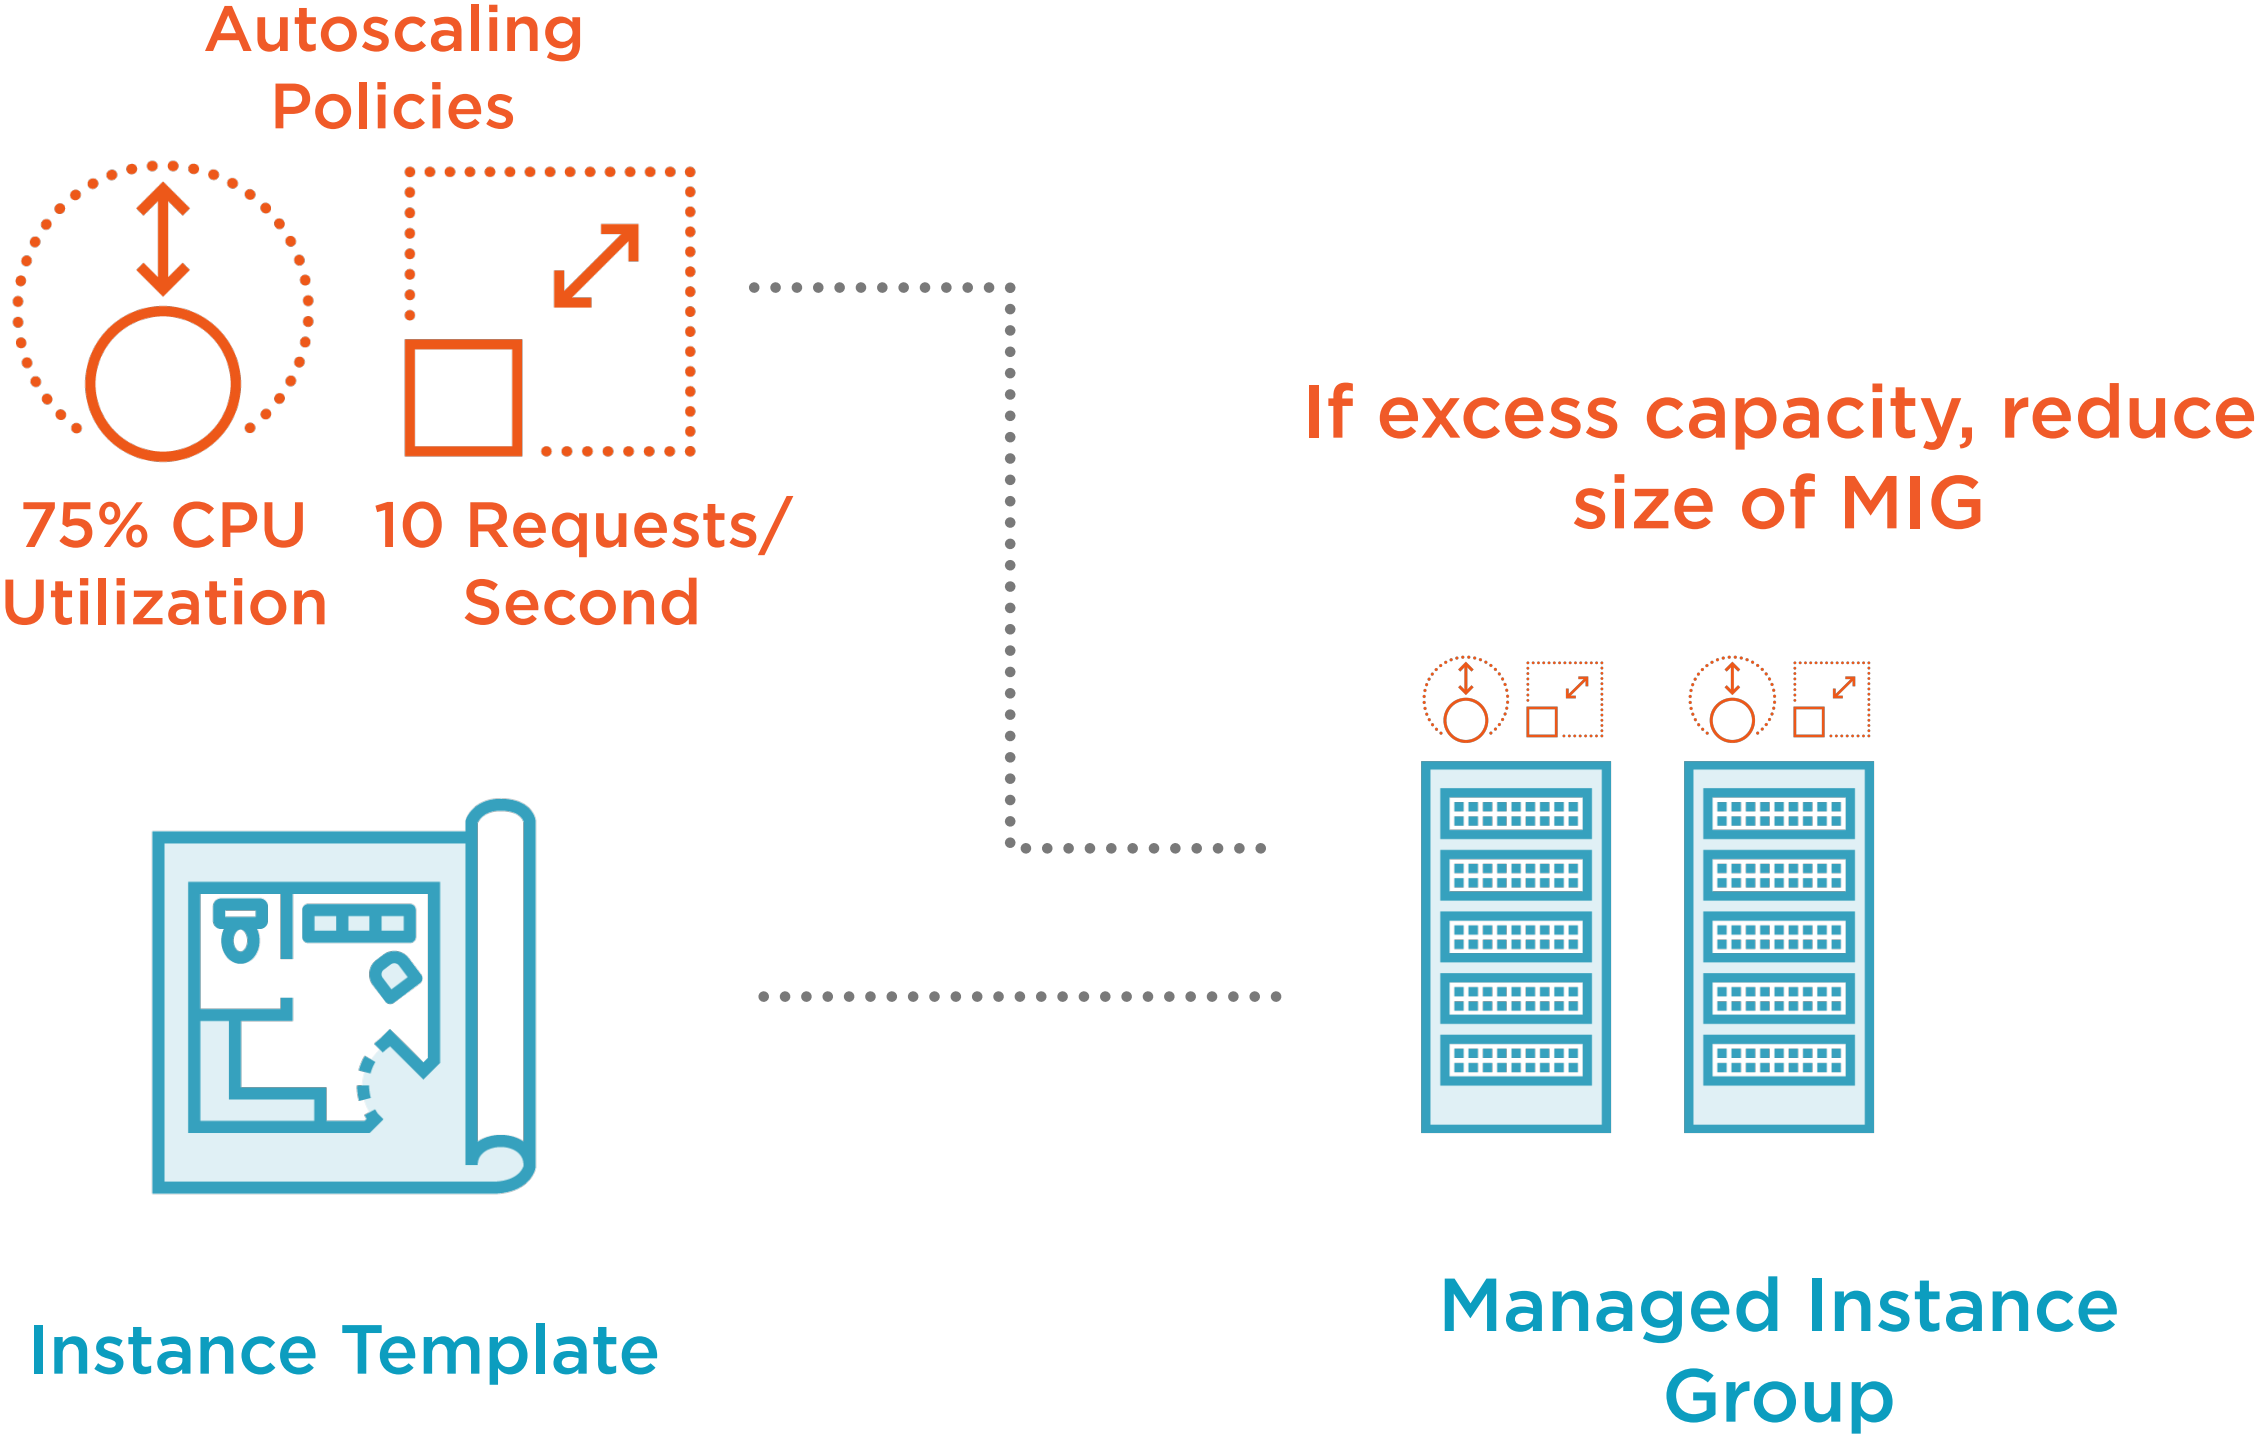
Instance Template

If excess capacity, reduce size of MIG



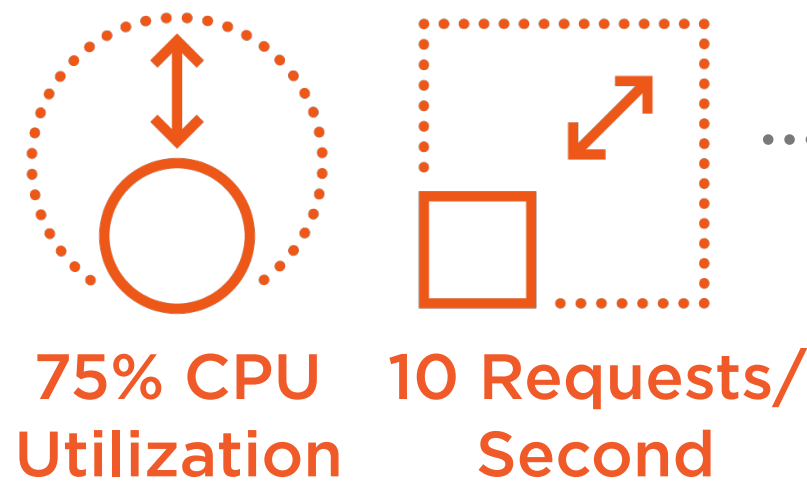
Managed Instance Group

# Autoscaling Policies




# Autoscaling Policies

## Autoscaling Policies



Instance Template

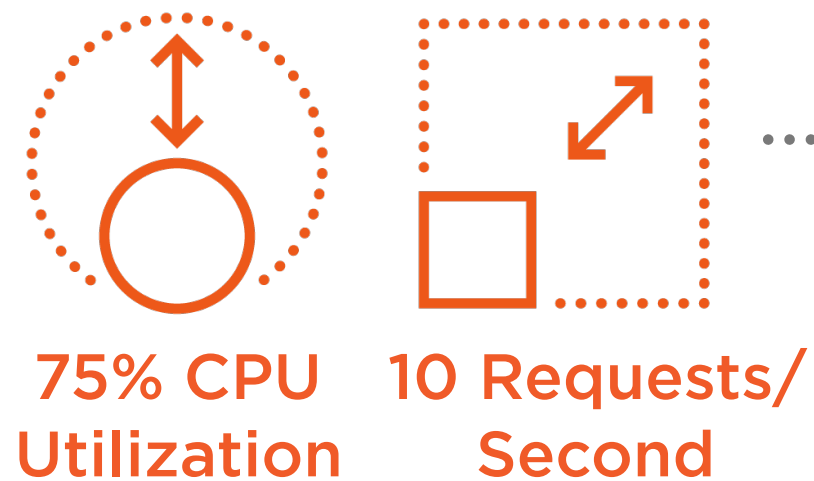
If insufficient capacity,  
scale up size of MIG



Managed Instance Group


# Autoscaling Policies

## Autoscaling Policies



Instance Template

If insufficient capacity,  
scale up size of MIG



Managed Instance Group

# Attractions of MIGs



## Autoscaling

Associate autoscaling policy with  
MIG



## Autohealing

Associate health check and  
autohealing policy with MIG



# Health Checks

Health Check



Associate Health  
Check with MIG



Instance Template



Managed Instance  
Group

# Health Checks

Health Check



Instance Template

Sends probes to check health  
of each instance in MIG



Managed Instance  
Group


# Health Checks

Health Check



Instance Template

Probes identify any unhealthy instance



Managed Instance Group

# Health Checks

Health Check



Instance Template

MIG then replaces it  
with a healthy one




Managed Instance  
Group

# HTTP(S) Load Balancing

---

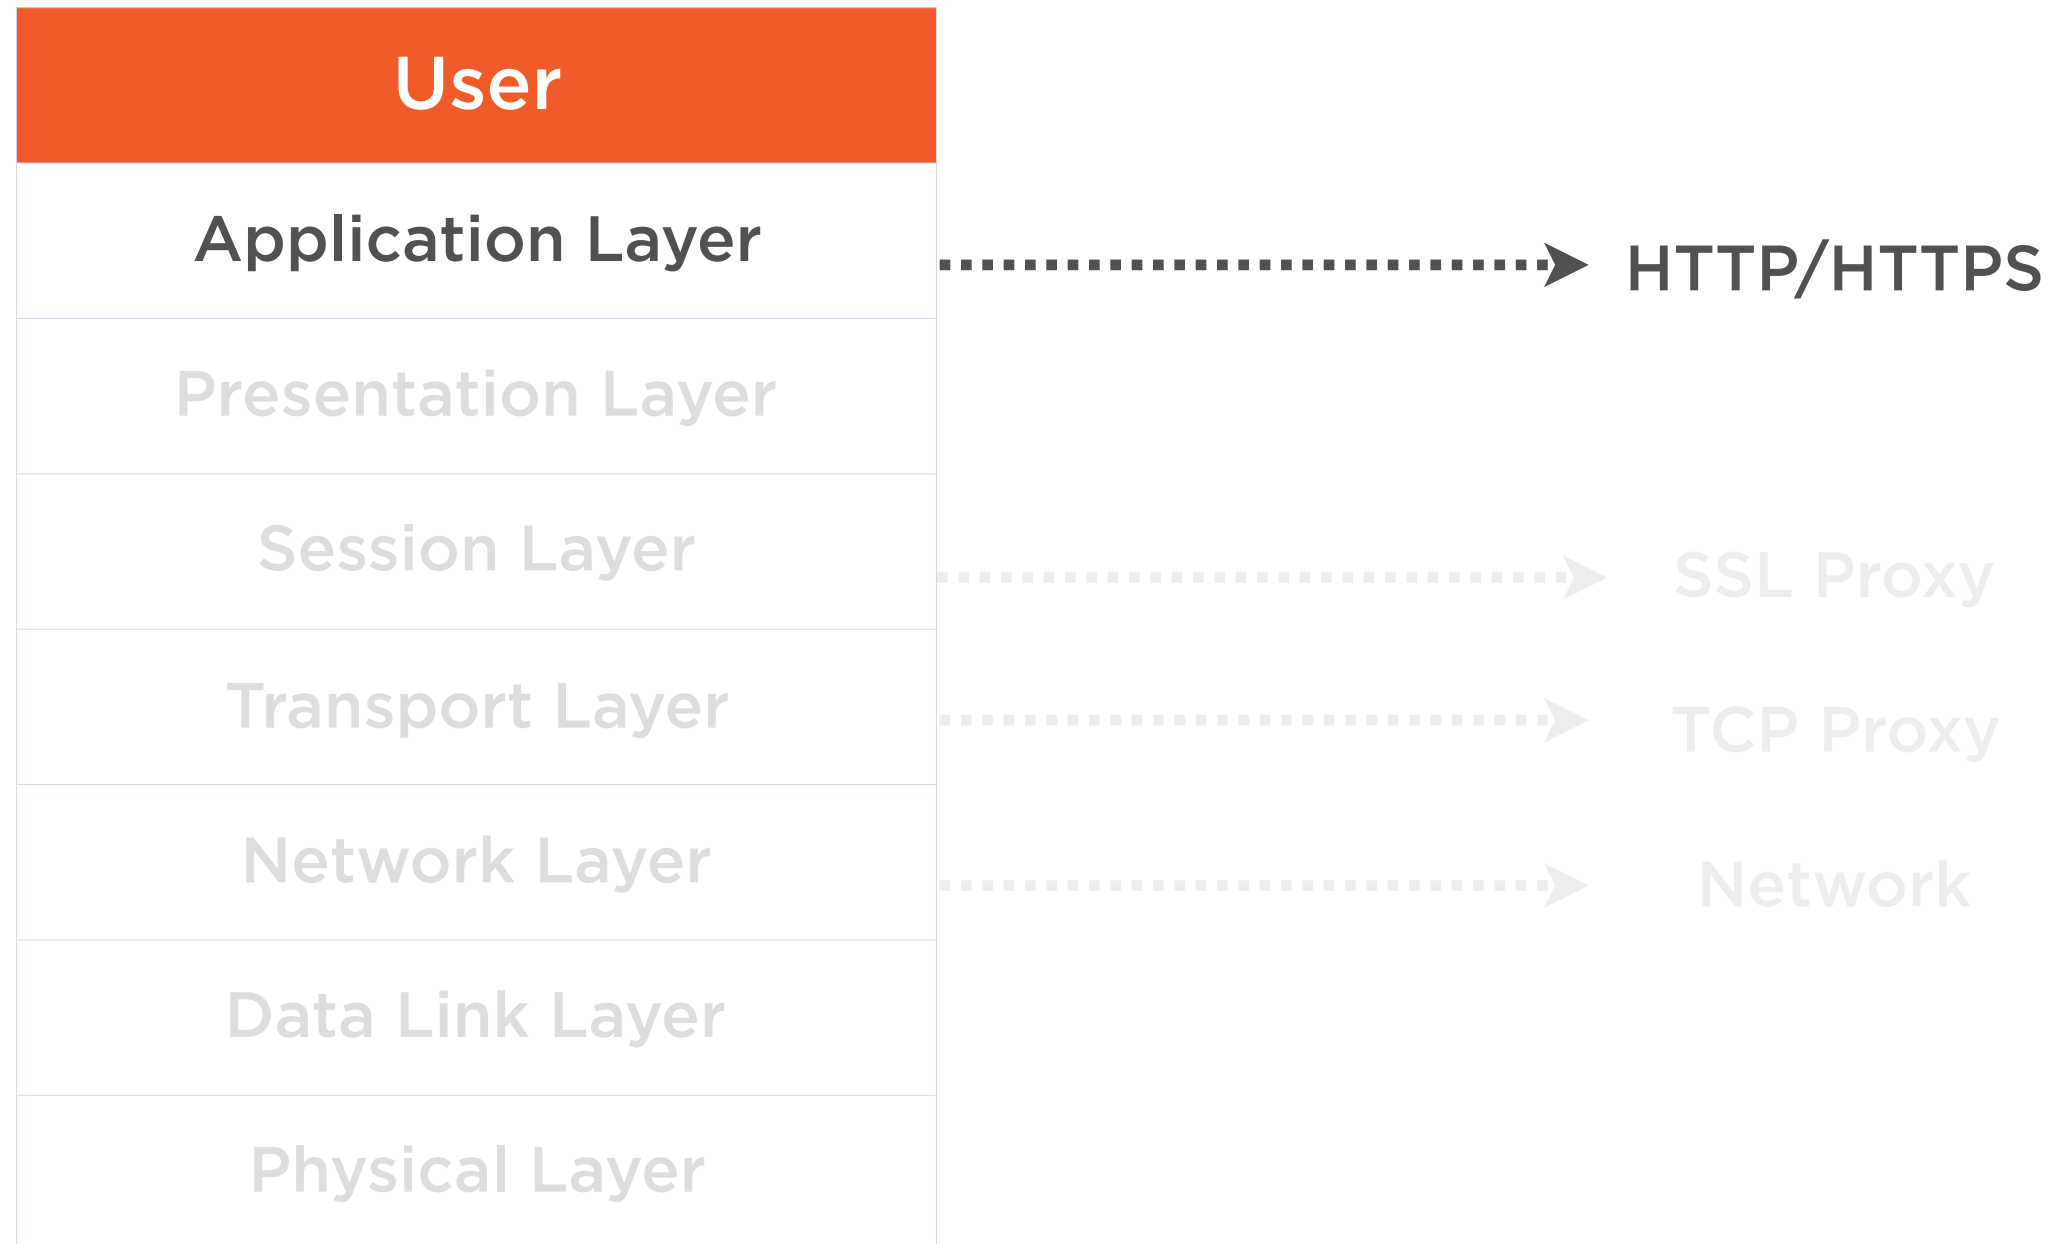
# Load Balancing



# HTTP(S) Load Balancing



# HTTP(S) Load Balancing



**HTTP(S) is used to balance global, external traffic**



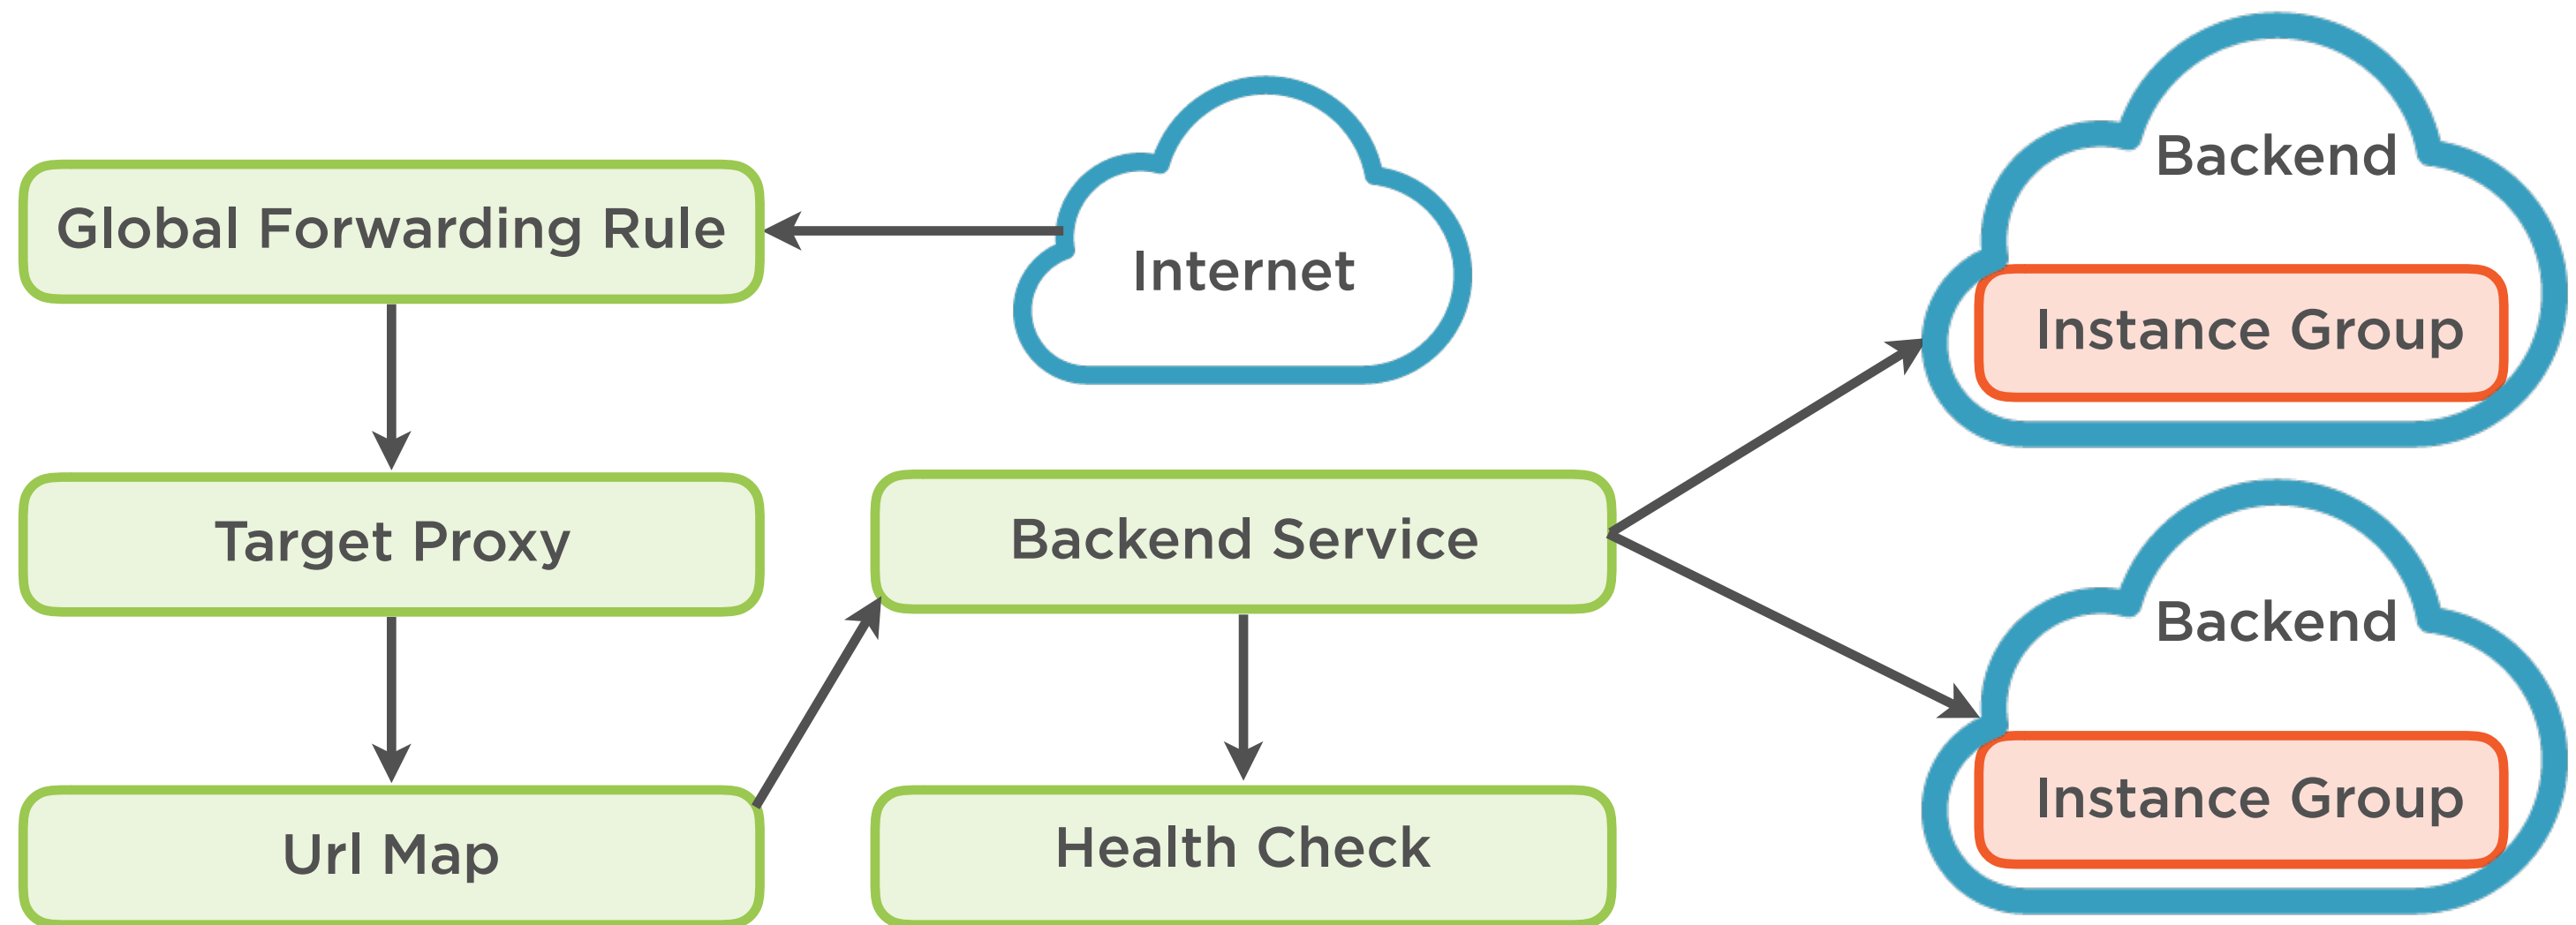
# HTTP(S) Load Balancing



**Distributes HTTP(S) traffic among groups of instances based on:**

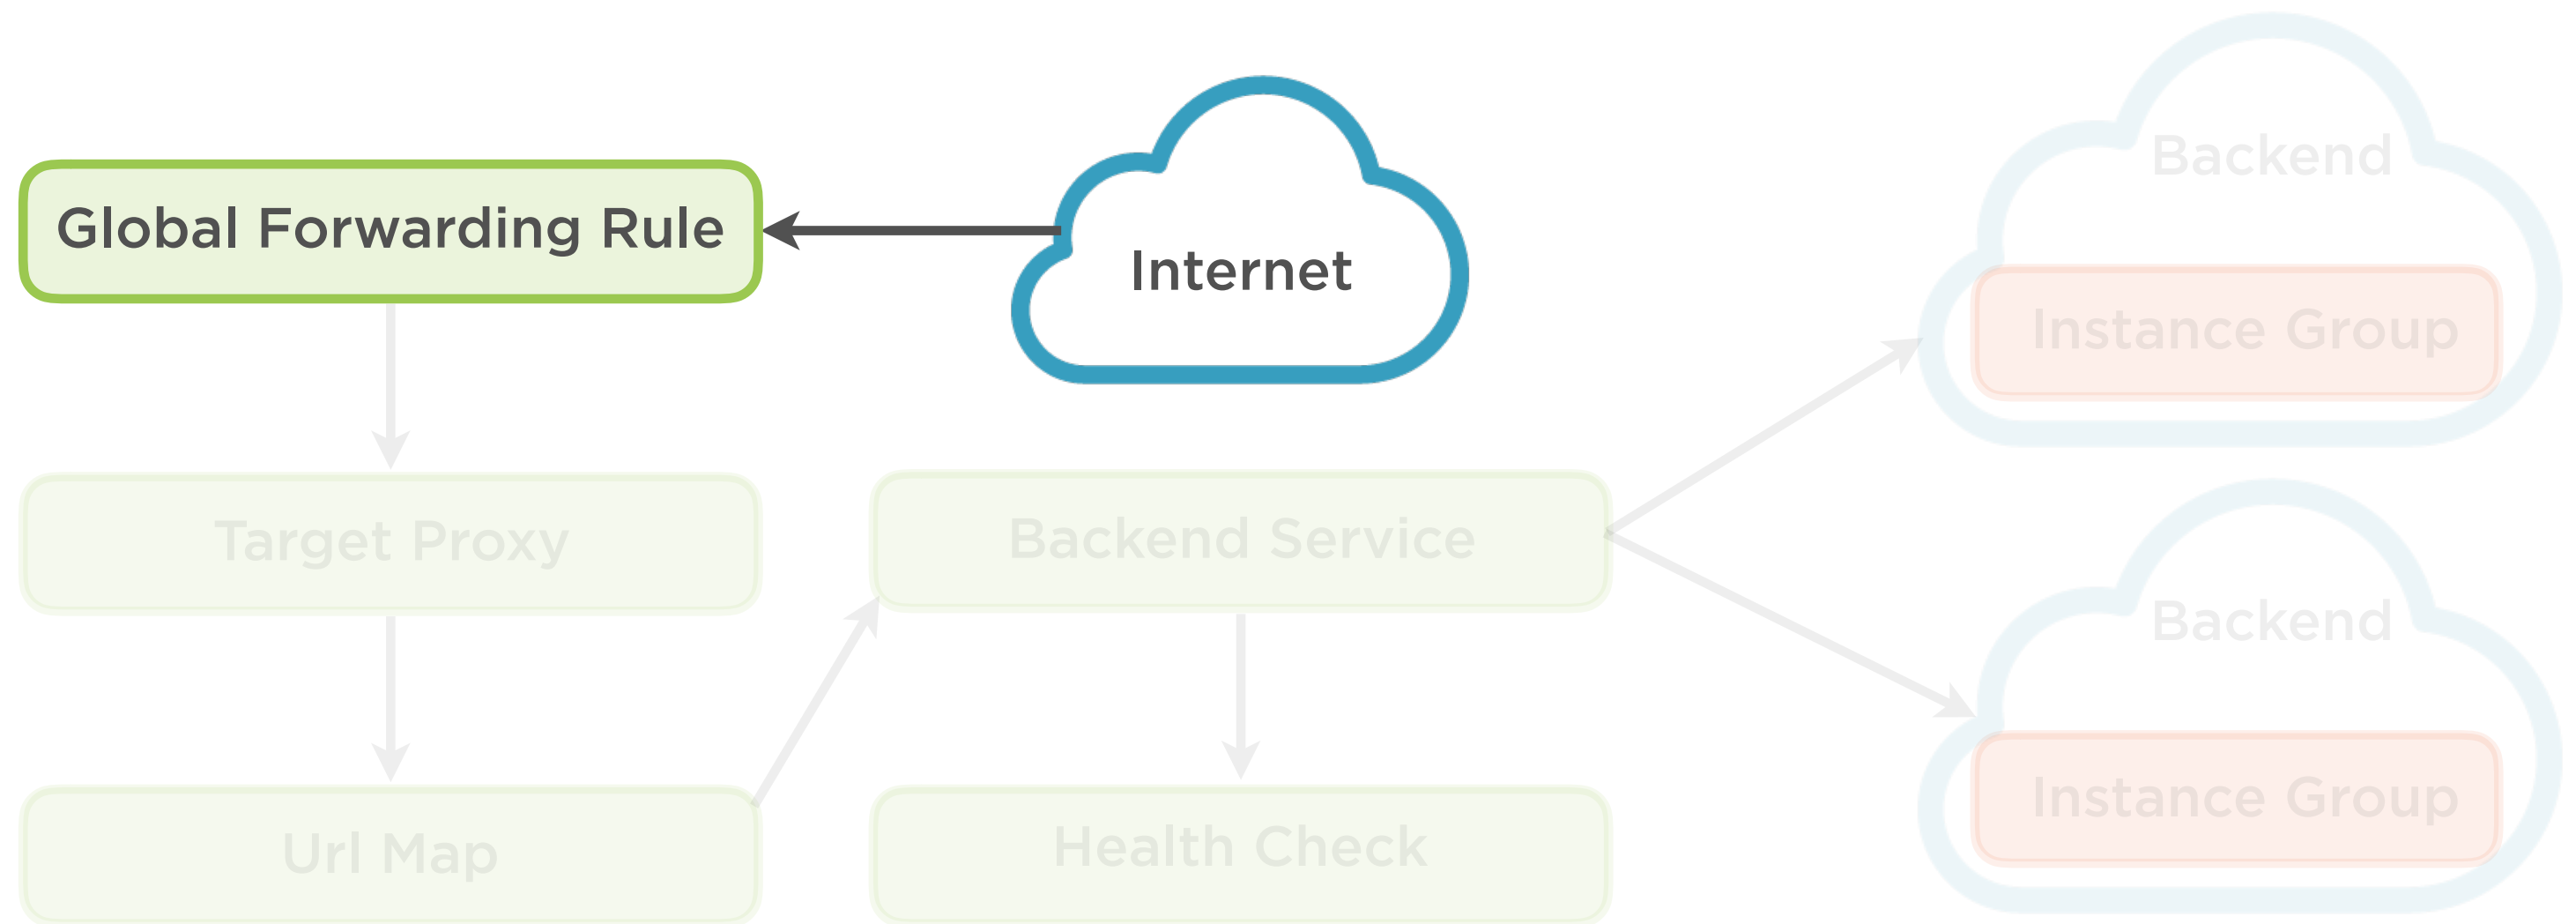
- Proximity to the user
- Requested URL
- Or both.

# HTTP(S) Load Balancing



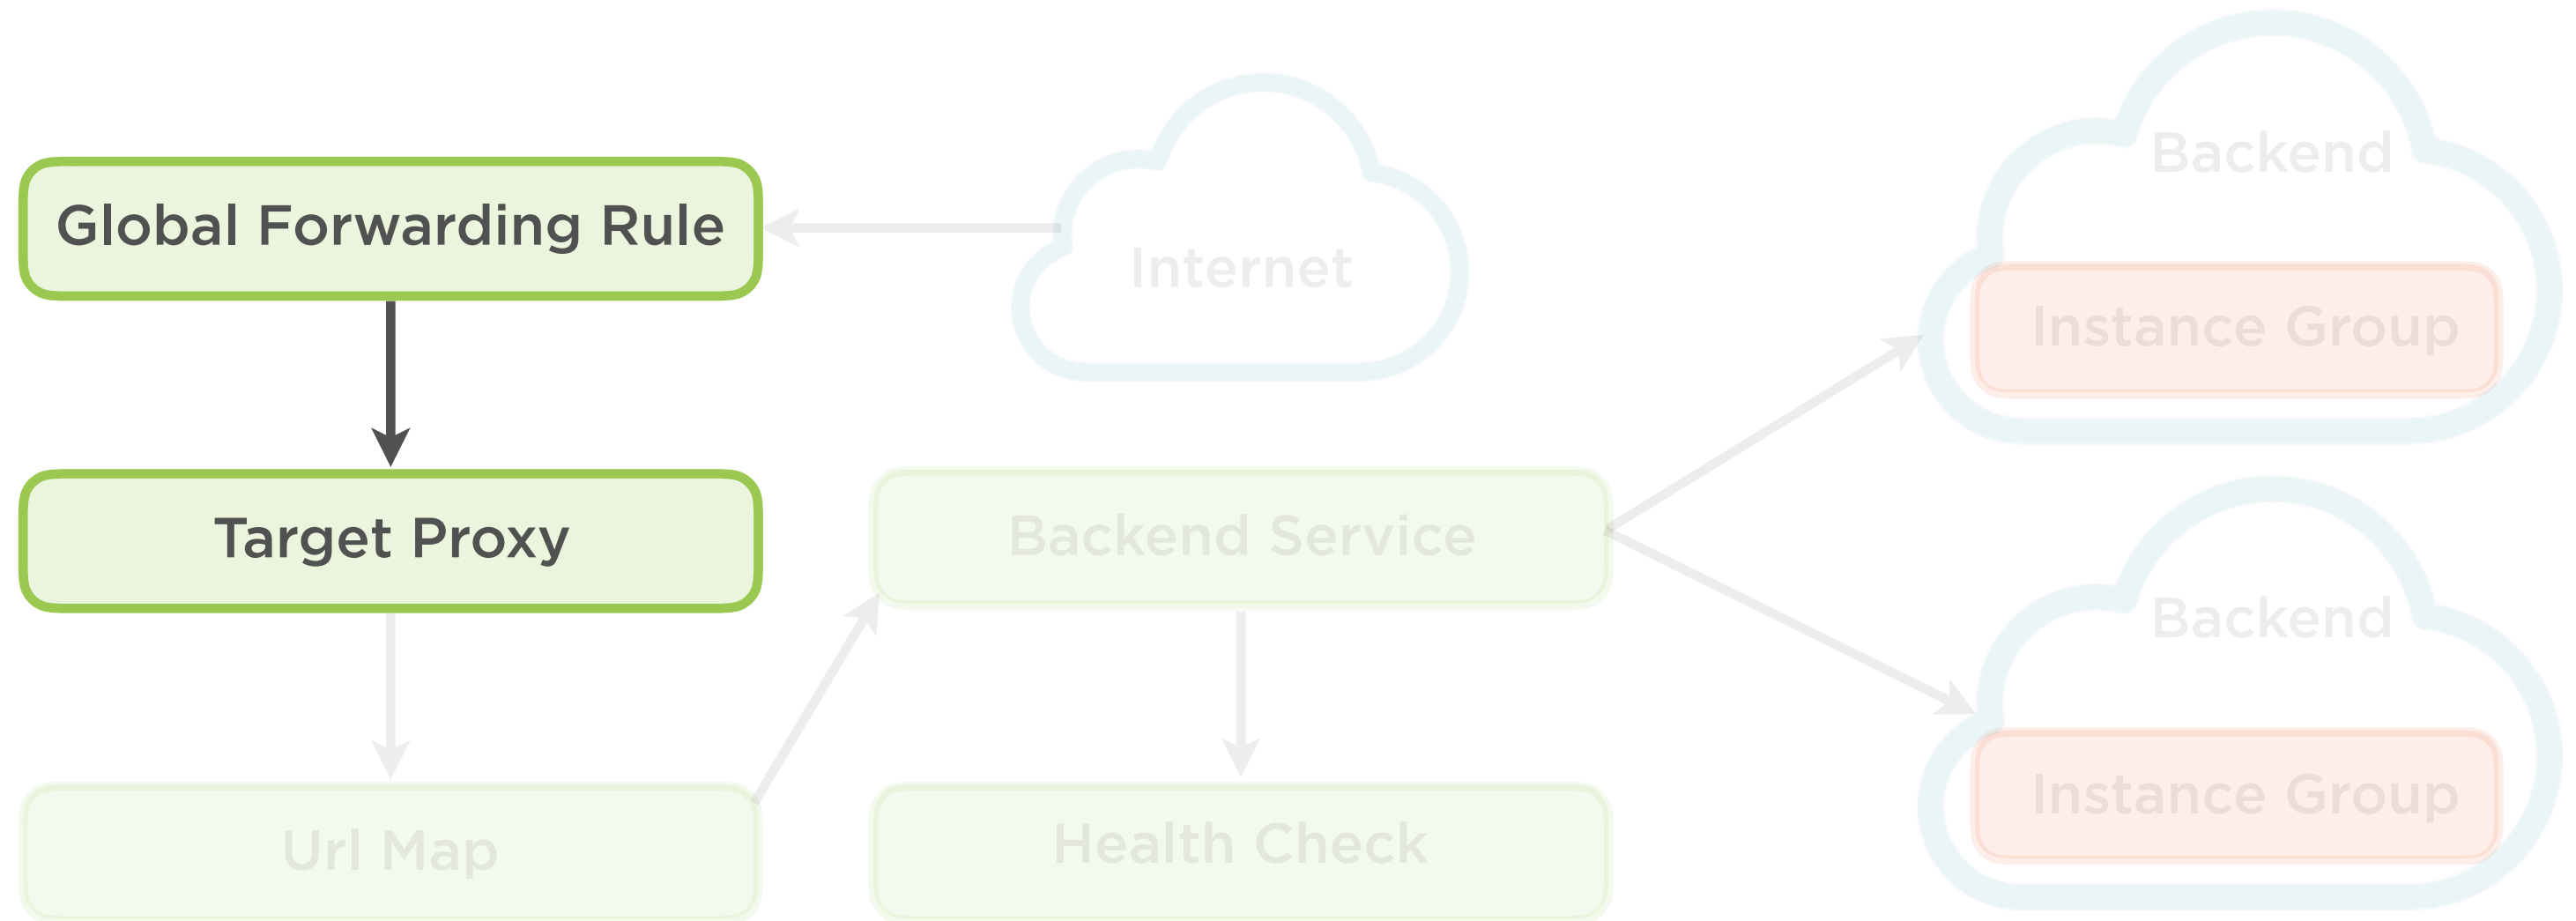
A global, external load balancing service offered on the GCP

# Global Forwarding Rule



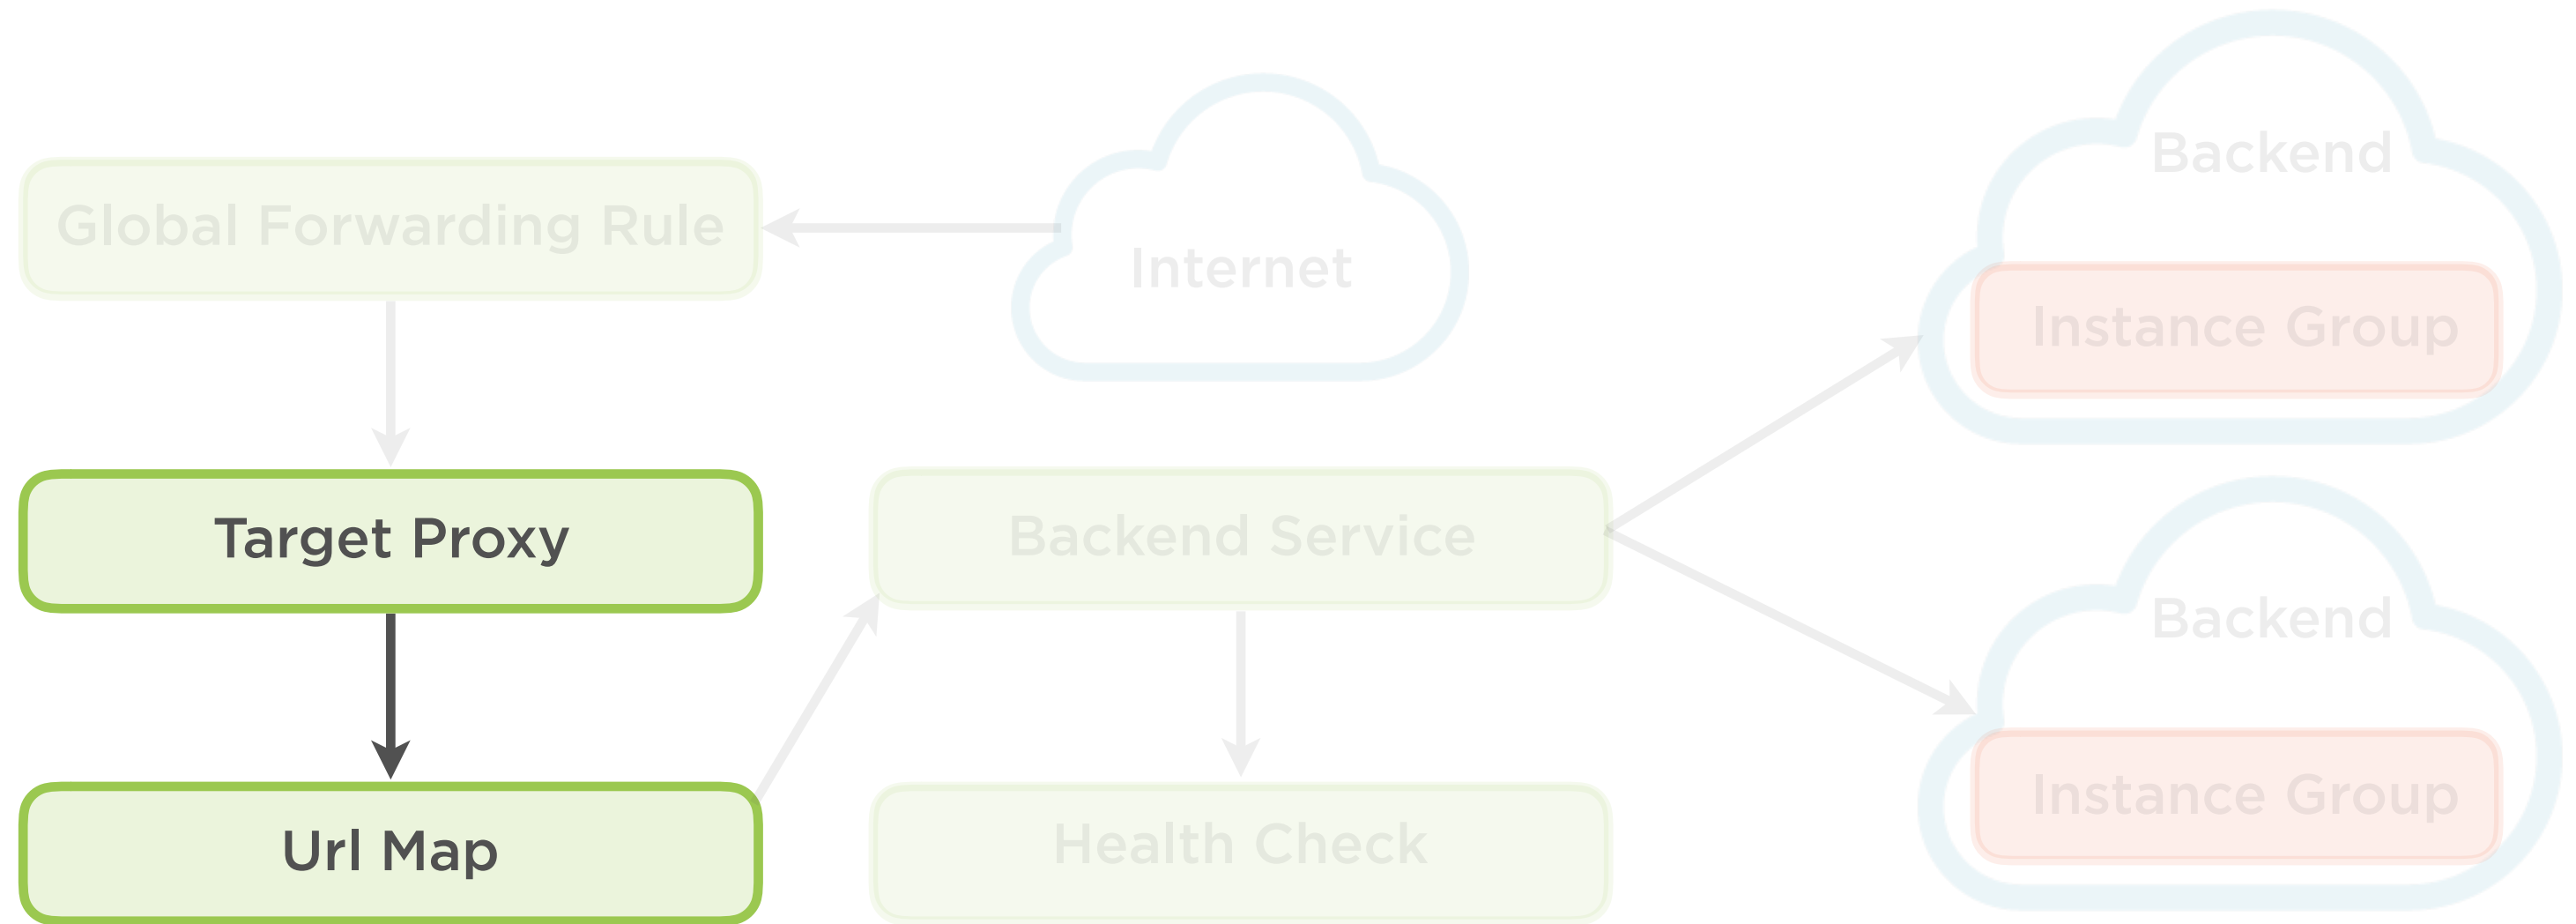
**Traffic from the internet is sent to a global forwarding rule - this rule determines which proxy the traffic should be directed to**

# Target Proxy




**The global forwarding rule directs incoming requests to a target HTTP proxy**

# URL Map



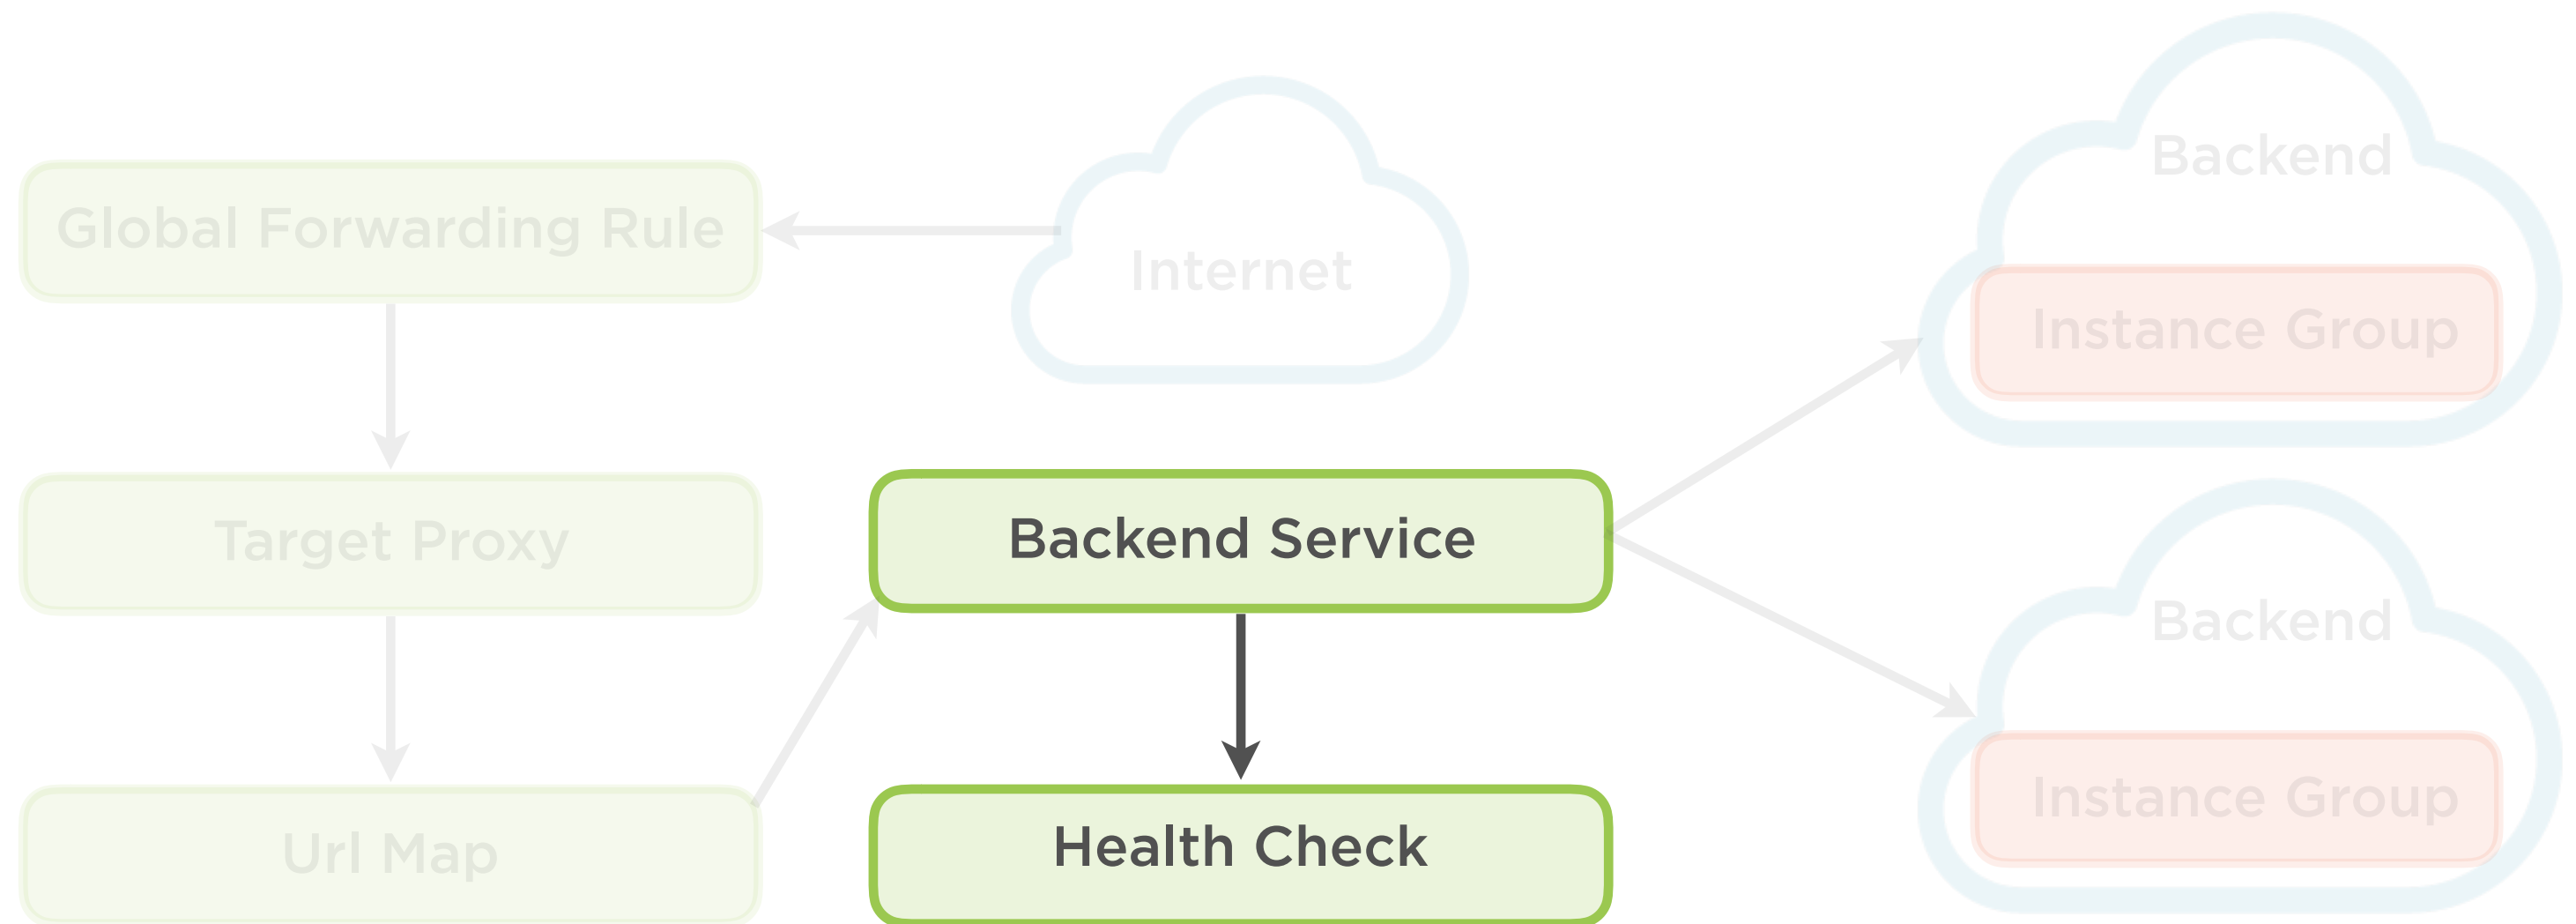
**The target HTTP proxy checks each request against a URL map to determine the appropriate backend service for the request**

# Backend Service



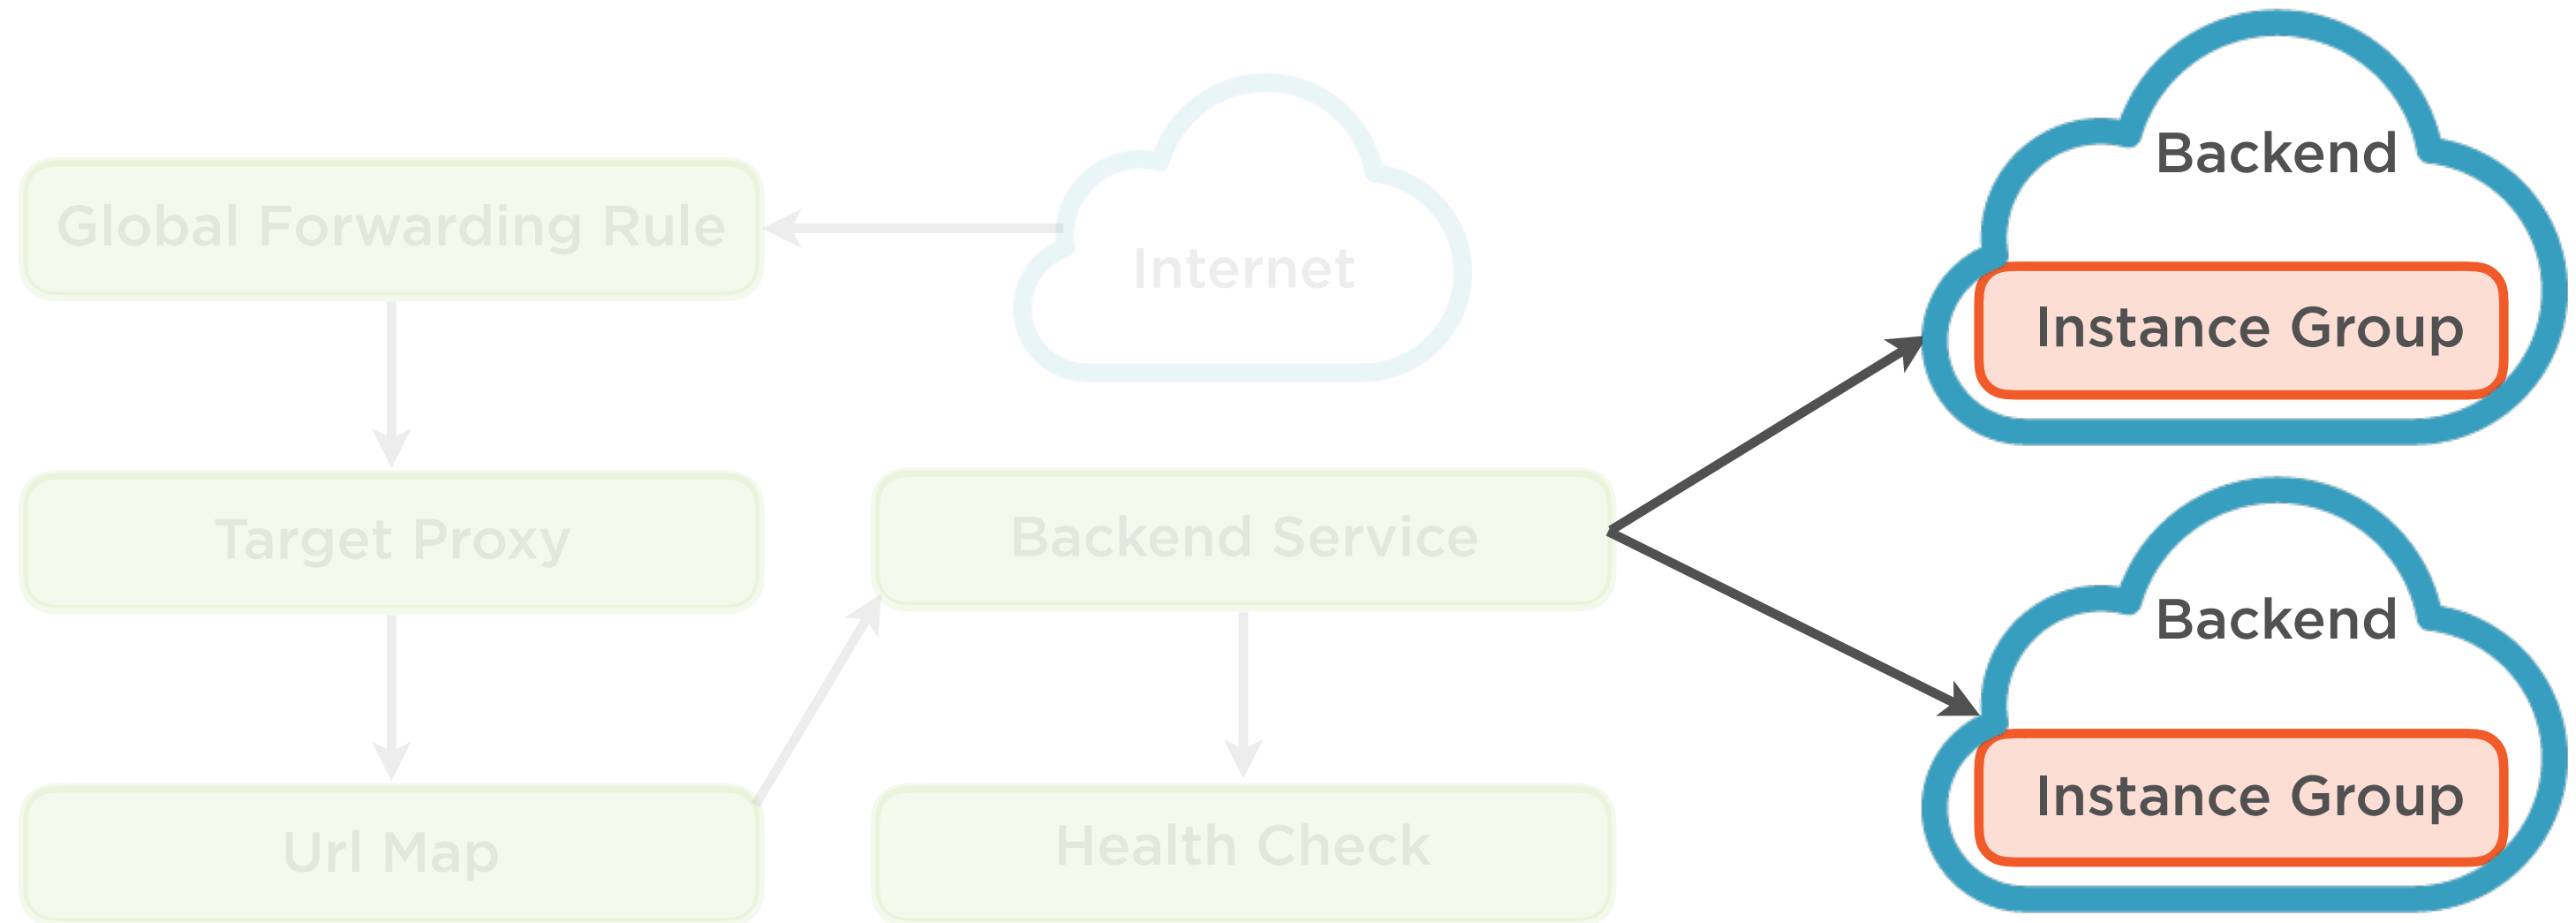
**The backend service directs each request to an appropriate backend based on serving capacity, zone, and instance health of its attached backends**

# Health Check



**The health of each backend instance is verified using either an HTTP health check or an HTTPS health check - if HTTPS, request is encrypted**

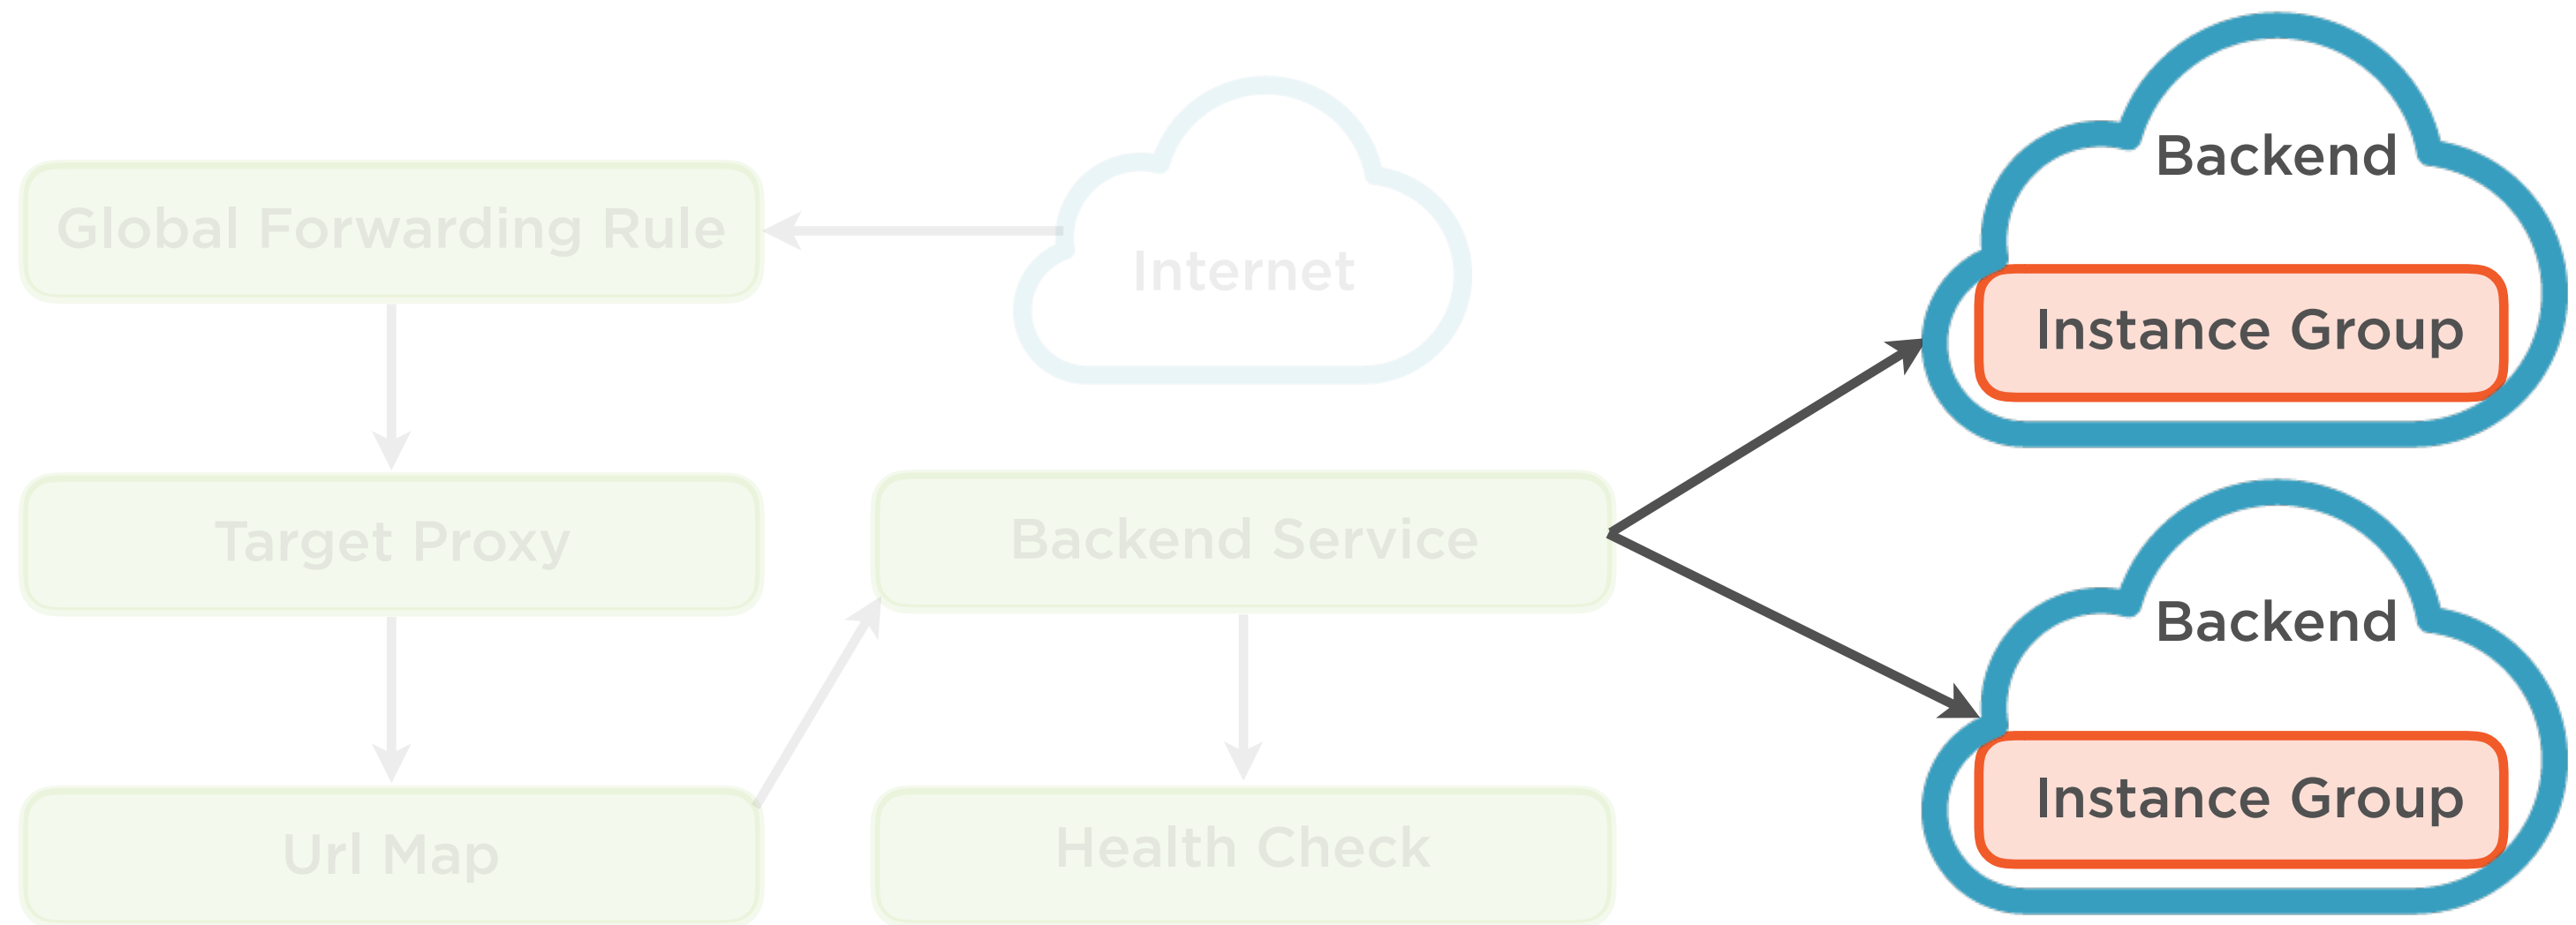
# Traffic Distribution to Backends



**Actual request distribution can happen based on CPU utilization, requests per instance**

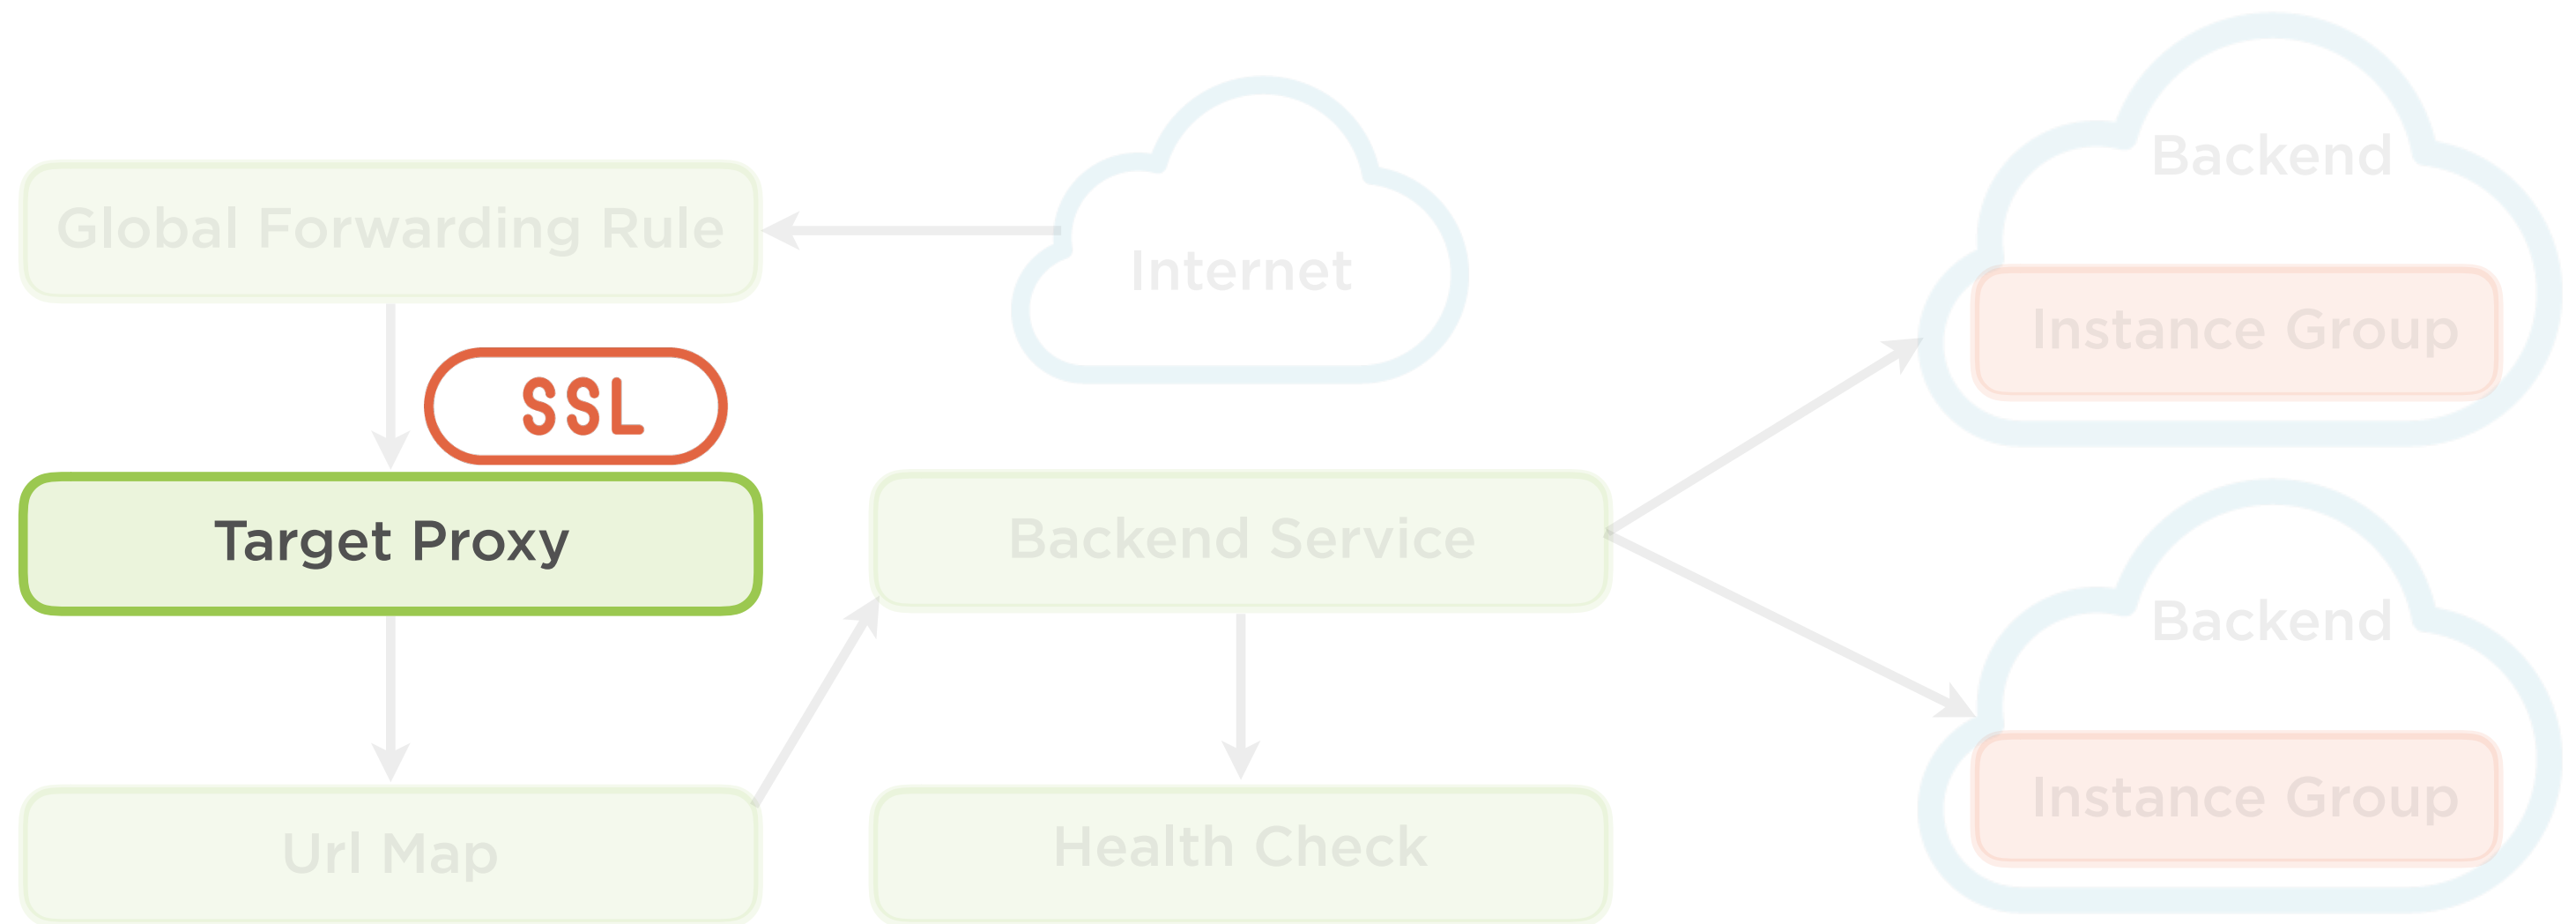


# Autoscaling and Autohealing



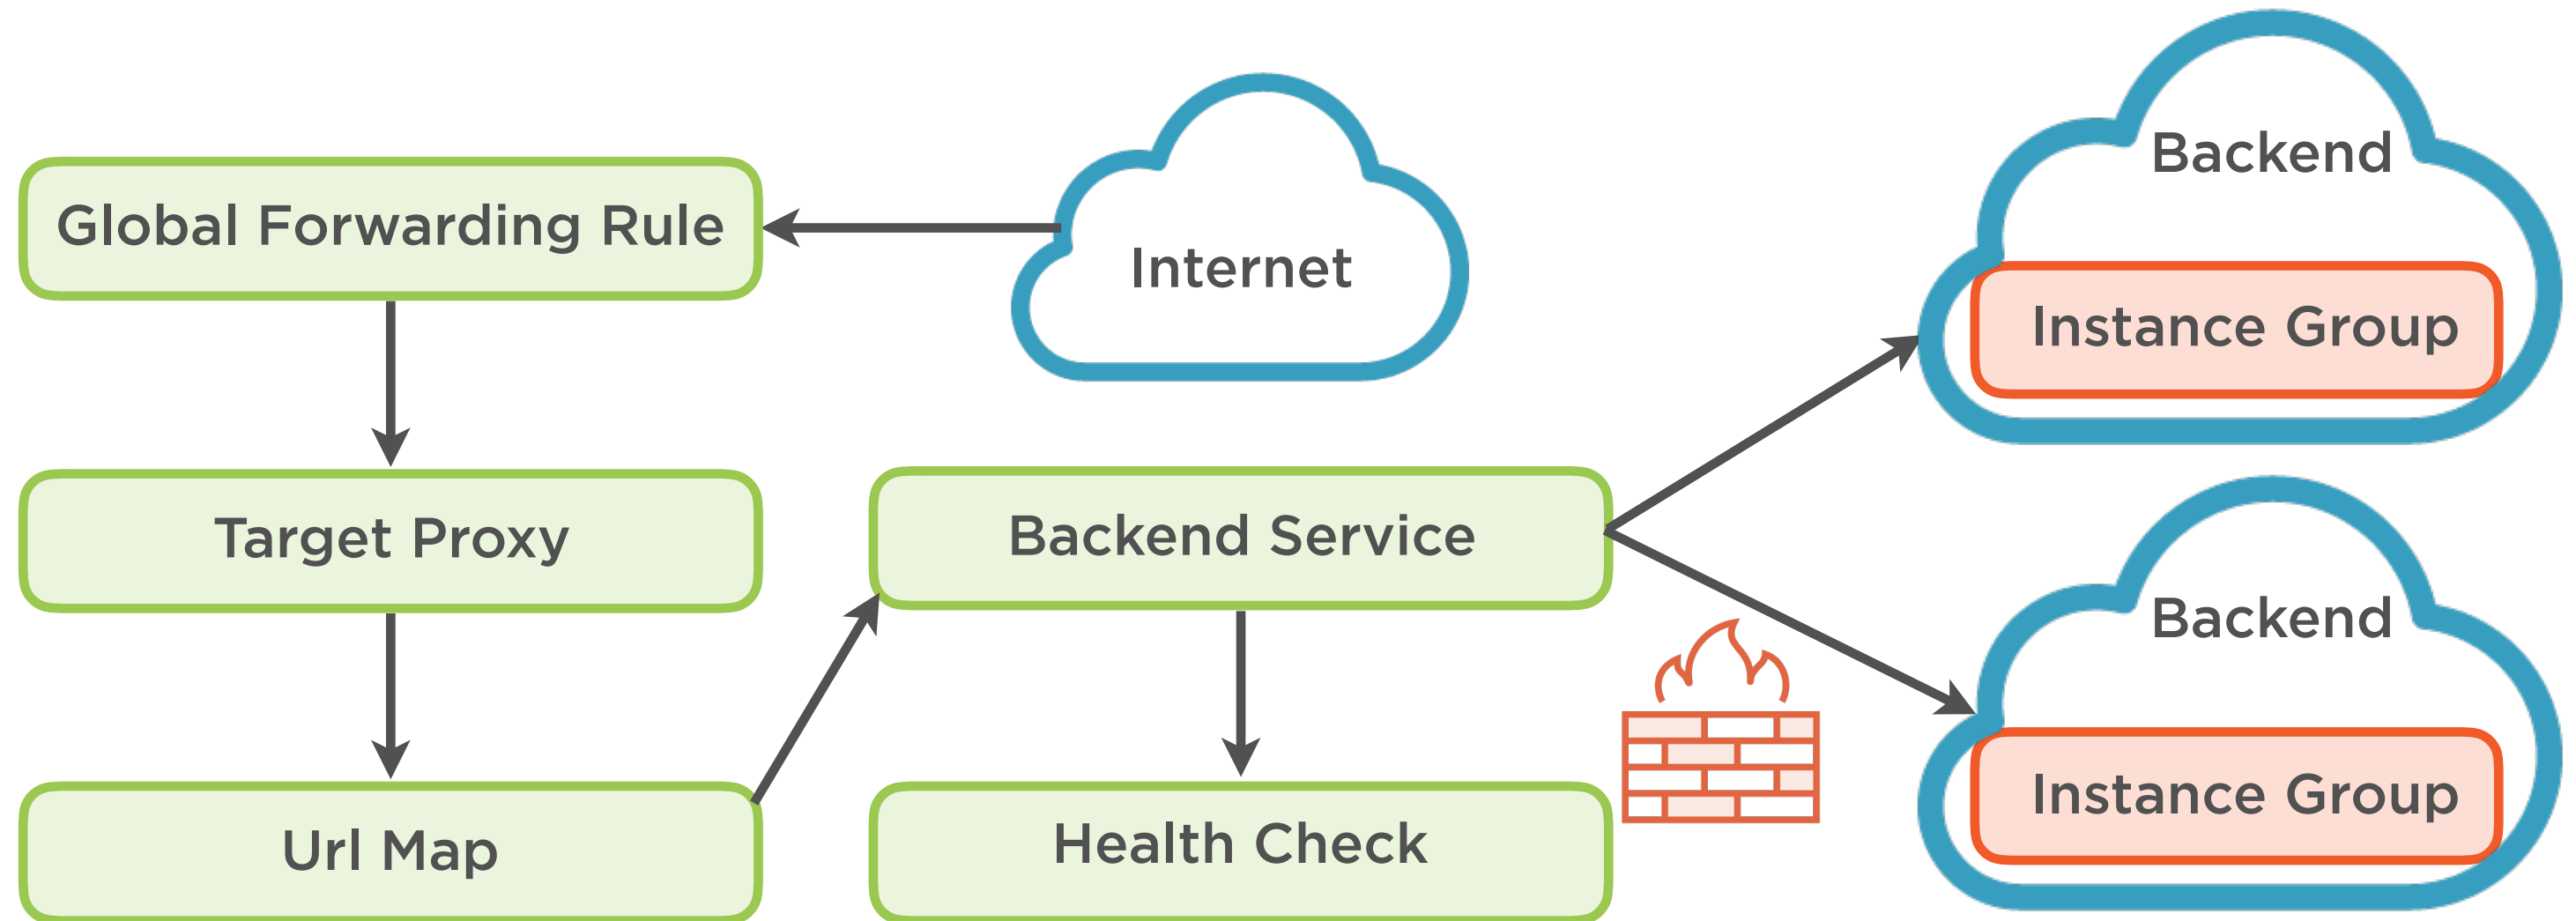
**If backends are managed instance groups instances  
can scale as the traffic scales**

# SSL Termination




**HTTPS load balancing requires the target proxy to have a signed certificate to terminate the SSL connection**

# Firewall Rules



**Must create firewall rules to allow requests from load balancer and health checker to get through to the instances**

# HTTP(S) Load Balancing



# Session Affinity



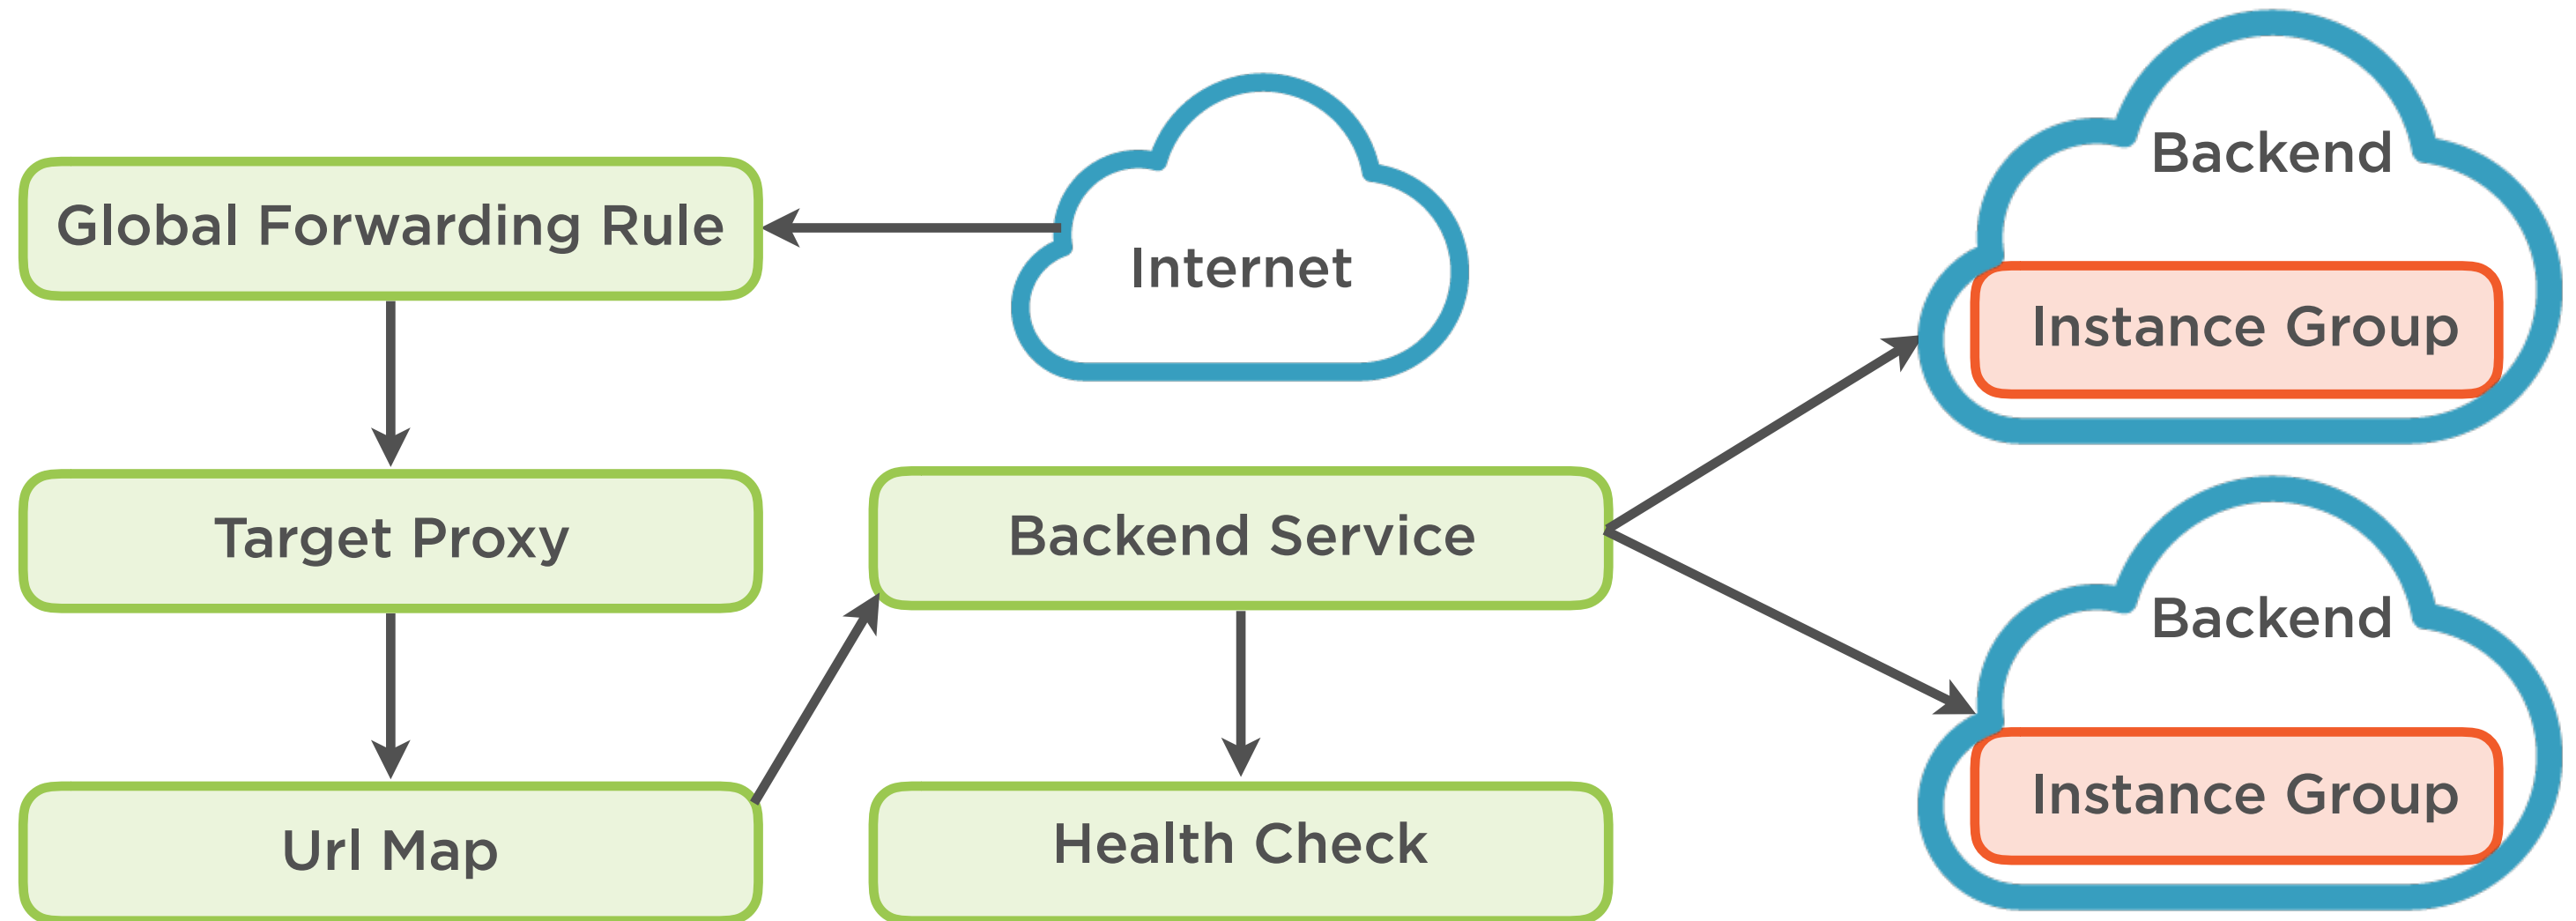
**Session affinity: All requests from same client to same server based on either**

- Client IP
- Cookie


# HTTP(S) Load Balancing Components

---

# HTTP(S) Load Balancing




# Global Forwarding Rule



**Traffic from the internet is sent to a global forwarding rule - this rule determines which proxy the traffic should be directed to**



# Global Forwarding Rule




**Route traffic by IP address, port and protocol to a load balancing proxy**

**Can only be used with global load balancing HTTP(S), SSL Proxy and TCP Proxy**

**Regional forwarding rules can be used with regional load balancing and individual instances**

# Target Proxy



**The global forwarding rule directs incoming requests to a target HTTP proxy**

# Target Proxy



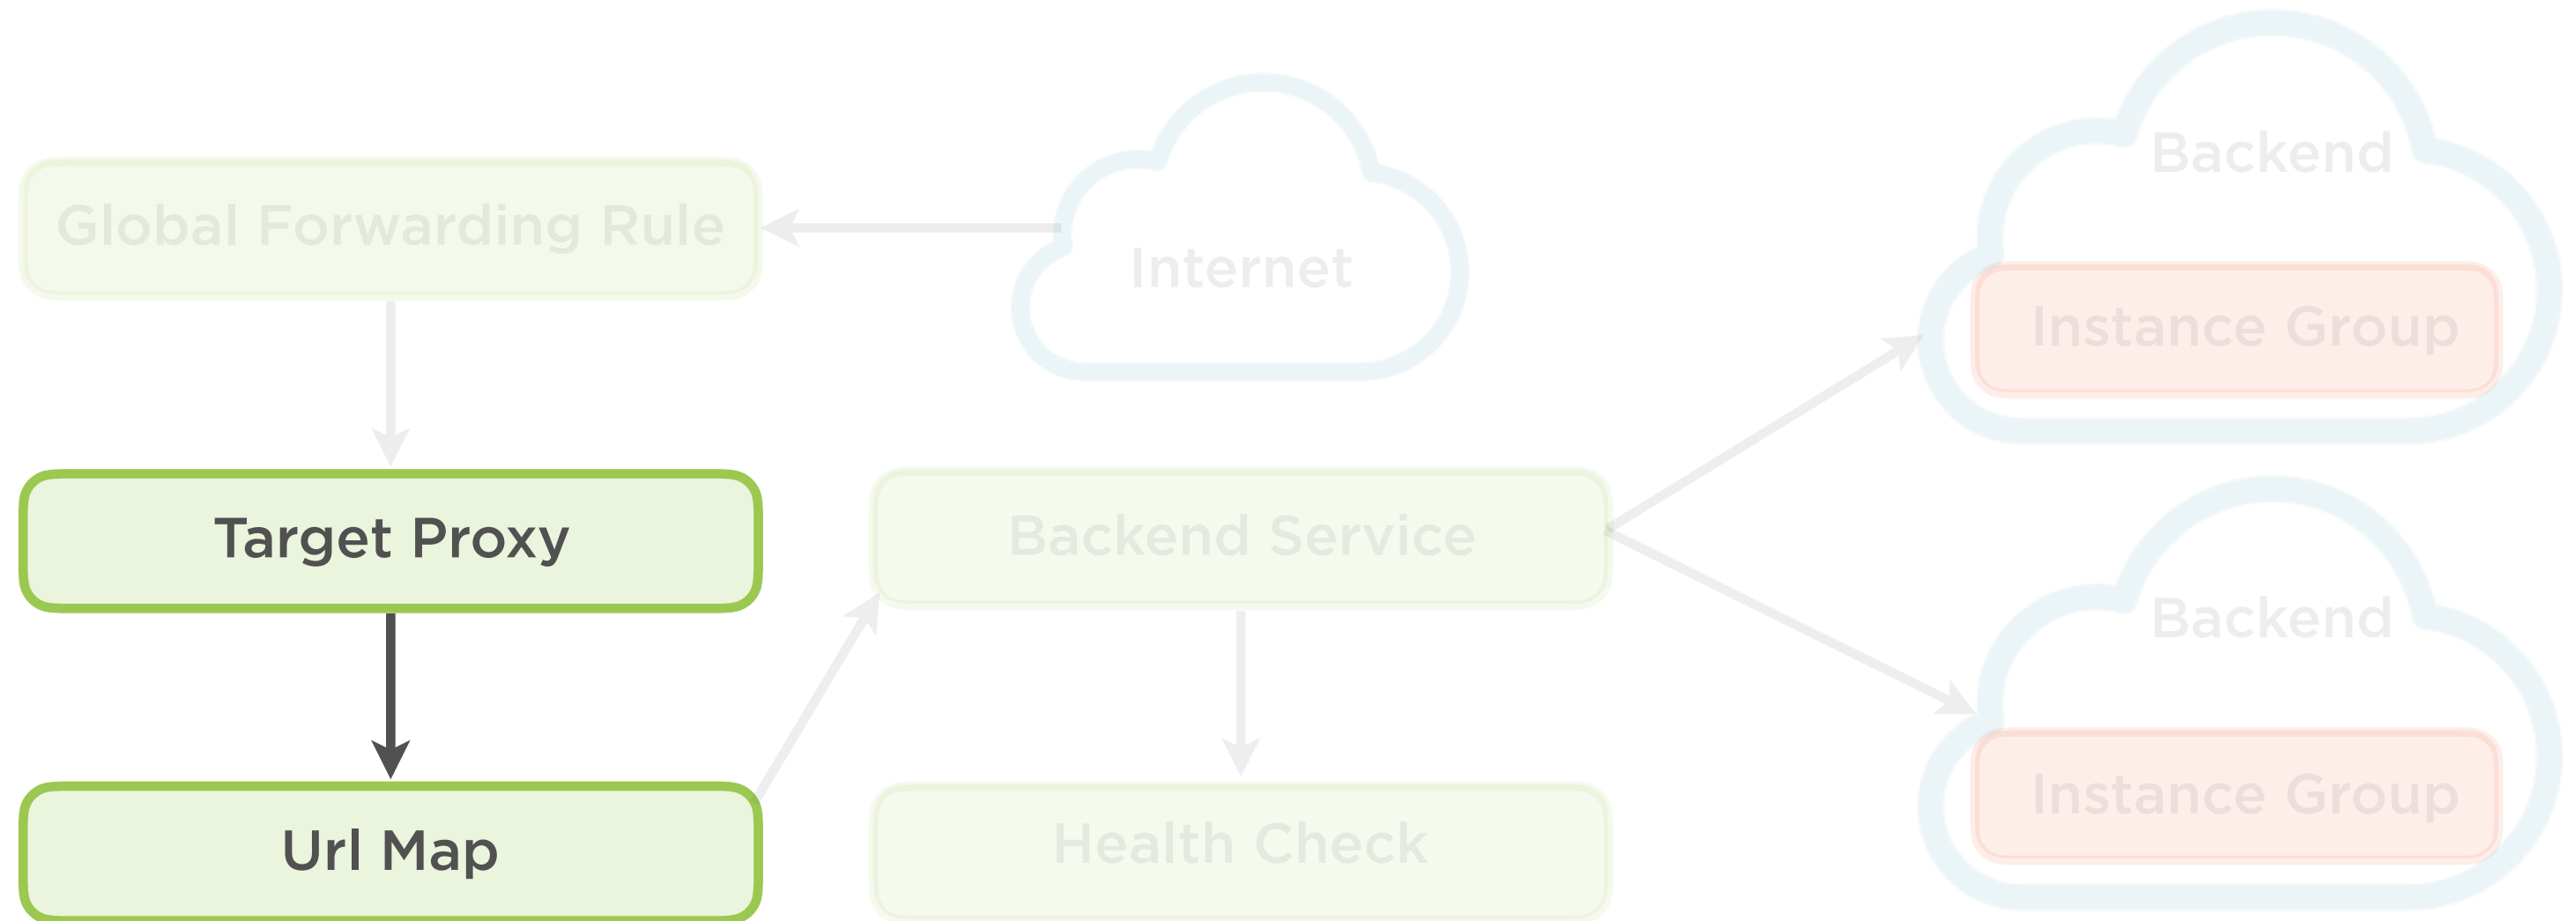
**Route the incoming requests to a URL map to determine where they should be sent**

**Specific to a protocol (HTTP, HTTPS, SSL and TCP)**

**Should have a SSL certificate if it terminates HTTPS connections**

**Can connect to backend services via HTTP or HTTPS**

# URL Map



**The target HTTP proxy checks each request against a URL map to determine the appropriate backend service for the request**

# URL Map

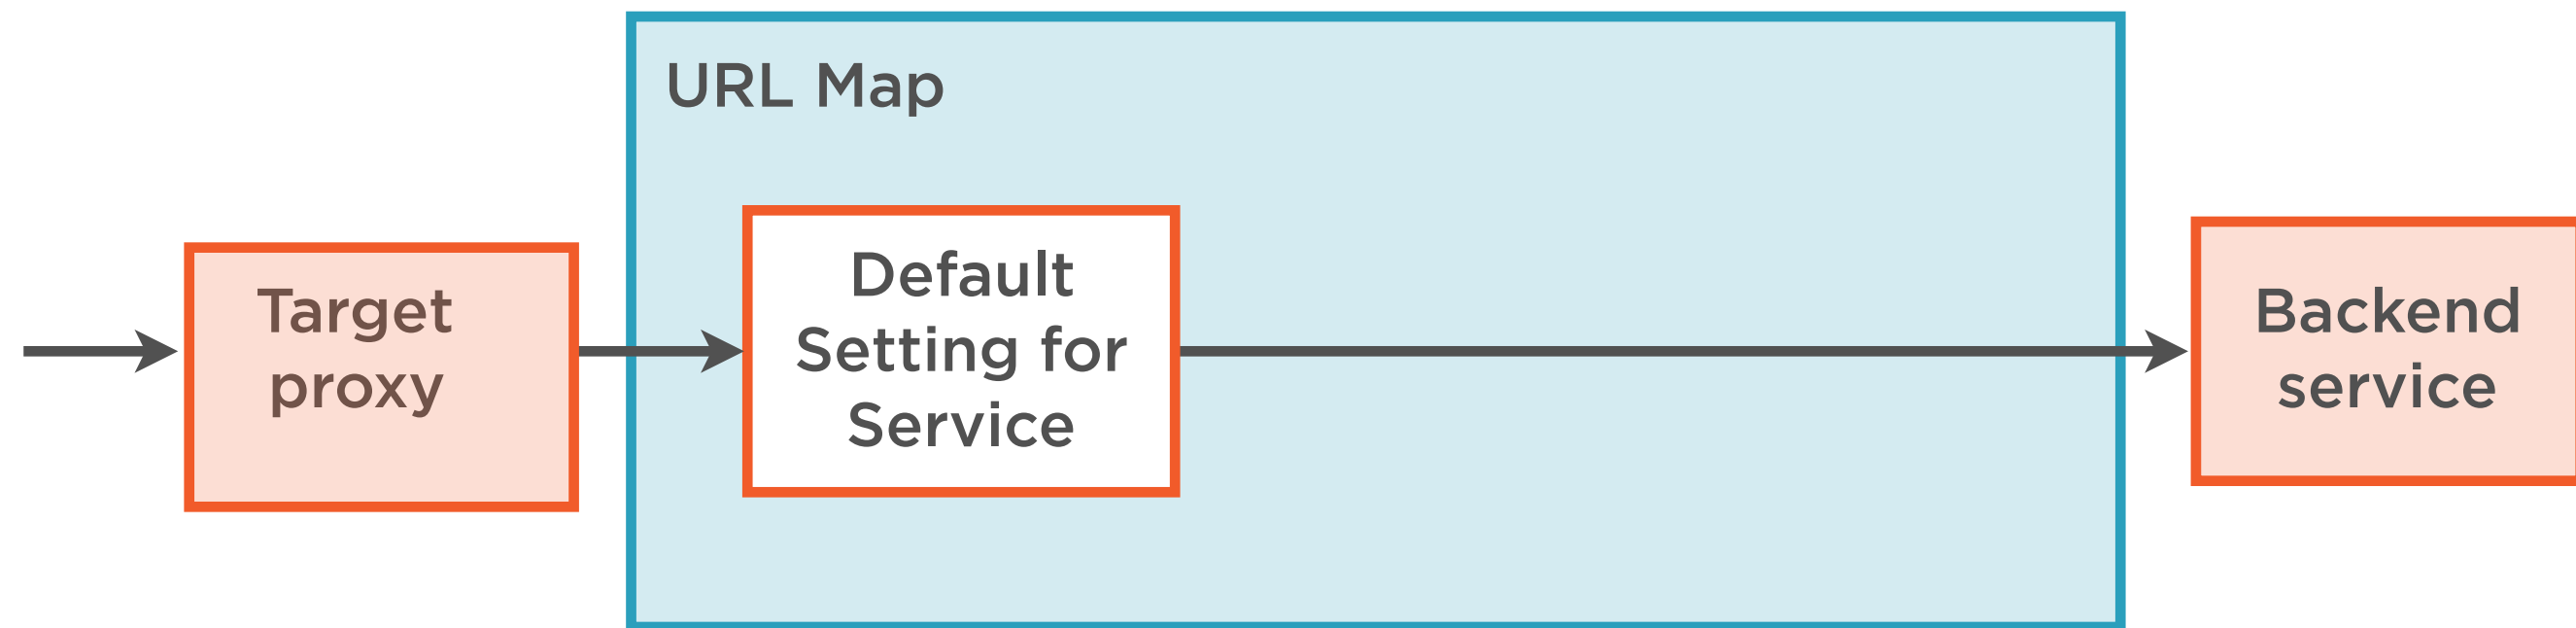


**Used to direct traffic to different instances based on the incoming URL**

- <http://www.example.com/audio> -> backend service 1
- <http://www.example.com/video> -> backend service 2

# URL Map

URL map with no rules except default

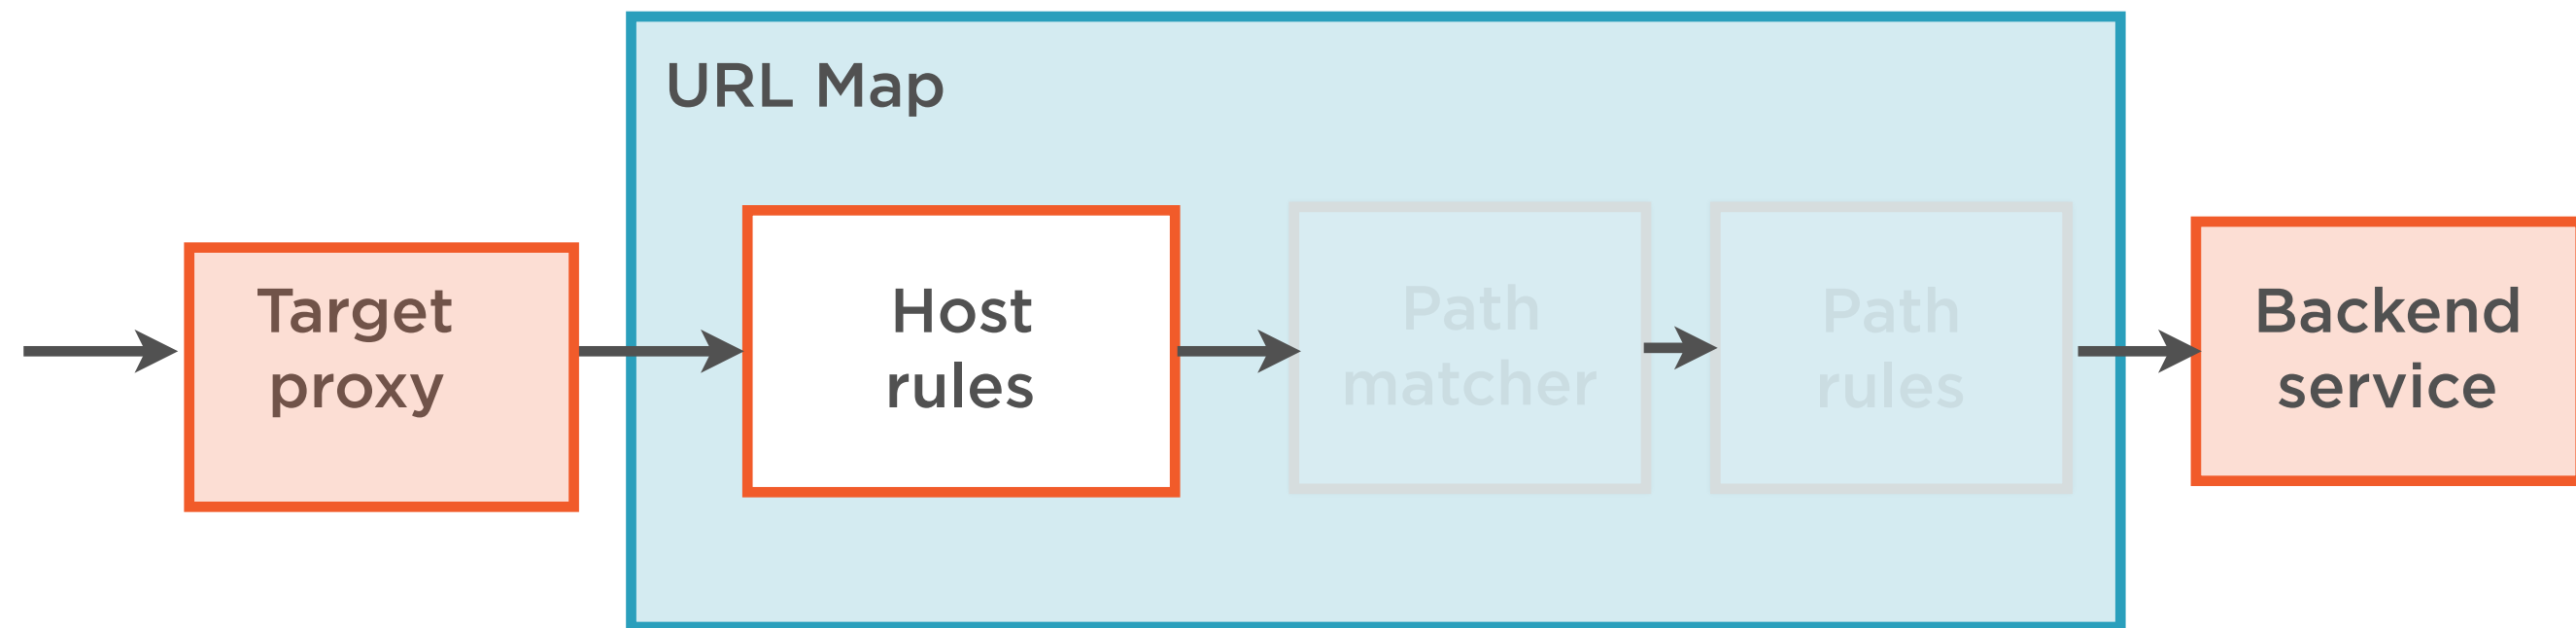


All traffic sent to the same groups of instances

Only the `/*` path matcher is created automatically and directs all traffic to the same backend service

# URL Map

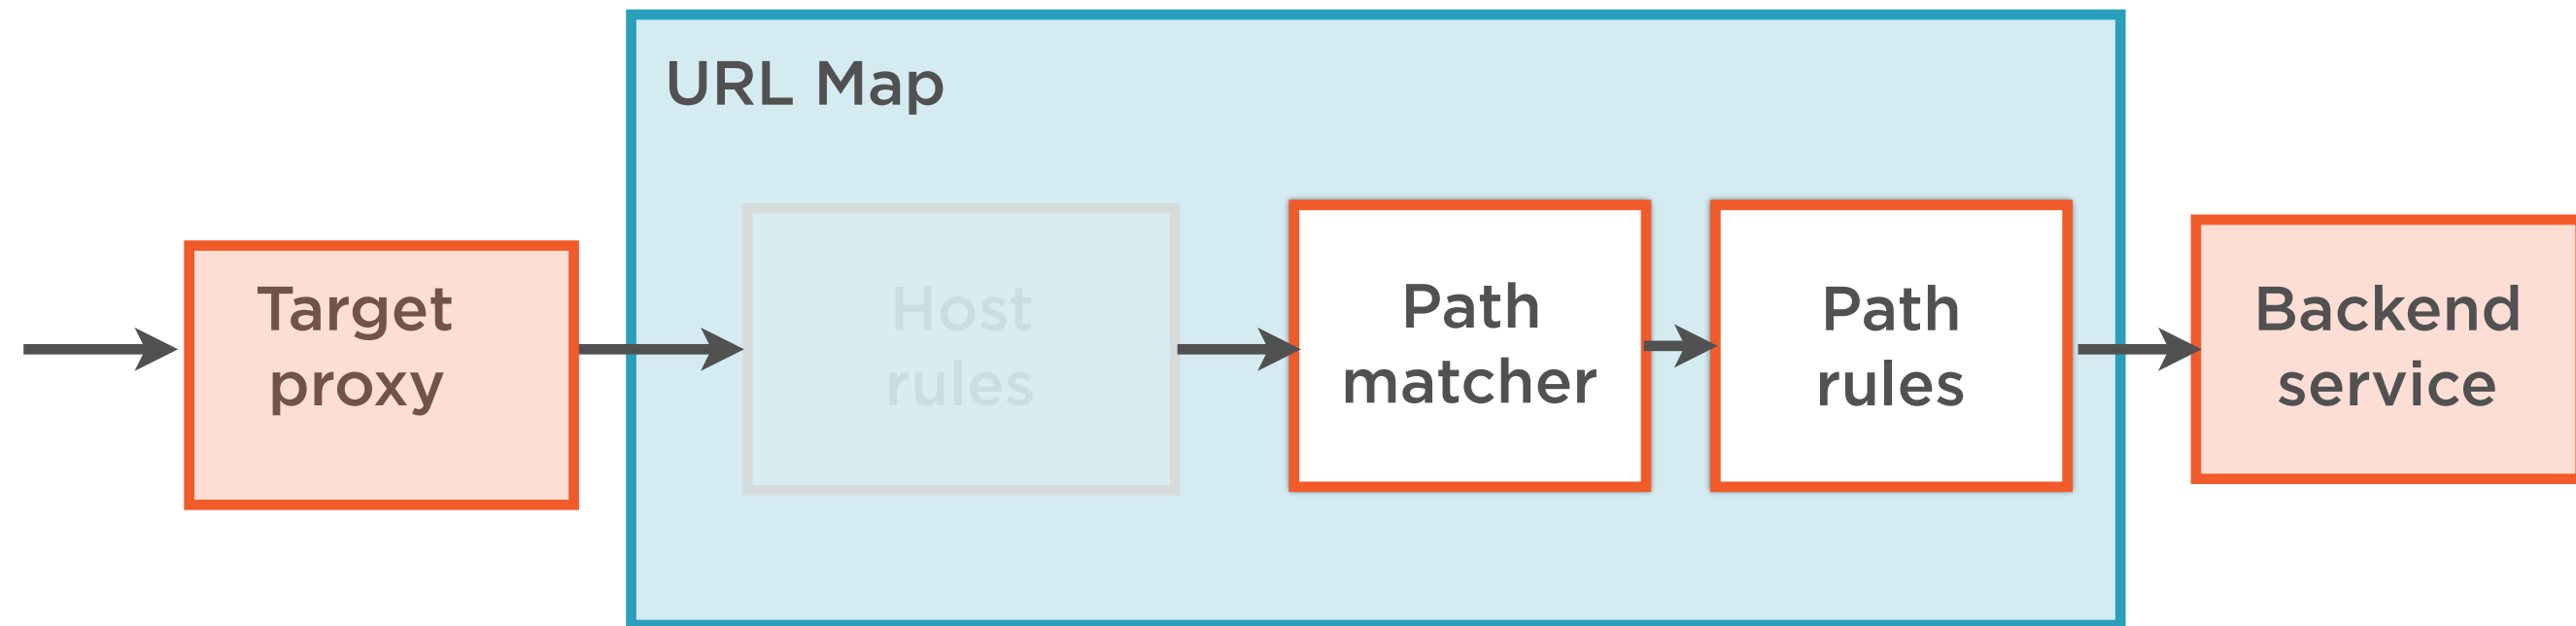
## Basic URL map flow



Host rules — example.com, customer.com

# URL Map

## Basic URL map flow

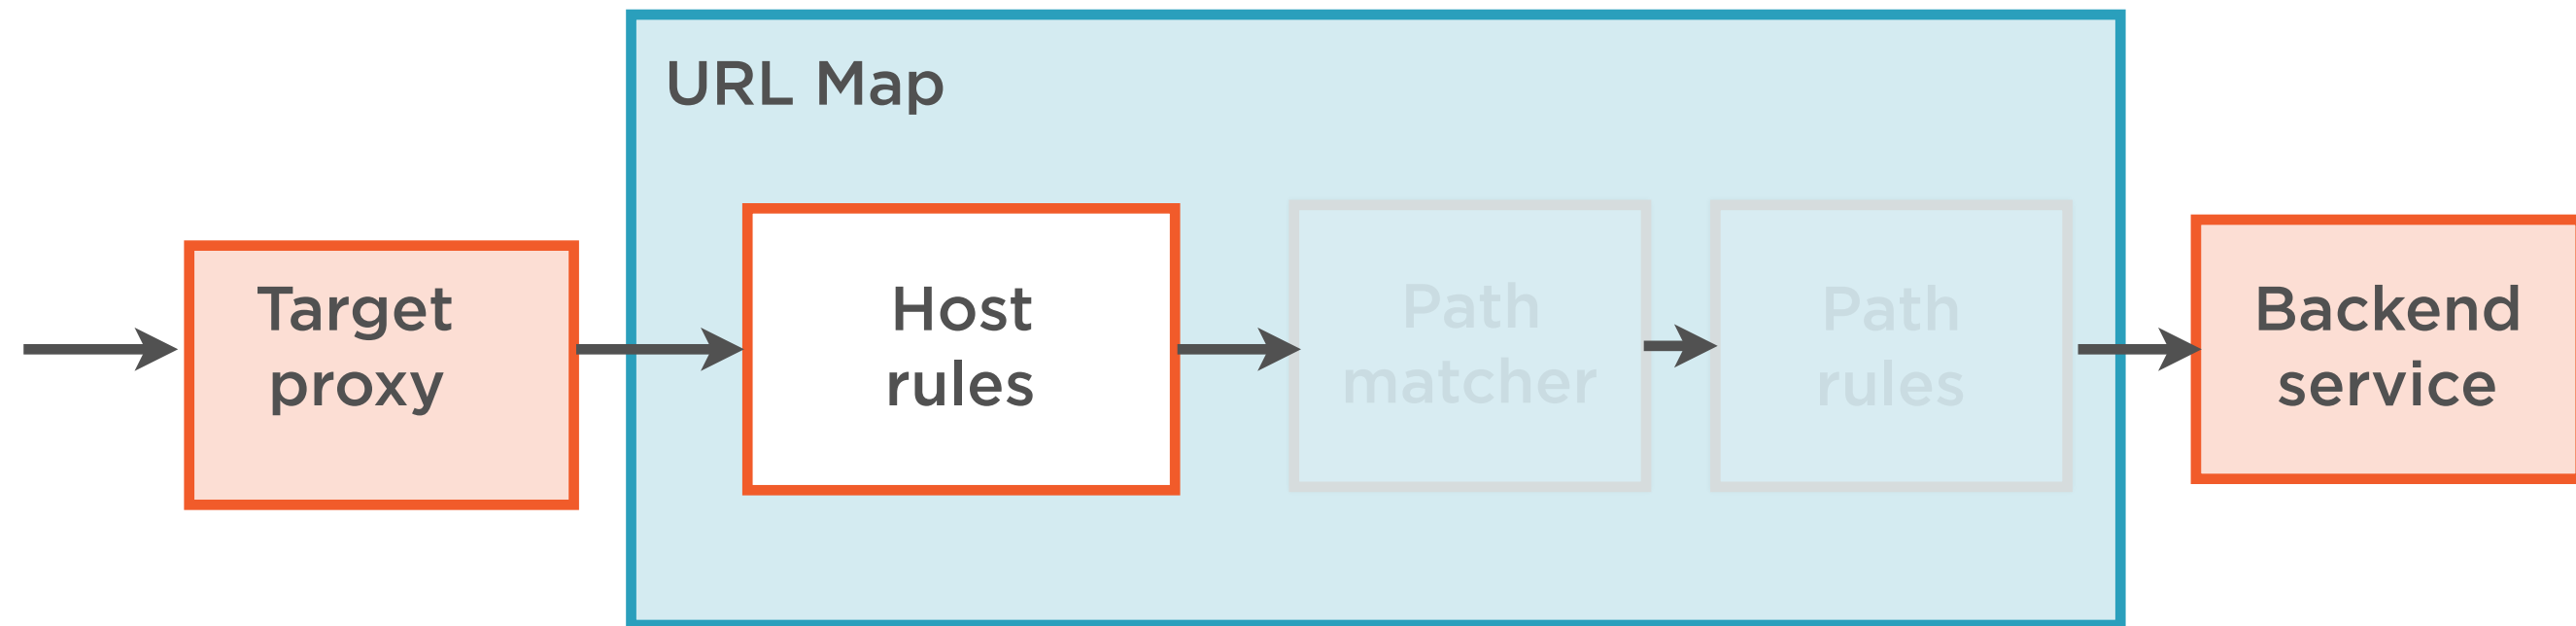


**Path rules** — /video, /video/hd, /video/sd



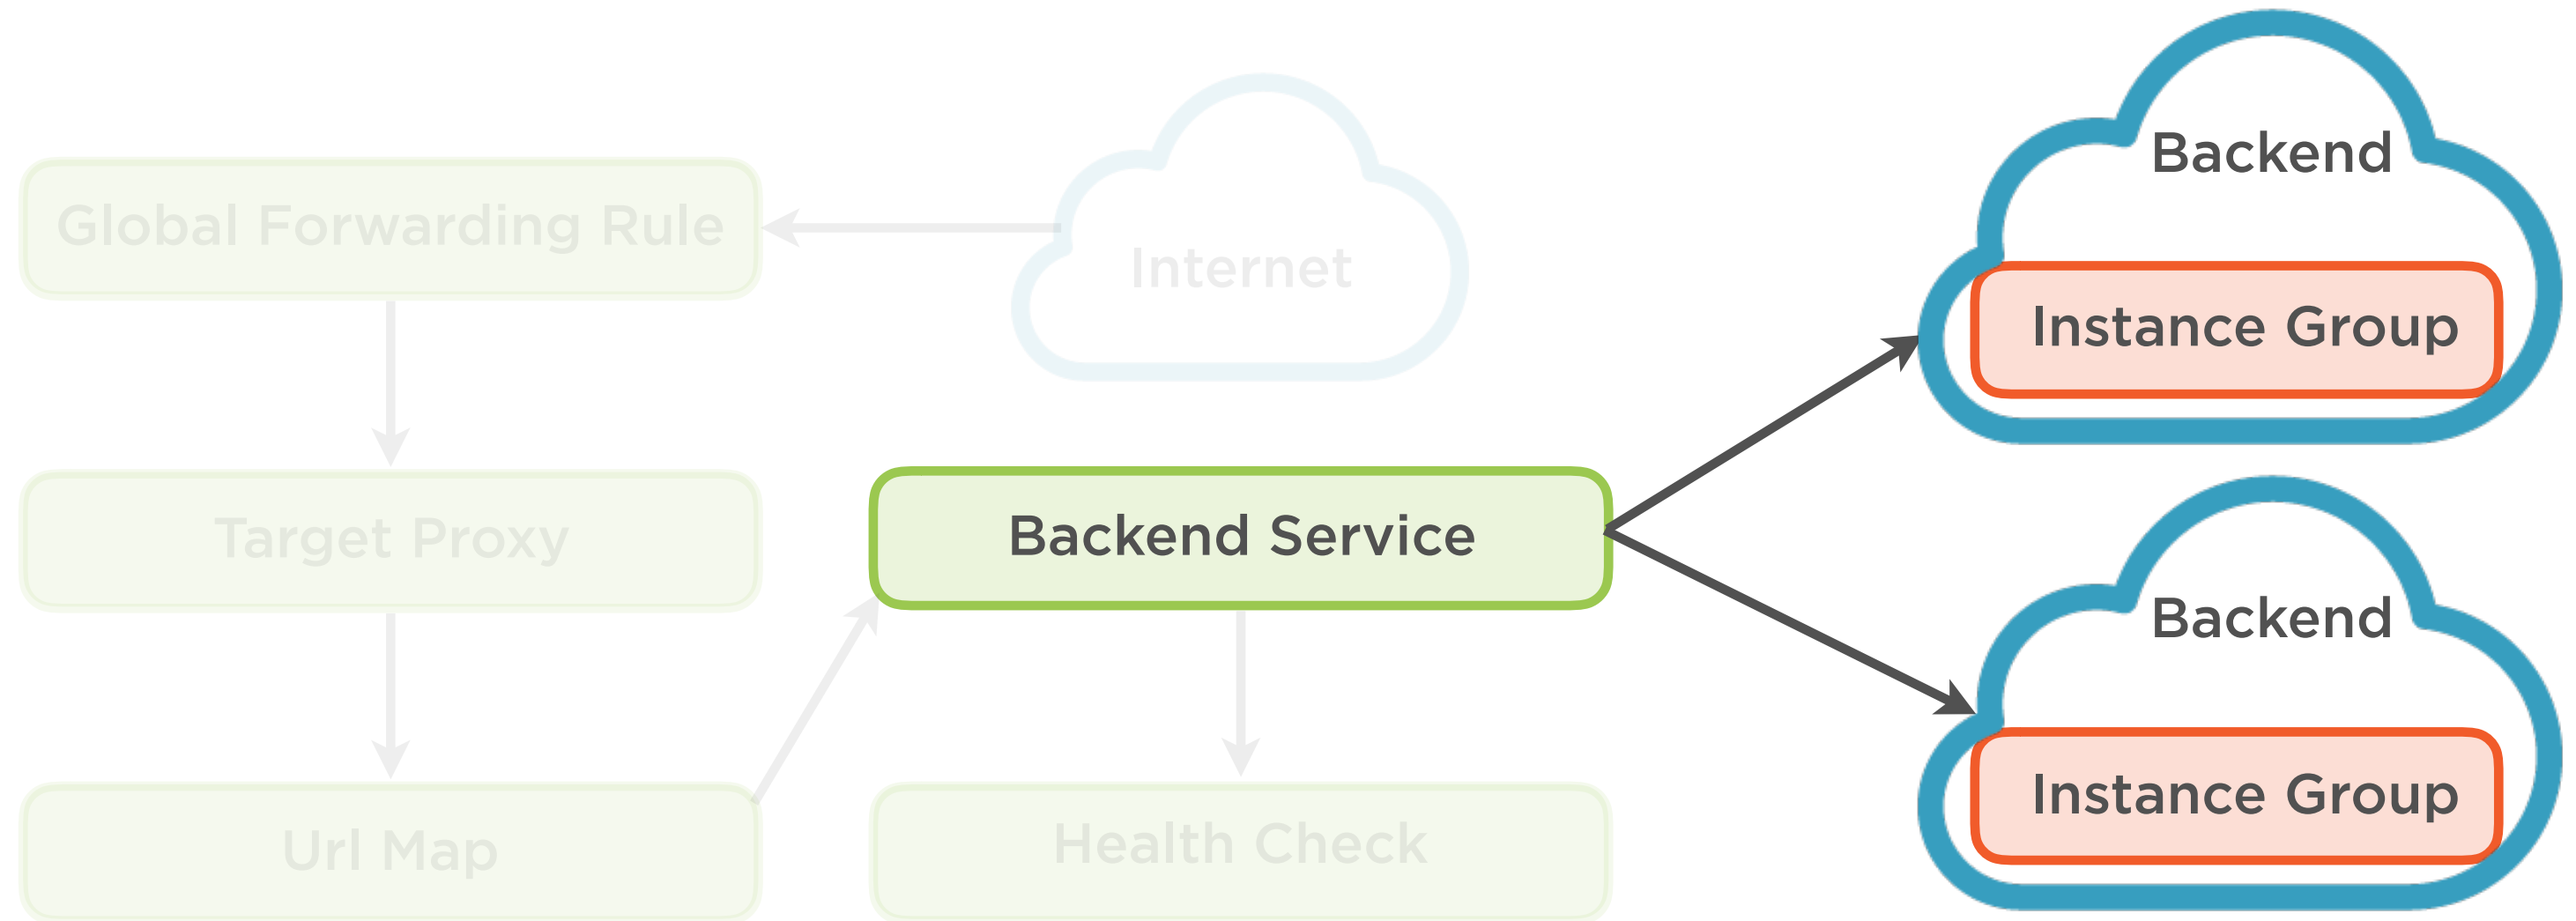
# URL Map

## Basic URL map flow



A default path matcher `/*` is created automatically. Traffic which does not match other path rules is sent to this default service

# Backend Service



**The backend service directs each request to an appropriate backend based on serving capacity, zone, and instance health of its attached backends**

# Backend Service



**Centralized service for managing backends**

**Backends contain instance groups which handle user requests**

**Knows which instances it can use, how much traffic they can handle**

**Monitors the health of backends and does not send traffic to unhealthy instances**

# Backend Service Components



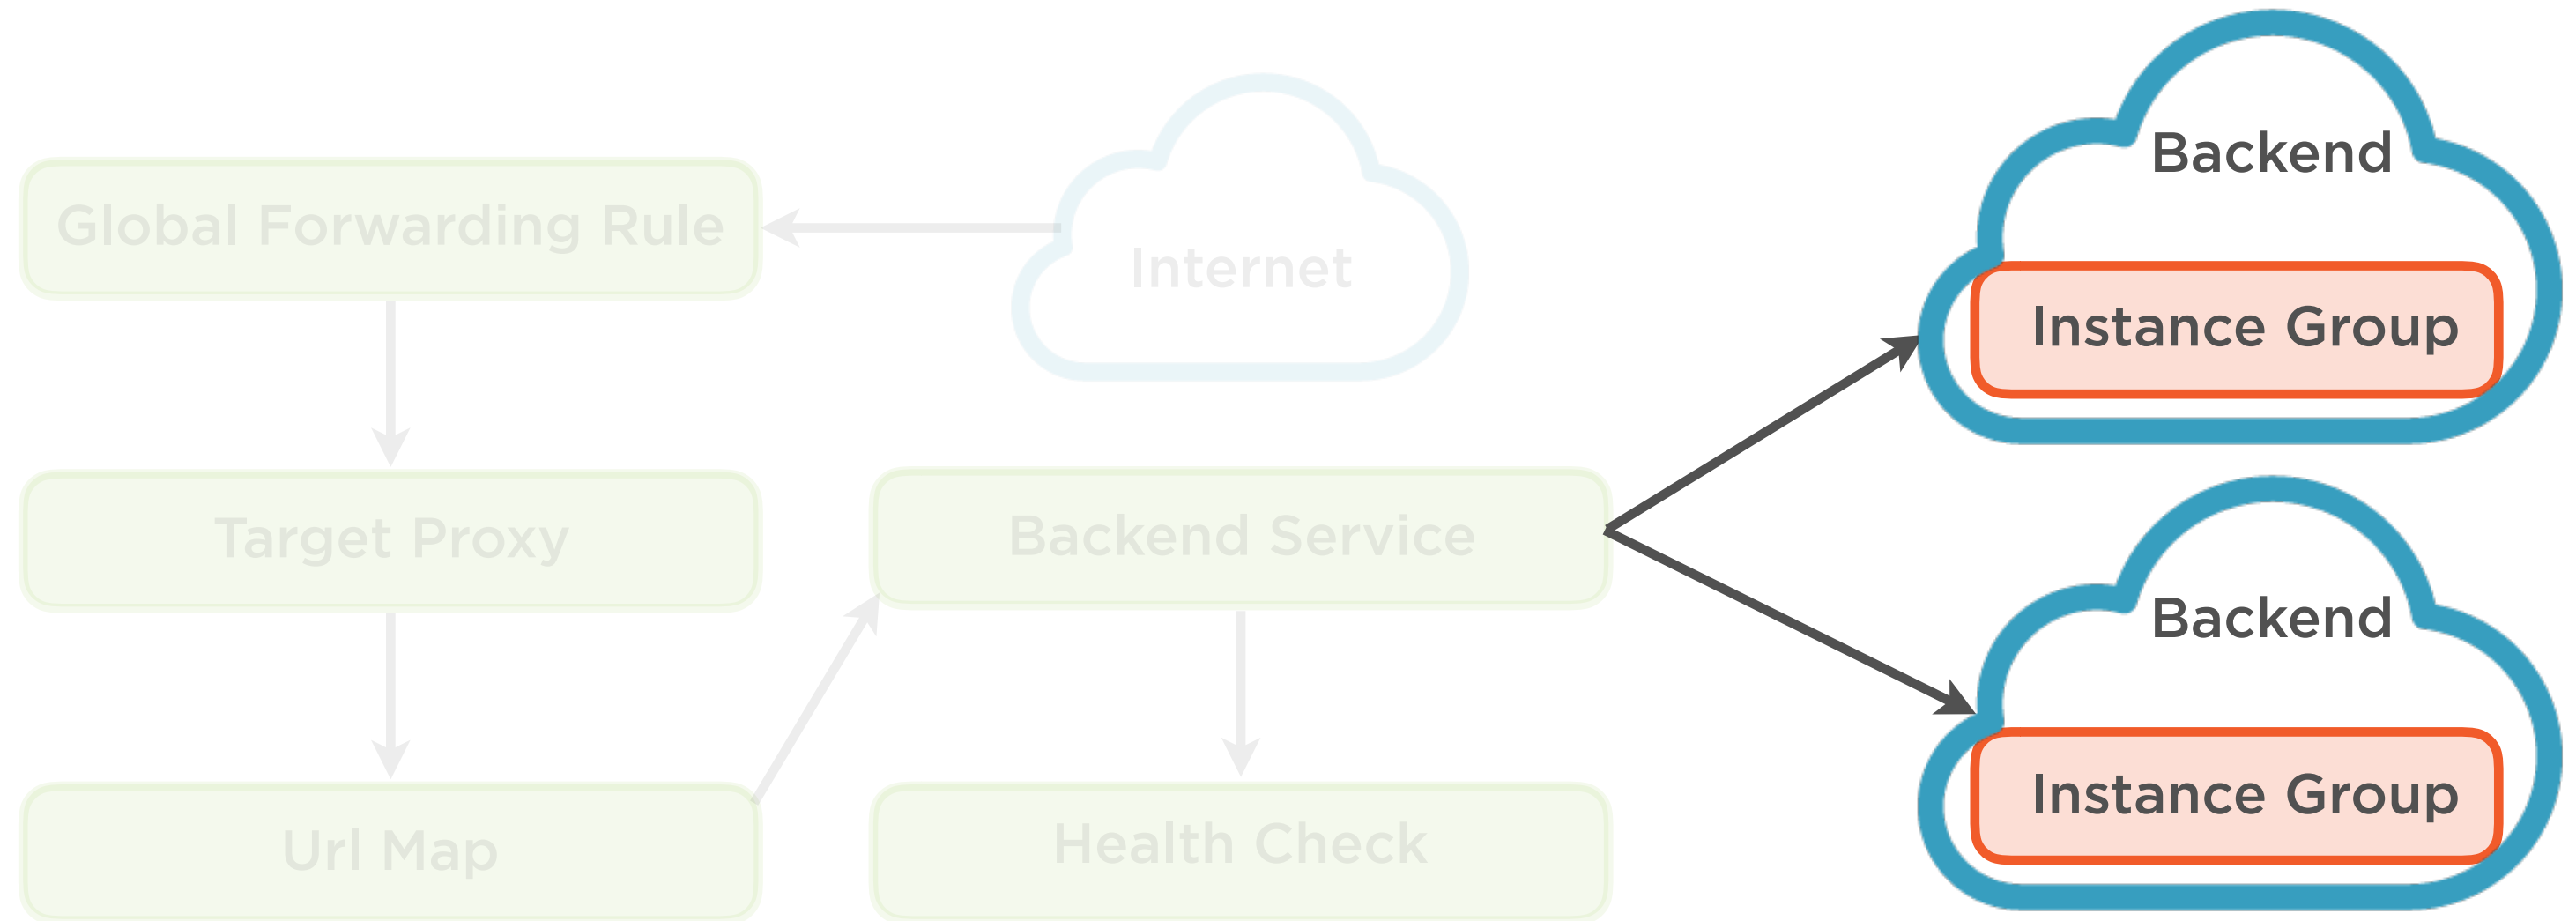
**Health Check:** Polls instances to determine which one can receive requests

**Backends:** Instance group of VMs which may be automatically scaled

**Session Affinity:** Attempts to send requests from the same client to the same VM


**Timeout:** Time the backend service will wait for a backend to respond

# Backends



**Actual request distribution can happen based on CPU utilization, requests per instance**

# Backends



**Instance group:** Can be a managed or unmanaged instance group

**Balancing mode:** Determines when the backend is at full usage

- CPU utilization, Requests per second

**Capacity setting:** A % of the balancing mode which determines the capacity of the backend

# Backend Buckets




**Allow you to use Cloud Storage buckets with HTTP(S) load balancing**

**Traffic is directed to the bucket instead of a backend**

**Useful in load balancing requests to static content**

# Load Distribution




Uses CPU utilization of the backend or requests per second as the **balancing mode**

Maximum values can be specified for both

Short bursts of traffic above the limit can occur



# Load Distribution



Incoming requests are first sent to the **region closest to the user**, if that region has capacity

Traffic distributed amongst zone instances based on capacity

Round robin distribution across instances in a zone

Round robin can be overridden by session affinity

# Demo

**Configure a cross-regional load balancer  
using unmanaged instance groups**

# Demo

**Configure a cross-regional load balancer using managed instance groups**

**Load test the load balancers to observe autoscaling and traffic distribution in action**

# Summary

**Introducing instance templates**

**Managed instance groups and unmanaged instance groups**

**Architectural overview of the HTTP(S) load balancer and its components**

**Cross-regional load balancers with unmanaged and managed instance groups**

**Autoscaling with managed instance groups**