

# Building Pipelines for Workflow Orchestration Using Google Composer

---

INTRODUCING GOOGLE COMPOSER



**Vitthal Srinivasan**

CO-FOUNDER, LOONYCORN

[www.loonycorn.com](http://www.loonycorn.com)

# Overview

**Workflow orchestration service on GCP**

**Define pipeline as DAG**

**Simple Python code**

**Great visualization support**

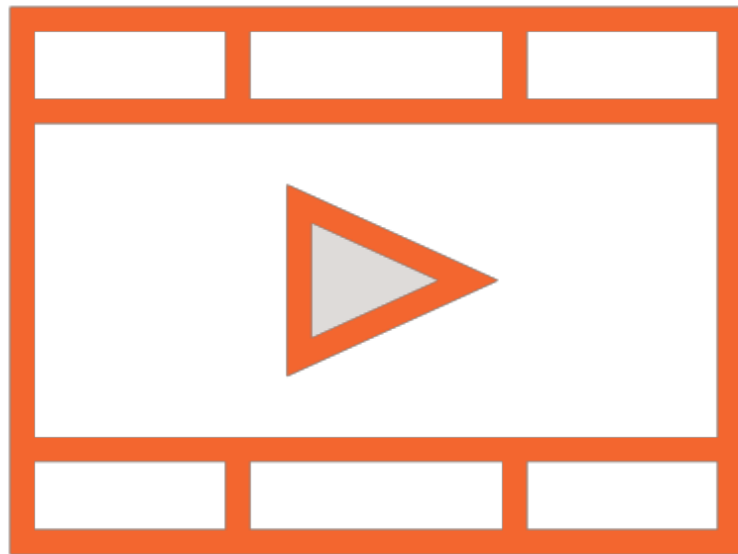
**Runs on Kubernetes**

**Host of useful operators**

# Prerequisites and Course Outline

---

# Prerequisites: Basic Cloud Computing



**Architecting Scalable Web Applications  
Using Google App Engine**

**Architecting Event-driven Serverless  
Solutions Using Google Cloud Functions**

# Course Outline



**Introducing Google Composer**

**Creating, configuring, and accessing environments**

**Managing and monitoring workflows**

# Scenarios: SpikeySales.com



## **Hypothetical online retailer**

- Flash sales of trending products
- Spikes in user traffic

## **SpikeySales on the GCP**

- Cloud computing fits perfectly
- Pay-as-you-go
- No idle capacity during off-sale periods

# Important Composer Concepts

---

Composer is a pipeline  
orchestration technology  
similar to Oozie or Azkaban



# Using Composer




**Write code  
for pipeline**



**Copy into GCS  
bucket**



**Airflow picks up  
and schedules**



**Pipeline  
parallelized and  
executed**

# Important Composer Concepts

**Airflow**

**Composer**

**Environment**

**DAG**

**Operator**

**Trigger**



## Airflow

**Apache Airflow (incubating)**

**Create workflows as DAGs**

**Airflow schedules and executes**

**Ensures dependencies satisfied**

A solid teal rectangle containing the word "Airflow" in white text.

**Airflow**

**Simple Python API**

**Scalable architecture**

**Powerful operators**

**Jinja templates**

A teal rectangular box containing the word "Airflow" in white text.

## Airflow

**Airflow workers**

**Web server to managed pipelines**

**Scheduler tracks and executes DAGs**

**Celery task queue to scale workers**

**Redis as message broker for Celery**

## Composer

**GCP managed service for Airflow**

**Airflow workers on GKE cluster**

**Airflow metadata on Cloud SQL**

**Airflow web server on App Engine Flex**

**GCS bucket for DAGs**

## **Environment**

**Self-contained Airflow deployment**

**Google-managed tenant project**

**Hosted entirely within a region**



**DAG**

**Directed-Acyclic-Graph**

**Defined in Airflow Python script**

- Must execute instantaneously
- Merely DAG definition file

**Contains operators and dependencies**

**Different tasks run on different workers**





**DAG**

**DAG defined in global namespace**

**Uses context manager**

**Contains dependencies**

**Copied to environment GCS bucket**

**Scheduled and executed periodically**



## Operator

**Corresponds to one step in pipeline**

**Each instantiation of an operator is one task instance**

**Each task instance exists inside pipeline**

**Used with Python context manager**

An orange rectangle with the word "Operator" in white text.

**Operator**

**BashOperator**

**PythonOperator**

**BranchPythonOperator**


**SendGrid**

**BigQuery**

**Dataprocc**

**Cloud Storage buckets**

# BashOperator




**Cloud Composer runs the provided commands in a Bash script on a worker**

**Worker is a Debian-based Docker container and includes several packages**

- gcloud
- bq
- gsutil
- kubectl


# PythonOperator



**Cloud Composer runs the Python code in a container that includes several packages**

- google-cloud-bigquery
- google-cloud-dataflow
- google-cloud-storage
- pandas
- Tensorflow

# GCP Operators



**Cloud Composer automatically configures an Airflow connection to the environment's project**

- BigQuery operators query and process data in BigQuery
- Cloud Dataflow operators run Apache Beam jobs in Cloud Dataflow

# EmailOperator

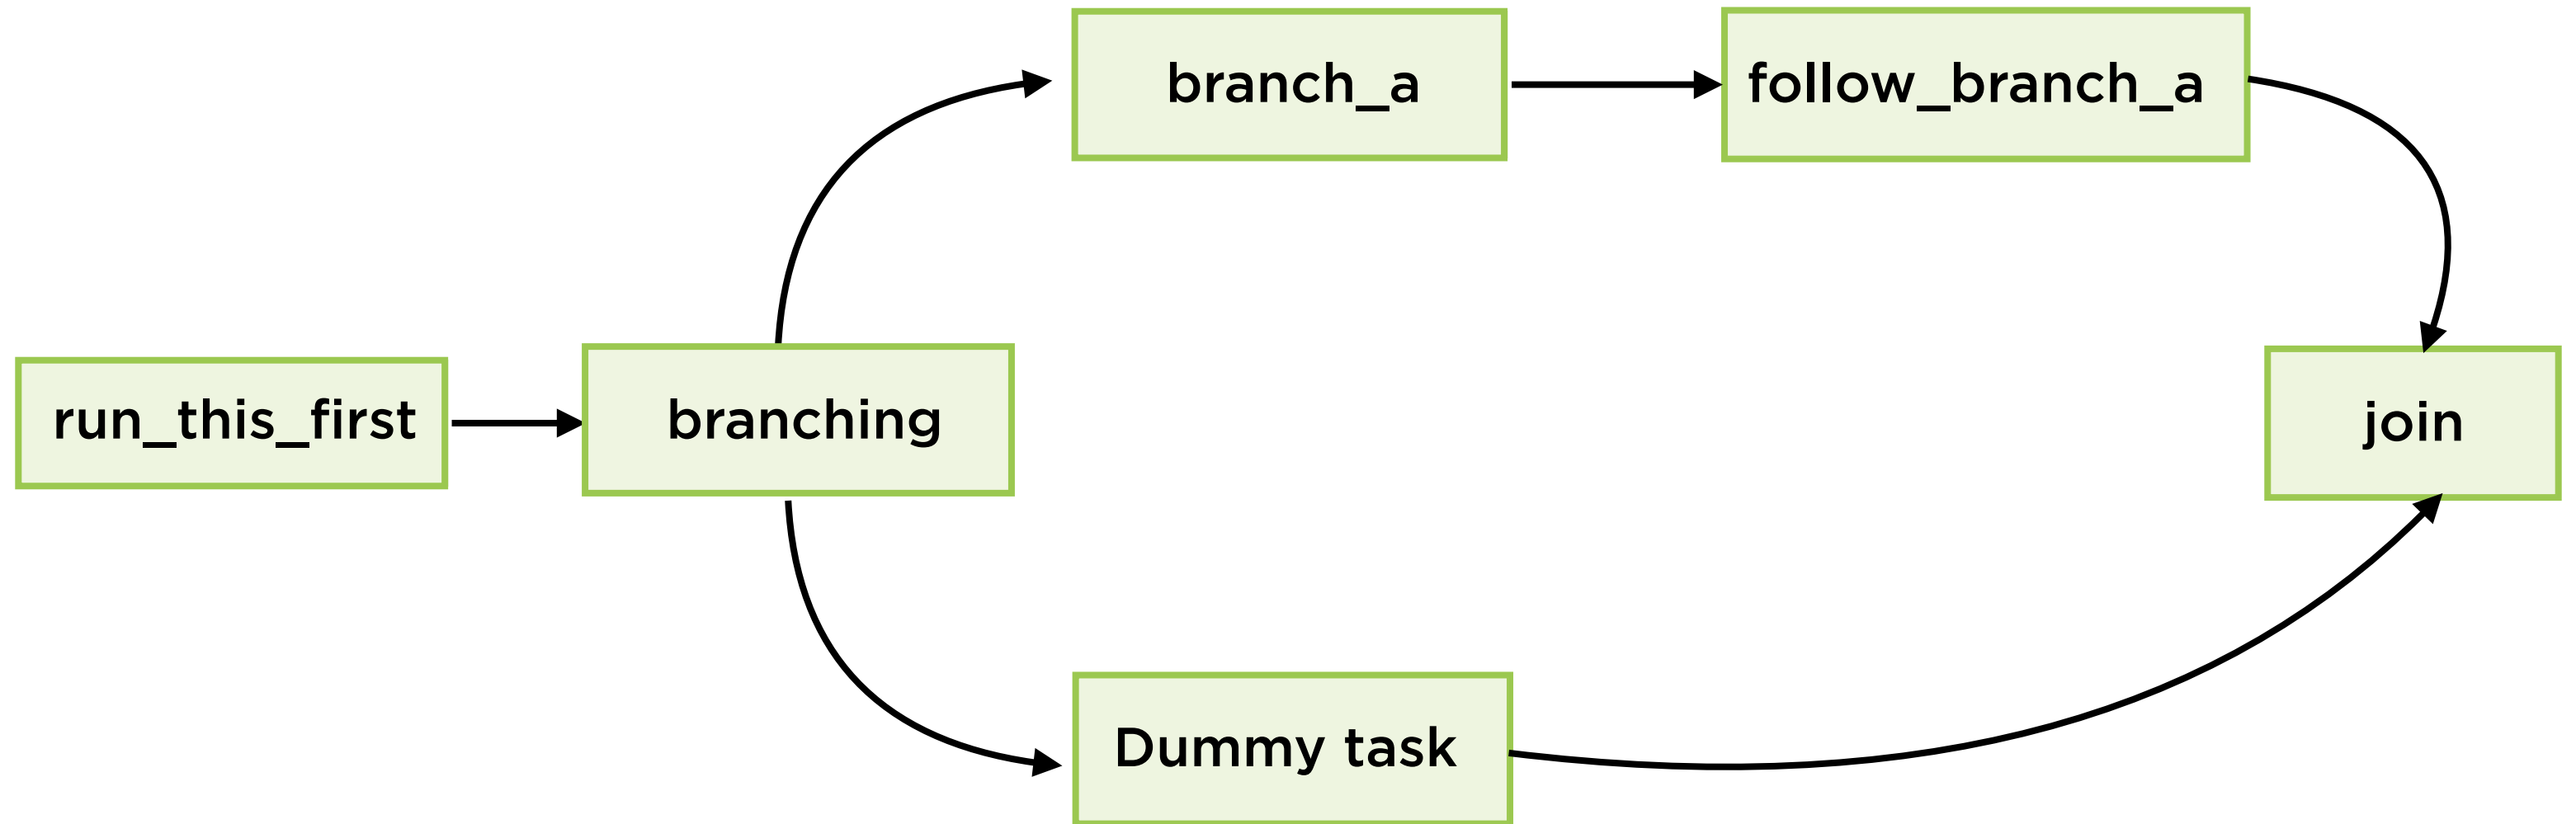
Use the EmailOperator to send email from a DAG. To send email from a Cloud Composer environment, you must configure your environment to use SendGrid.

# Branching

BranchPythonOperator expects a `python_callable` that returns a `task_id`; The `task_id` returned is followed, and all of the other paths are skipped.

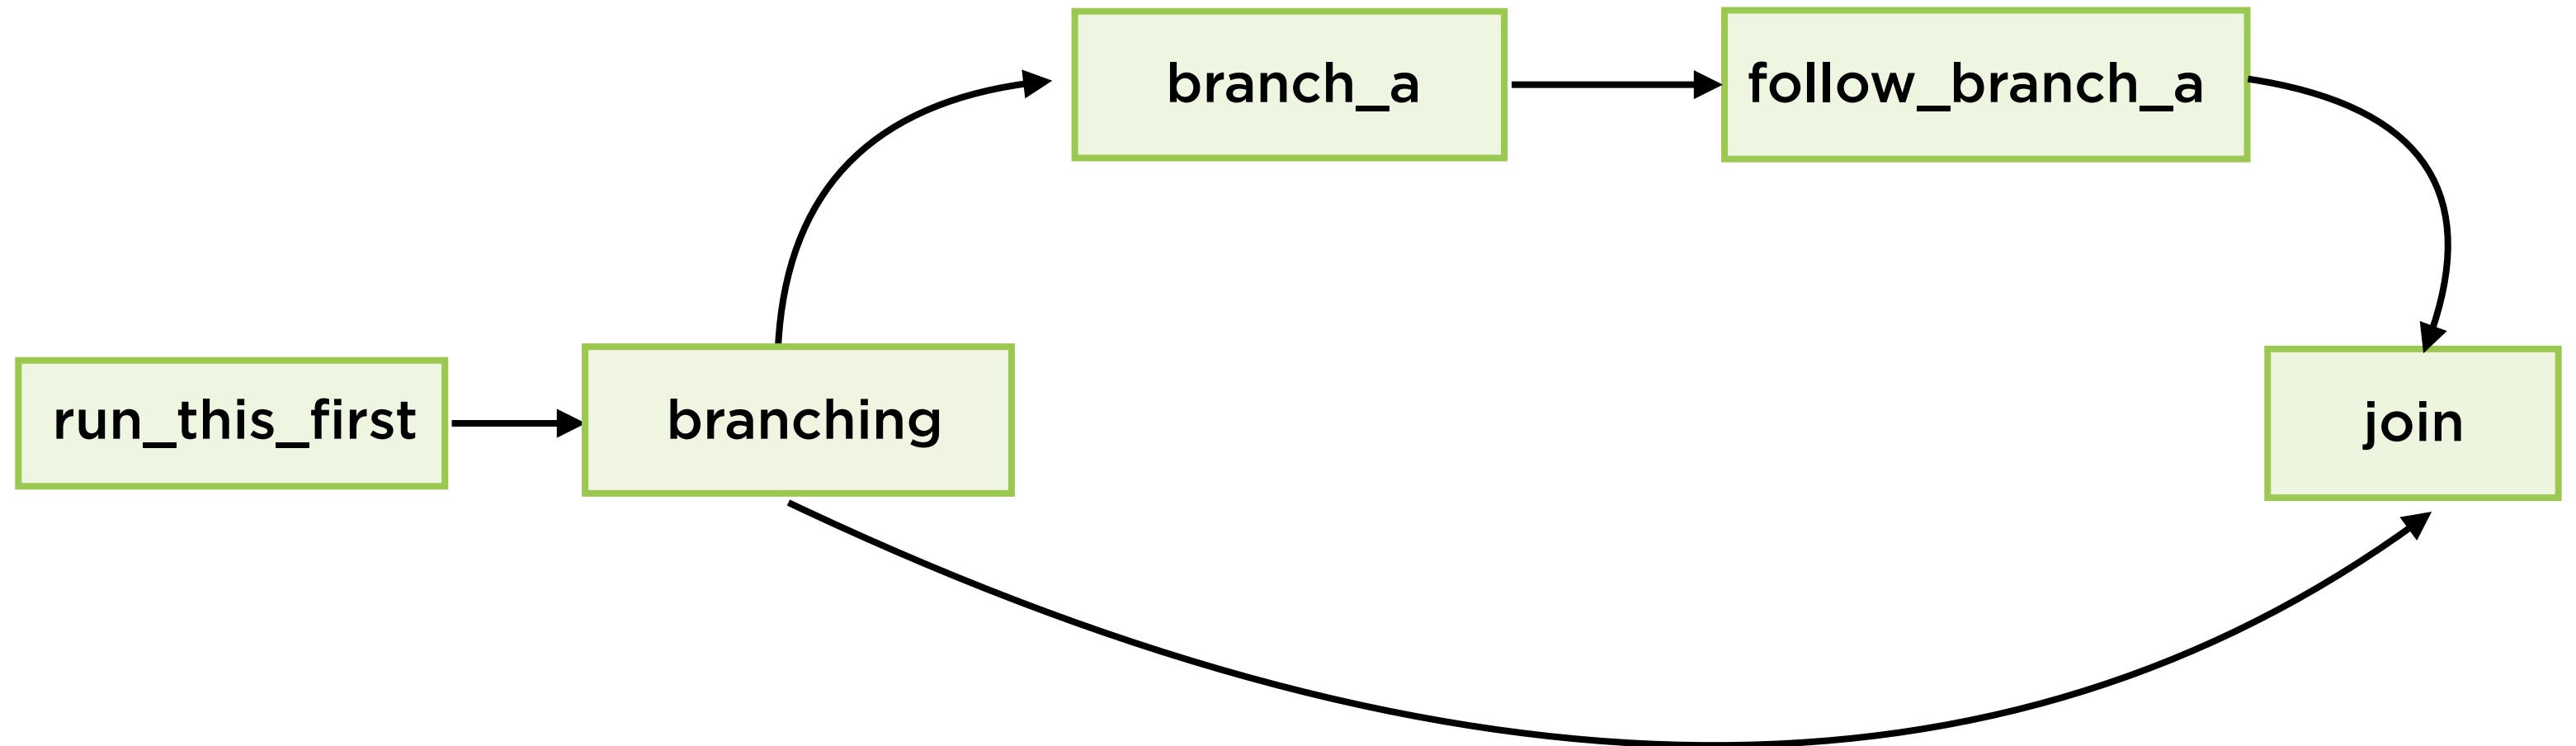


# Branching Needs Dummy Operators



Never leave an empty path after a branch - use a dummy task instead

# Branching Done Wrong



**Here the join task is skipped**

## Trigger

**Airflow scheduler monitors DAGs**

**Triggers task instances whose dependencies have been met**

**Runs each DAG one `schedule_interval` after start date**

A solid green rectangle with the word "Trigger" centered inside it in white text.

Trigger

## Several trigger rules

- all\_success
- all\_failed
- one\_failed
- one\_success
- dummy

Composer is a pipeline  
orchestration technology  
similar to Oozie or Azkaban

# Dataflow vs. Composer

## Dataflow

Data processing pipelines

Focus on windowing and streaming

Complex to implement

Cumbersome to trigger

Visualization UI exists but not central

Few specialized operators

## Composer

General purpose pipelines

Focus on scripting and Python

Simple to implement

Trivial to trigger

Fundamentally UI-based

Many helpful specialized operators

# Dataflow vs. Composer

## Dataflow

**Serverless, no clusters provisioned**

**No access to compute nodes**

**Apache Beam API**

## Composer

**Runs on Kubernetes cluster**

**Compute nodes in GKE cluster**

**Apache Airflow API**



# Using Dataflow




**Write code  
for pipeline**



**Submit job for  
execution**



**Dataflow assigns  
workers to  
execute**




**Pipeline  
parallelized and  
executed**

# Using Composer




**Write code  
for pipeline**



**Copy into GCS  
bucket**



**Airflow picks up  
and schedules**



**Pipeline  
parallelized and  
executed**

# Demo

**Enable APIs and create Composer environment**

# Demo

**Write and run DAG on Composer environment**

# Pricing

---

# Google Composer Pricing



**Based on the size of Cloud Composer environment**

**Per-minute billing**

**In addition to charges for**

- Google Kubernetes Engine
- Cloud Storage

# Indicative Pricing

<https://cloud.google.com/composer/pricing>

Item	Price (USD)
Web core hours	\$0.074/vCPU hour
Database core hours	\$0.125/vCPU hour
Web and database storage	\$0.273 per GB/month
Network egress	\$0.156/GB

# Environment Sizing

Settings	Default	Adjustable
Storage (GB)	20	No
Database vCPUs	2	No
Web server vCPUs	2	No
Worker machine type	n1-standard-1	Yes
Worker nodes	3	Yes
Worker storage (GB per worker)	100	Yes



# Google Composer Pricing Example



## Create a Cloud Composer environment

### Assume

- Default specs
- Environment used for 25% of time
- Equivalent to 182 hours/month
- 6.5 GB of egress

# Google Composer Pricing Example

Resource	Total cost
Database core hours	\$45.66
Web core hours	\$26.91
Web and database storage	\$1.37
Network egress	\$1.44
Total Cloud Composer cost	\$76.06

# Summary

**Workflow orchestration service on GCP**

**Simple Python code**

**Easy to visualize**

**Runs on Kubernetes**

**Host of useful operators**