

Architecting Scalable Web Applications Using Google App Engine

INTRODUCING GOOGLE APP ENGINE



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Hosted applications with built-in load balancing and autoscaling

Services and versions

Instances, instance classes and scaling

Choosing between the standard and flexible environment

Deploying and scaling a simple App Engine application

Prerequisites and Course Outline

Prerequisites




Basic understanding of cloud computing

Understanding how web applications are built, GET and POST HTTP requests

Working with NoSQL databases

Prerequisites: Basic Cloud Computing



**Choosing and Implementing Google
Cloud Compute Engine Solutions**

**Architecting Schemaless Scalable NoSQL
Databases Using Google Datastore**

Course Outline



Introducing Google App Engine

- App Engine vs. other compute options
- Standard and Flexible environments
- Services, versions, instances and scaling

App Engine Standard Environment

- Traffic migration and splitting
- Integrating apps with Cloud Datastore
- Integrating with Memcache for lower latency

App Engine Flexible Environment

- Using custom containers to deploy apps
- Integrating applications with other GCP services such as Pub/Sub

Scenarios: SpikeySales.com



Hypothetical online retailer

- Flash sales of trending products
- Spikes in user traffic

SpikeySales on the GCP

- Cloud computing fits perfectly
- Pay-as-you-go
- No idle capacity during off-sale periods


Introducing Google App Engine

Choices in (Any) Computing



Compute

Where and how does code
run?

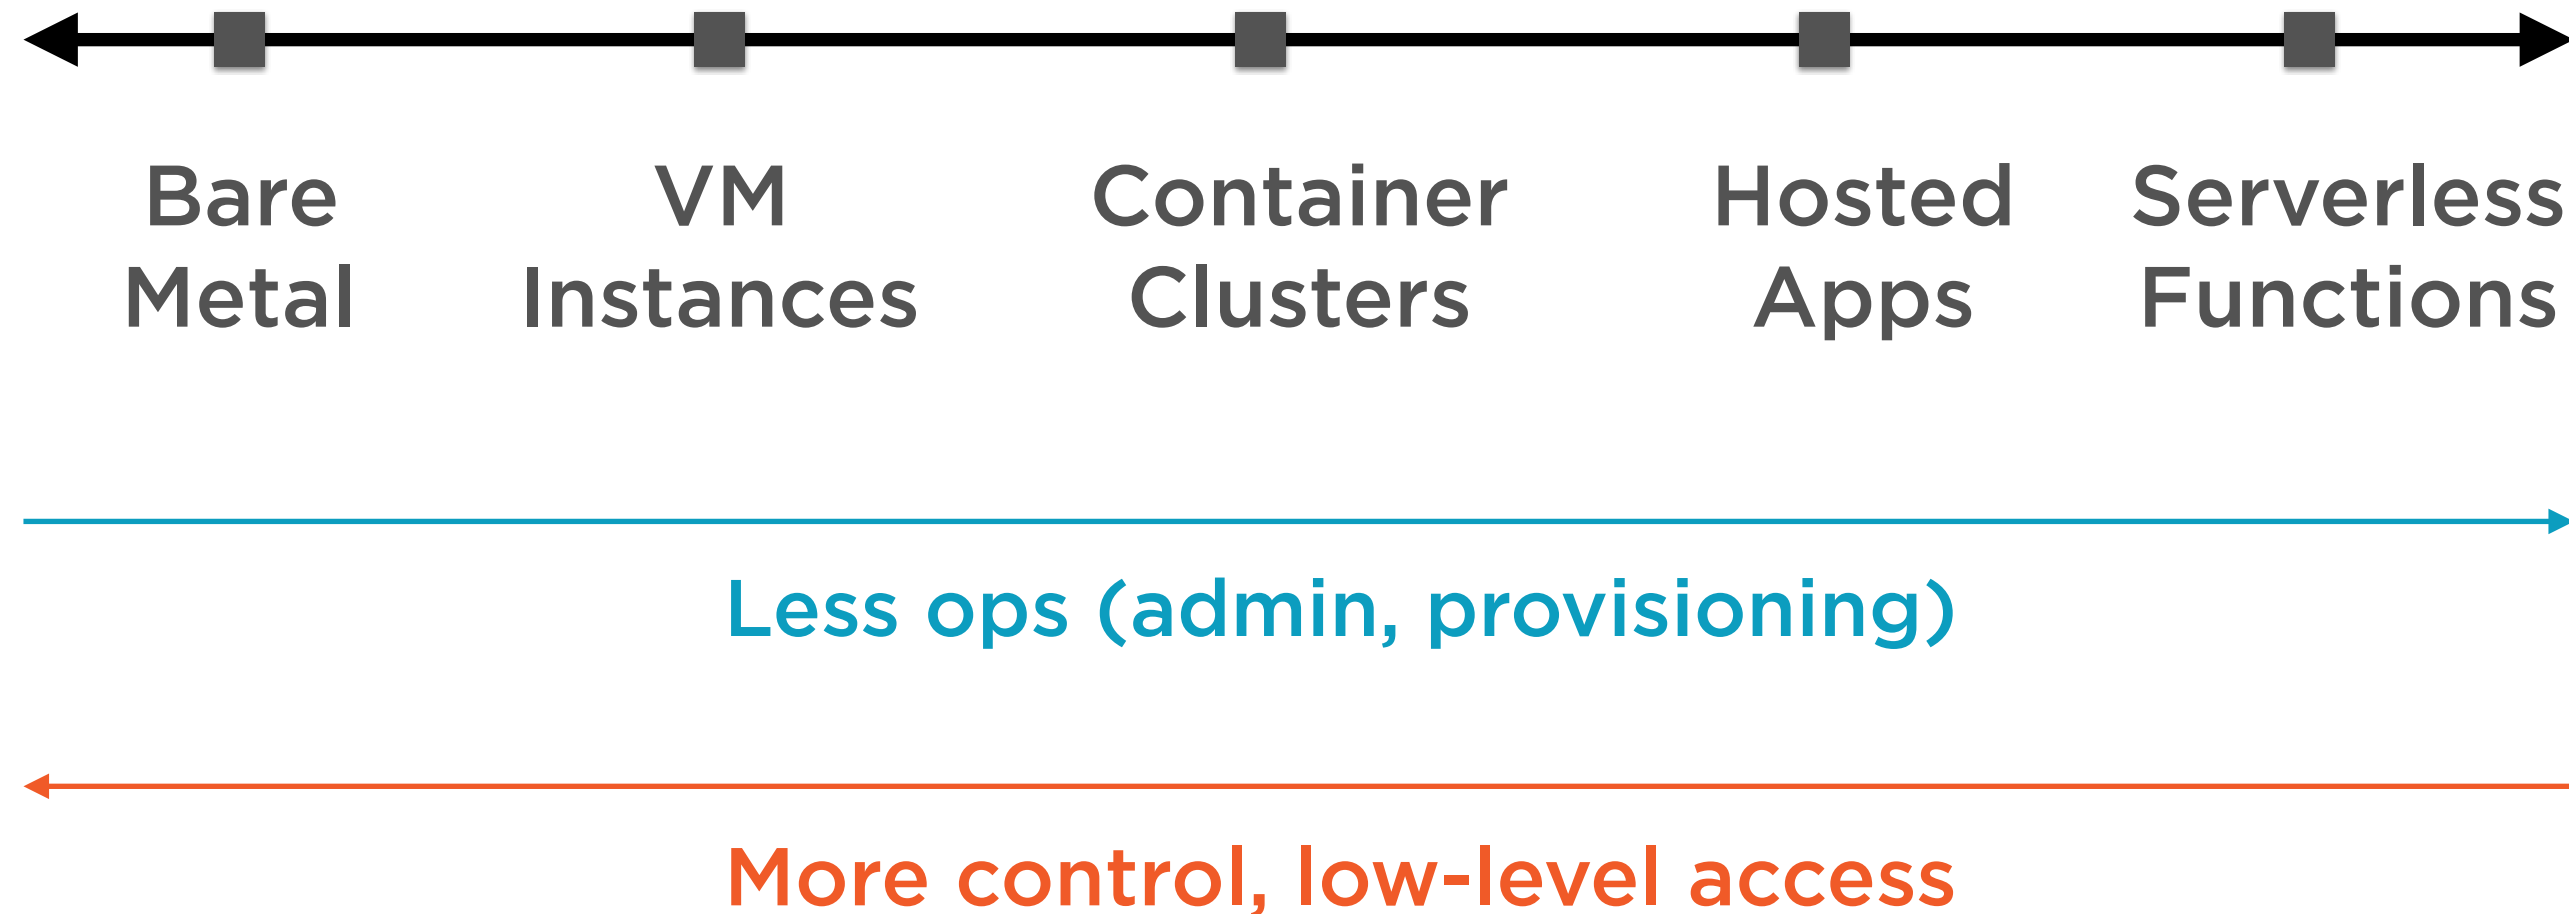


Storage


Where and how is the data
stored?

Other choices - networking, logging etc. - are
less important

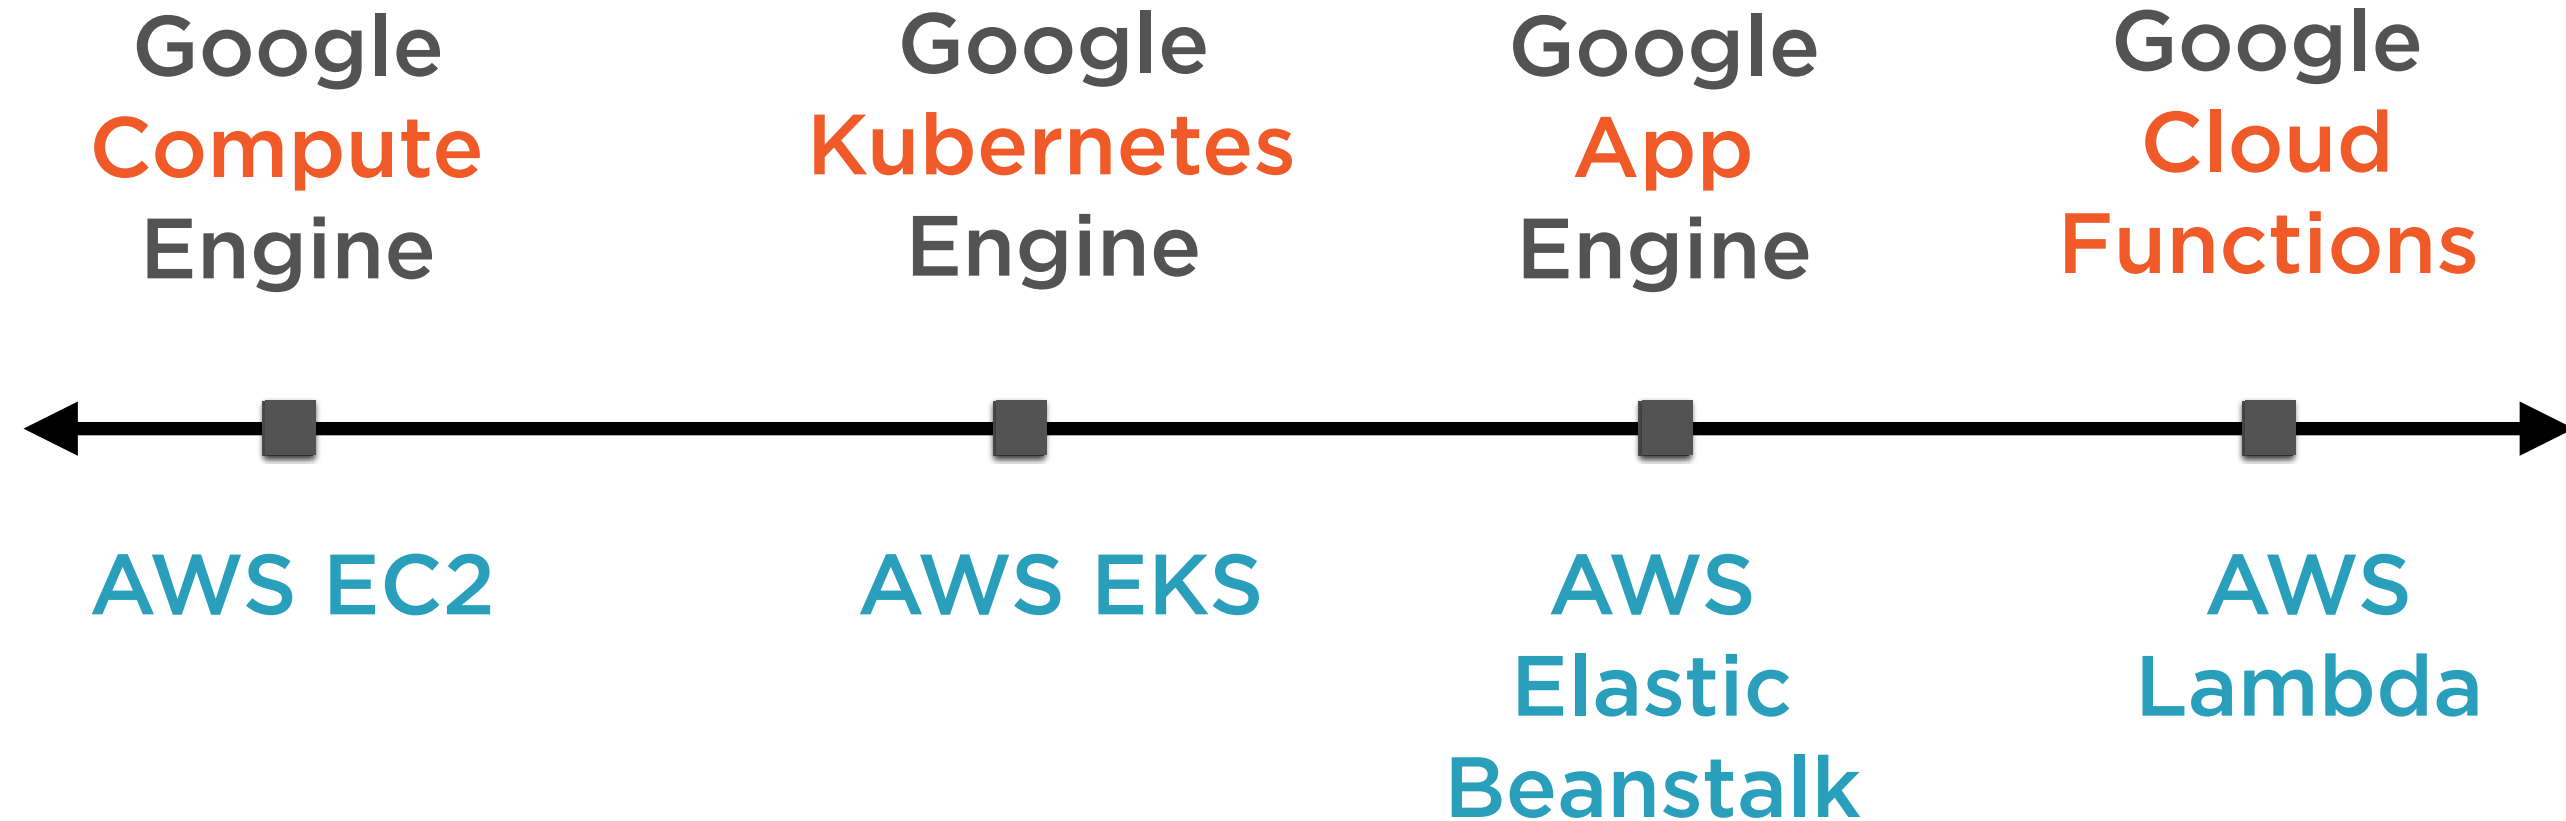
Compute Choices



GCP Compute Choices

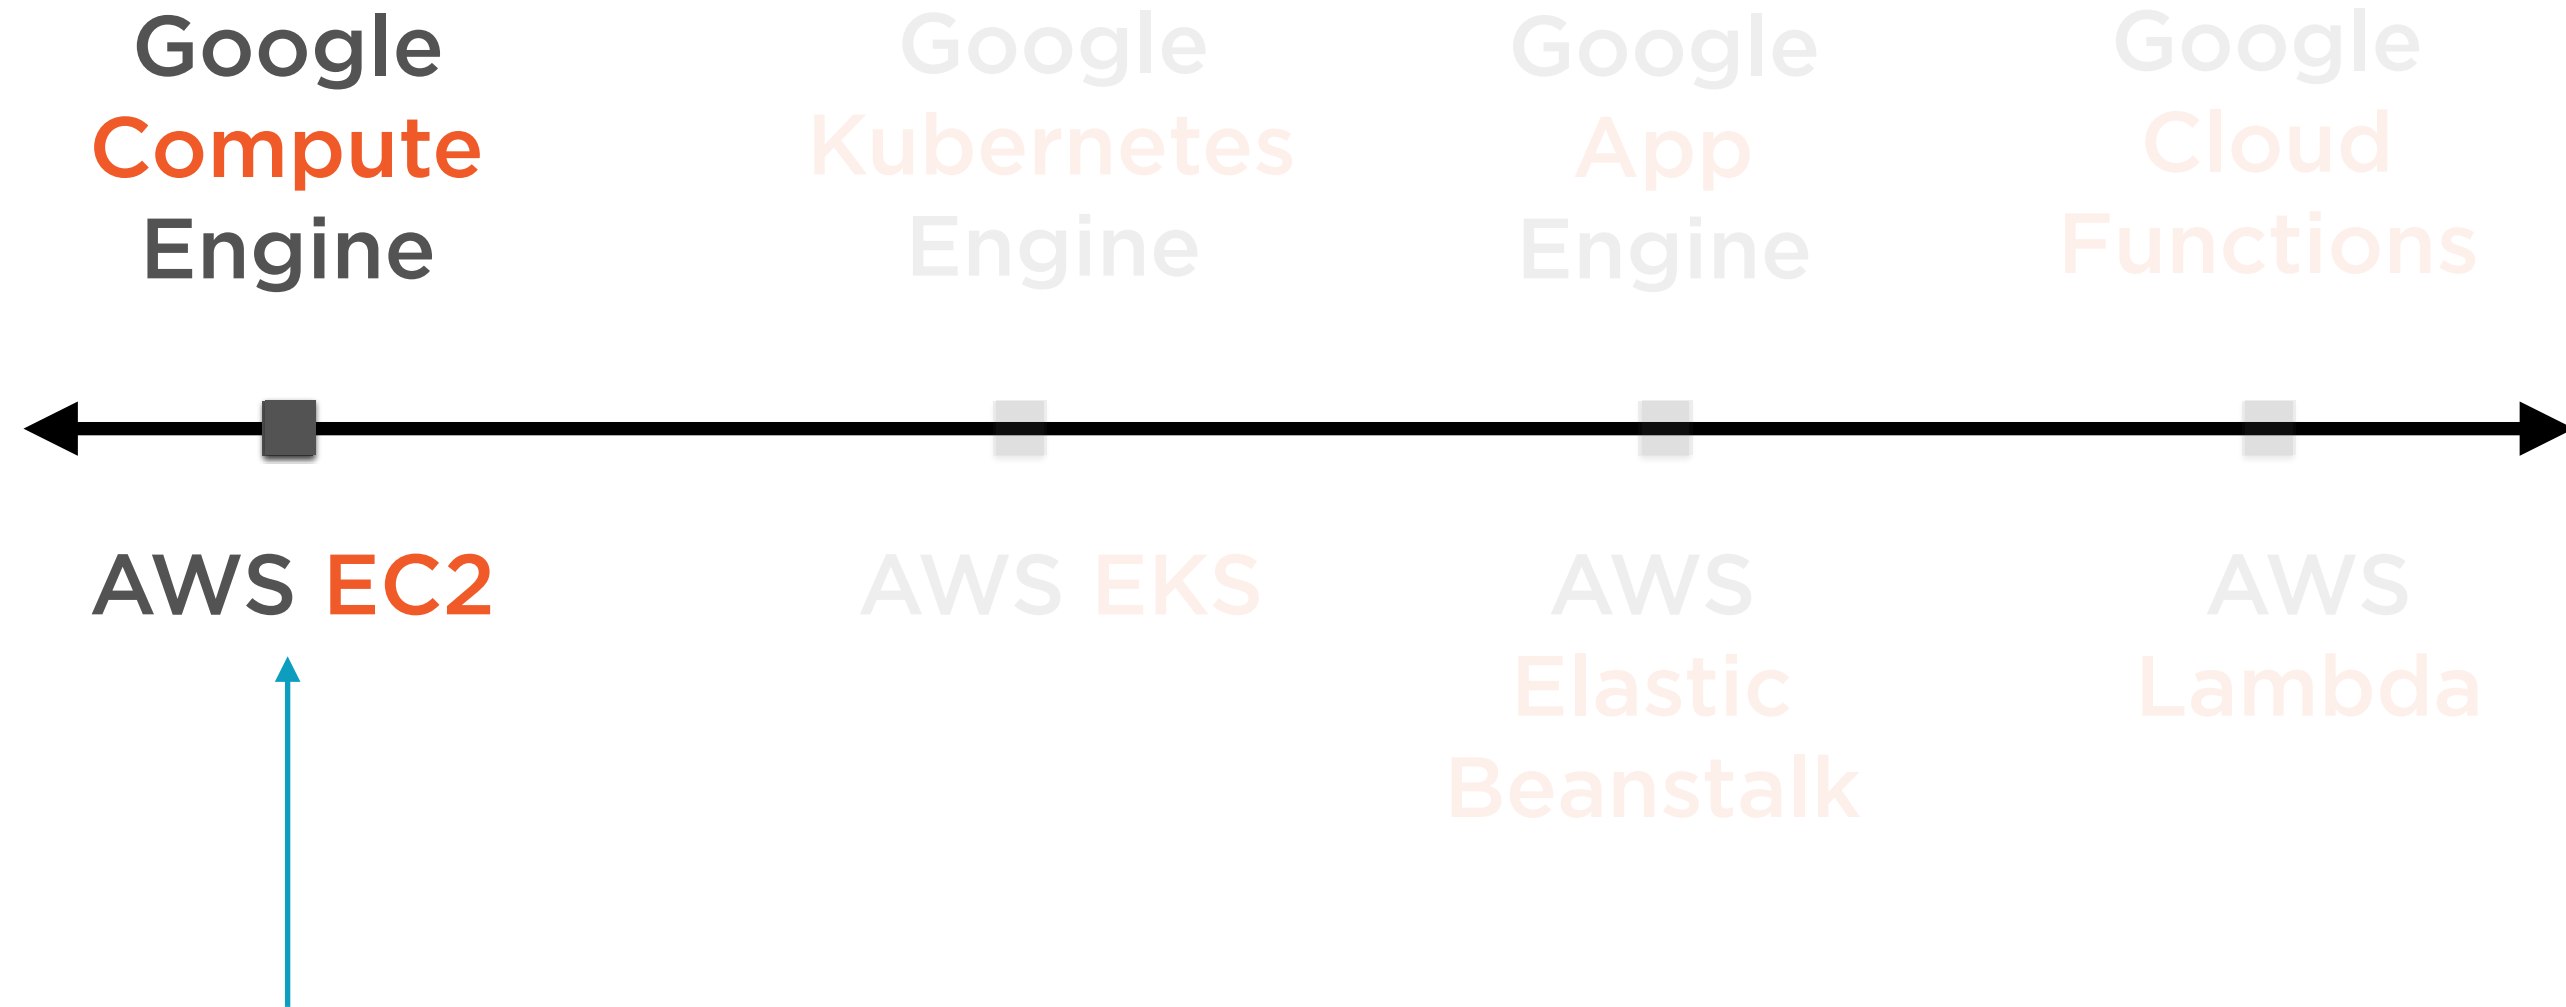


GCP Compute Choices



Every major cloud platform supports the same
range of compute choices

GCP Compute Choices



“Infrastructure-as-a-Service” (IaaS)

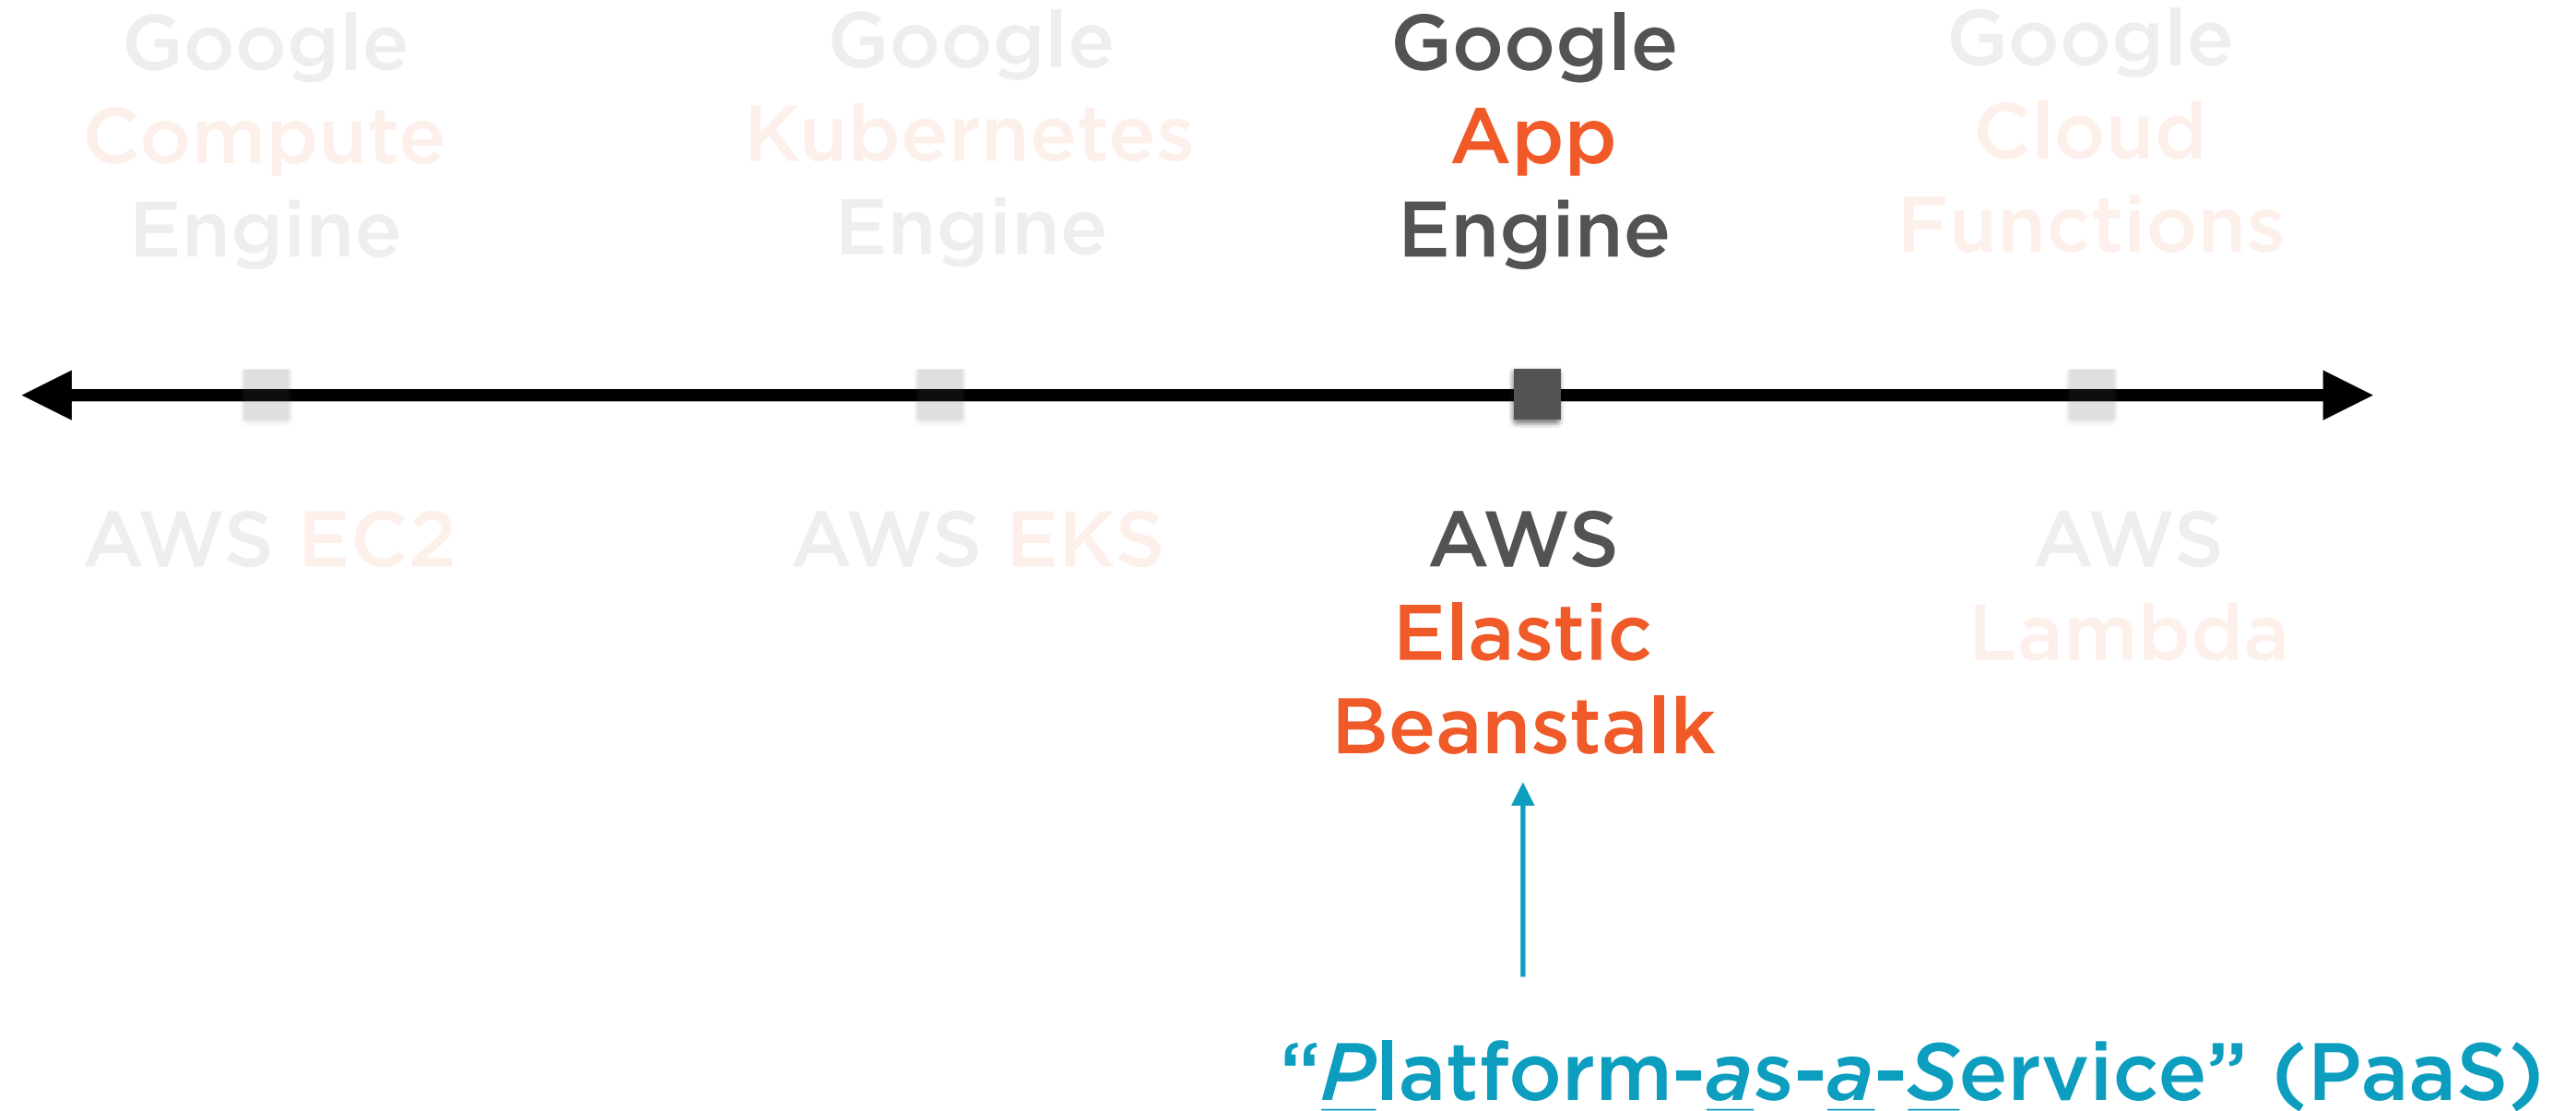
GCP Compute Choices



GCP Compute Choices



GCP Compute Choices



Google App Engine

Web framework and platform for hosting web applications
on the Google Cloud Platform

App Engine Environments

Standard Environment

Flexible Environment

App Engine Environments

Standard

App runs in a proprietary sandbox

Code in few languages/versions only

No other runtimes possible

Apps cannot access Compute Engine resources

Flexible

Runs in Docker container on GCE VM

Code in far more languages/versions

Custom runtimes possible (Docker)

Apps can access Compute Engine resources and some OS packages

App Engine Environments

Standard

Instance startup in seconds

Scaling: Manual, Basic or Automatic

No background processes

No SSH debugging

Scale to zero

No installation of third-party binaries

Flexible

Instance startup in minutes

Scaling: Manual or Automatic

Background processes supported

SSH debugging supported

No scaling to zero (minimum 1 instance)

Installation of third-party binaries allowed

App Engine Environments

Standard

Apps that experience traffic spikes

Usually stateless HTTP web apps

All instances in same zone (moved in case of zone outage)


Flexible

Apps that experience consistent traffic

General purpose apps

Instances are part of regional Managed Instance Group

App Engine Standard Runtimes



Python 2.7, Python 3.7


Java 8, Java 7

Node.js 8 (beta), 10 (beta)

PHP 5.5, 7.2 (beta)

Go 1.9, 1.11 (beta)

App Engine Flexible Runtimes



Python 2.7, Python 3.6

Java 8

Node.js

Go 1.9, 1.10, 1.11

Ruby

PHP 5.6, 7.0, 7.1, 7.2

.NET

Custom runtimes

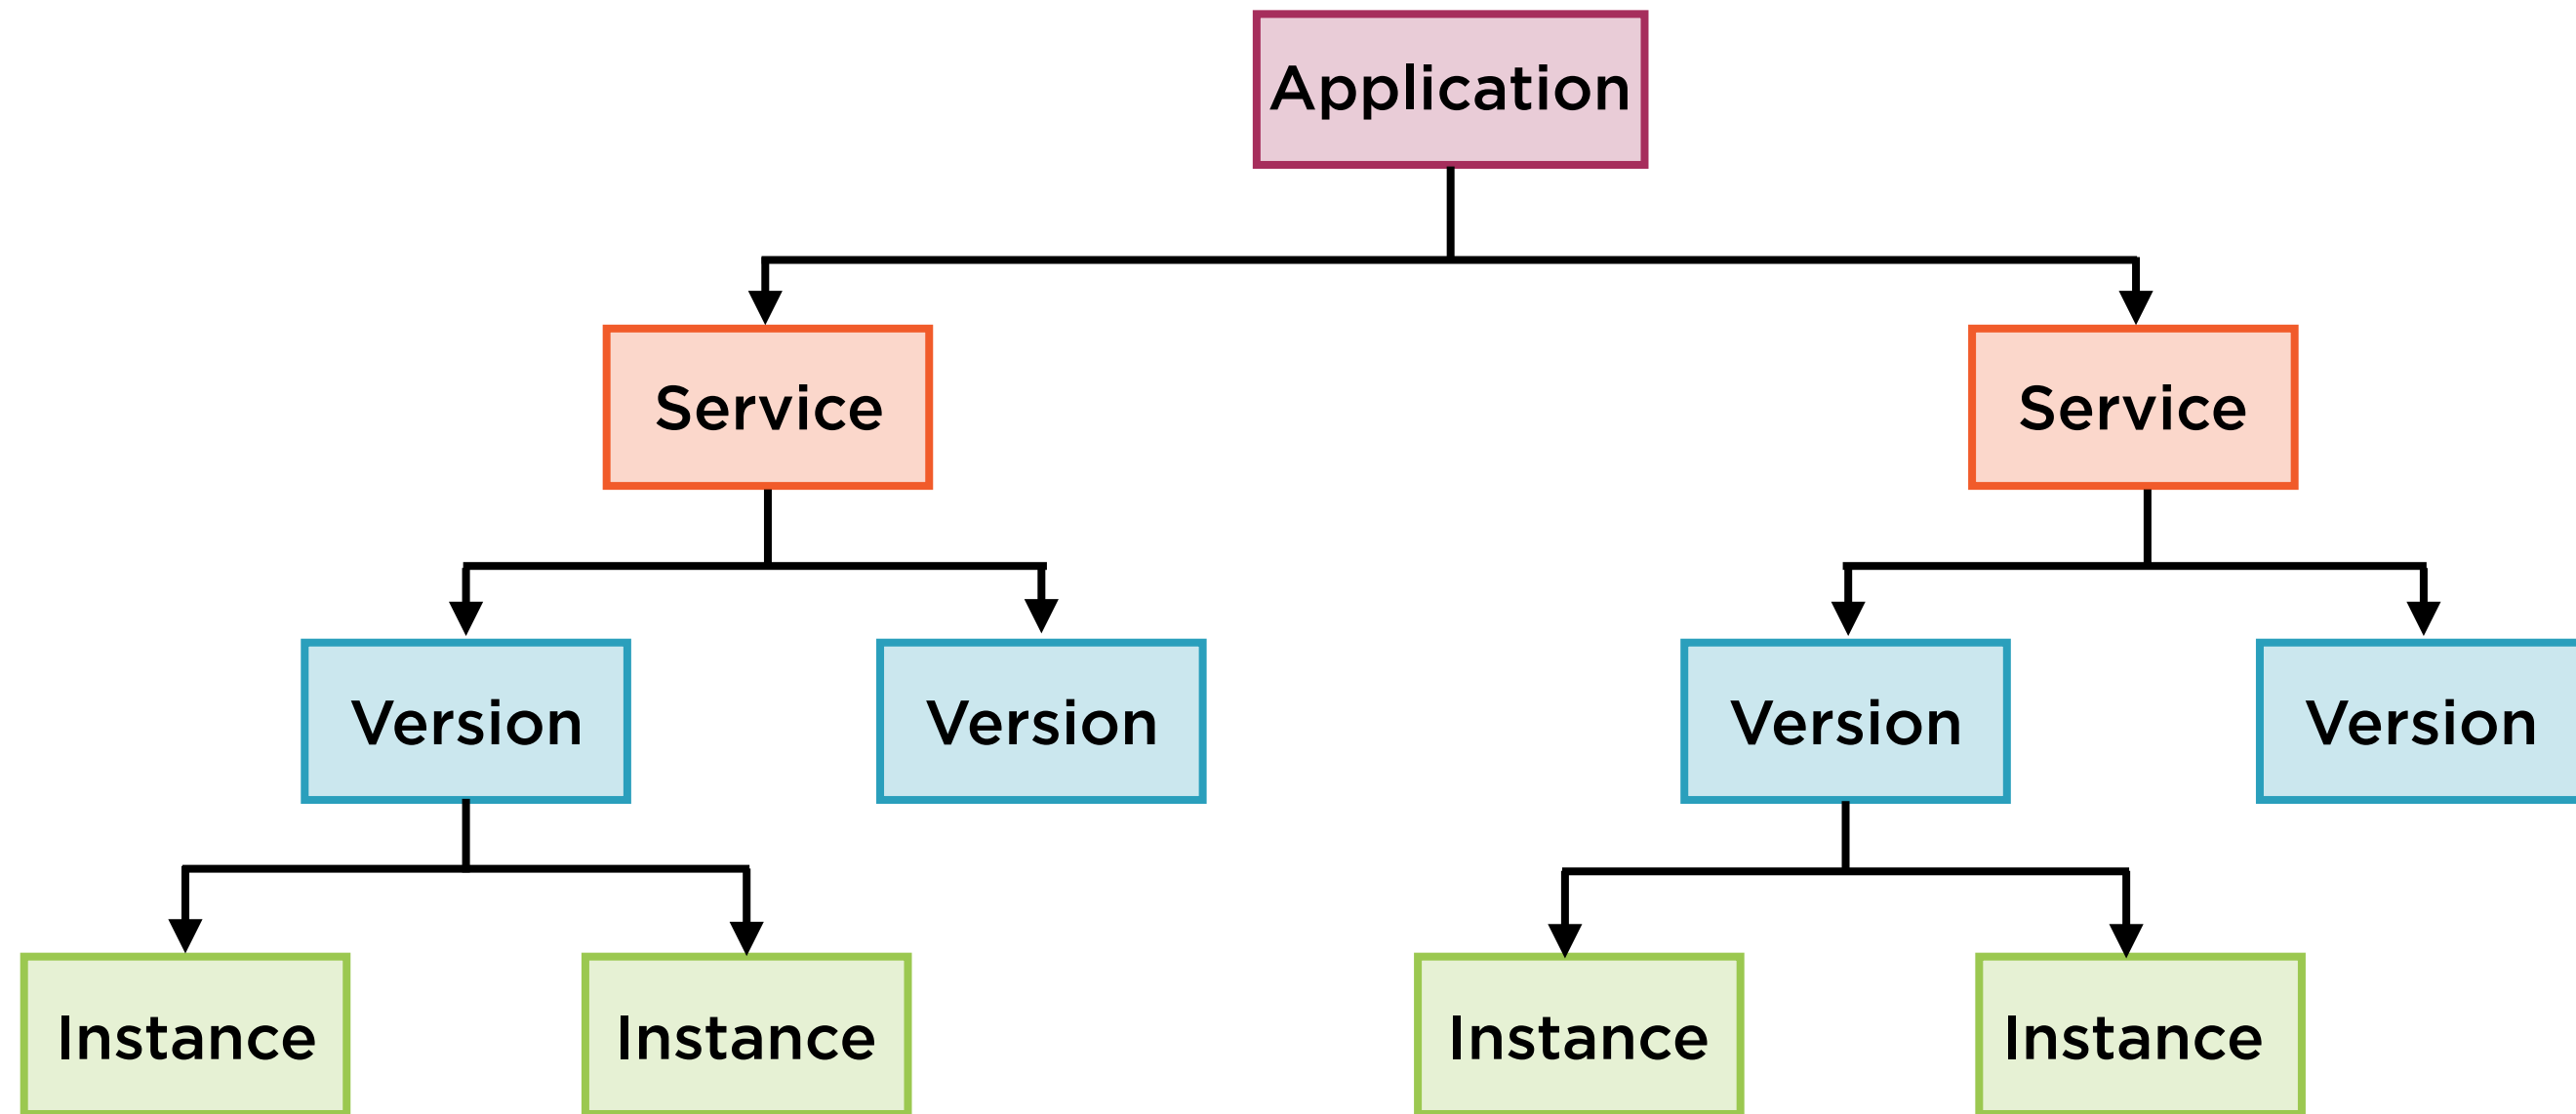
Choose the flexible environment for applications that require custom runtimes or third-party libraries not supported in the standard environment

App Engine Components

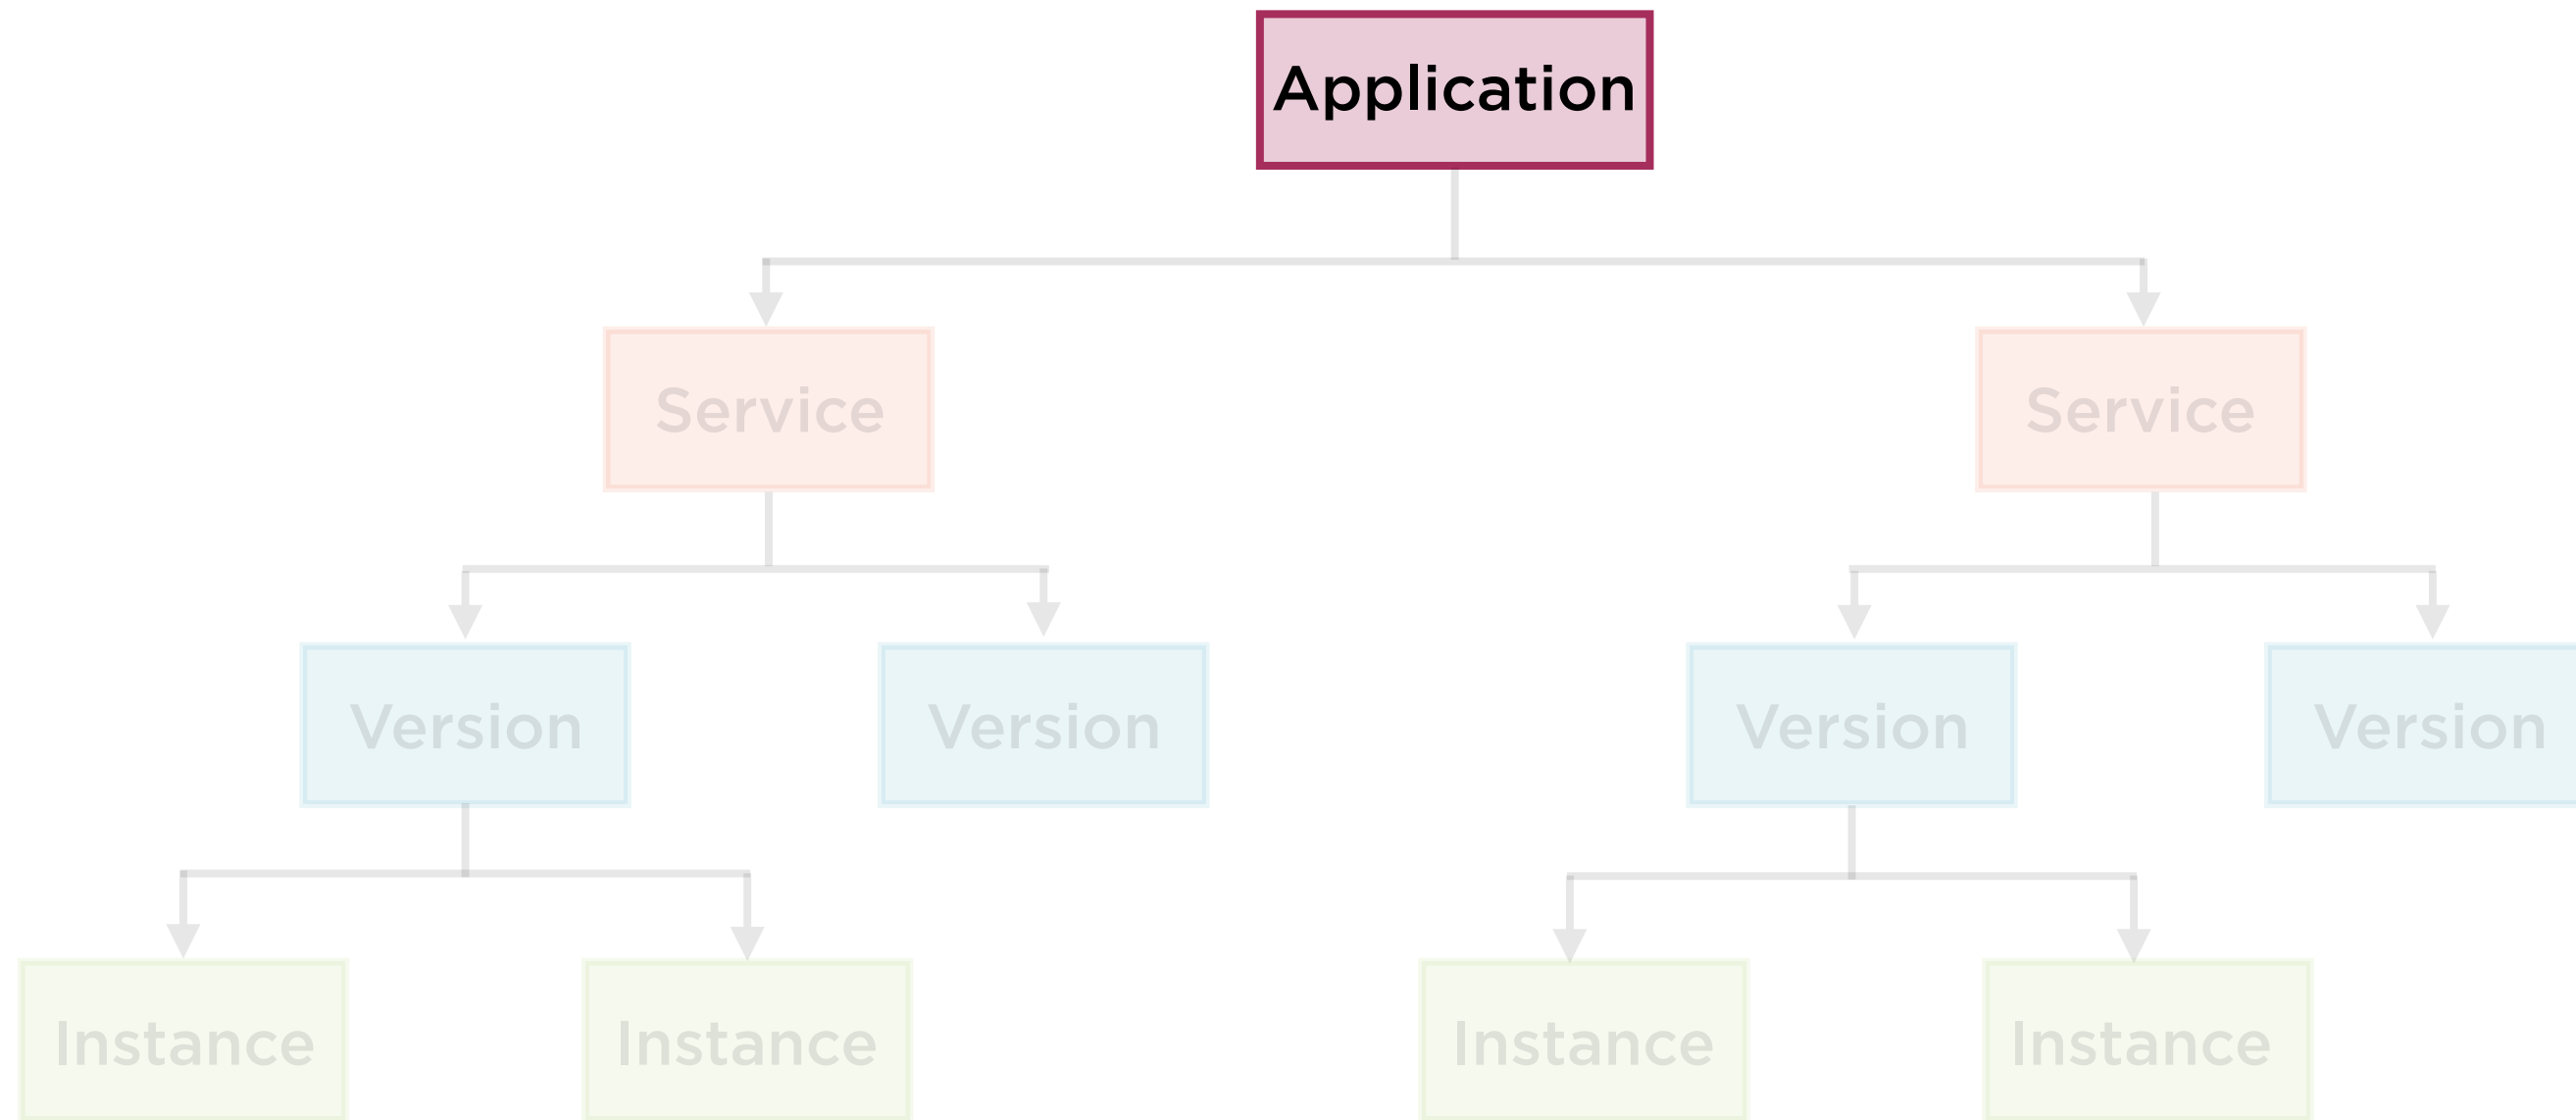
App Engine App

Single regional application resource consisting of hierarchy of services, versions and instances

Components of an Application



Application

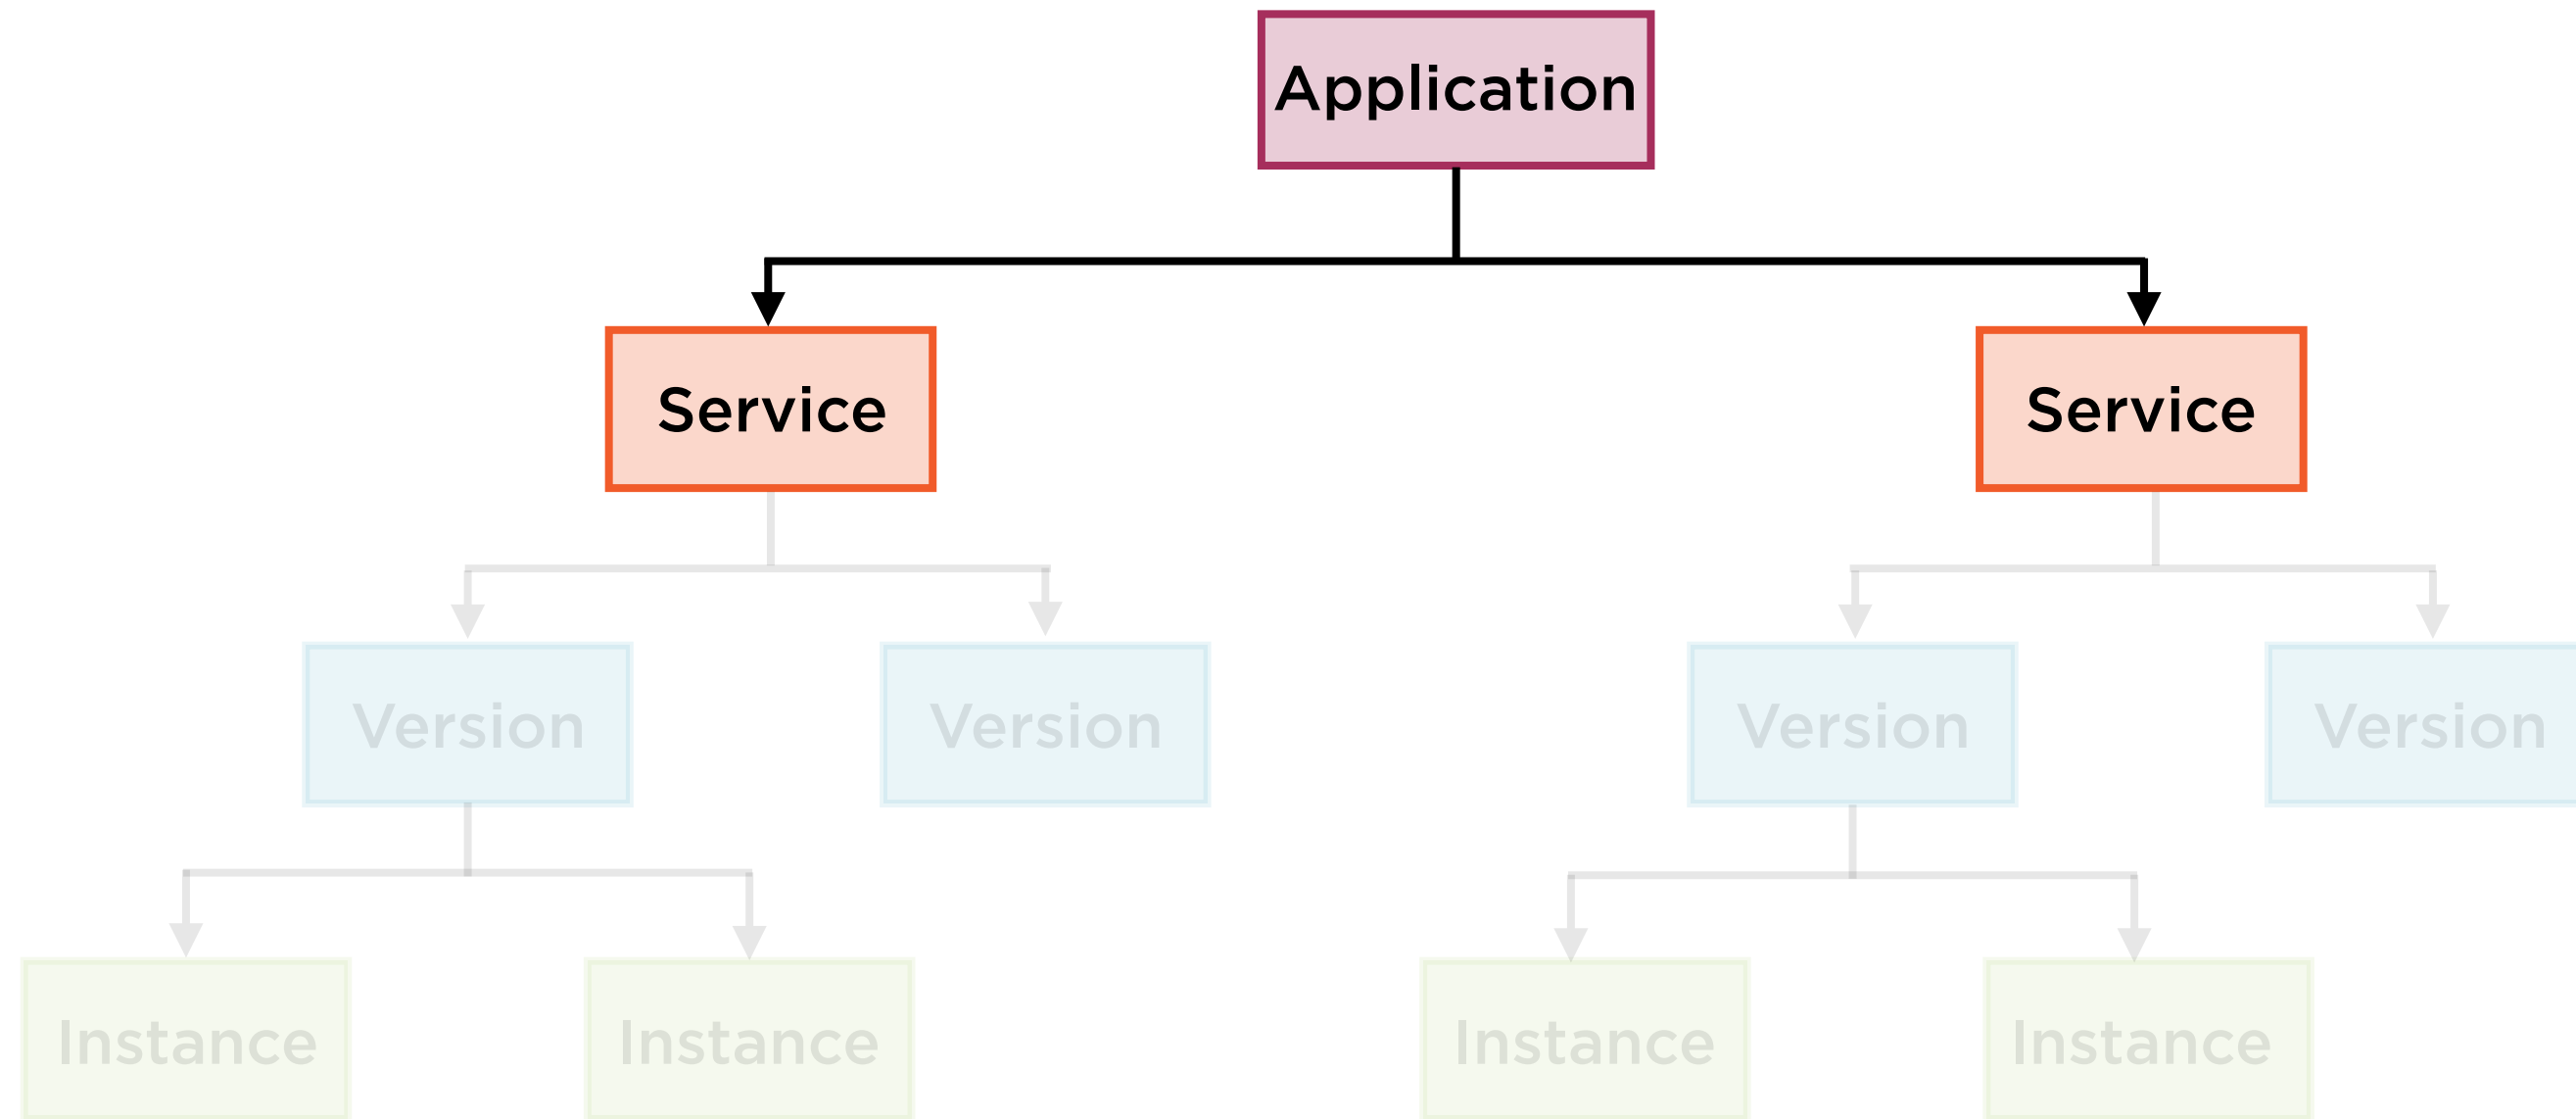


Top-level container for multiple services, their versions and instances

Services


Each service in an App Engine app behaves like a microservice and scales independently

Single Application, Multiple Services



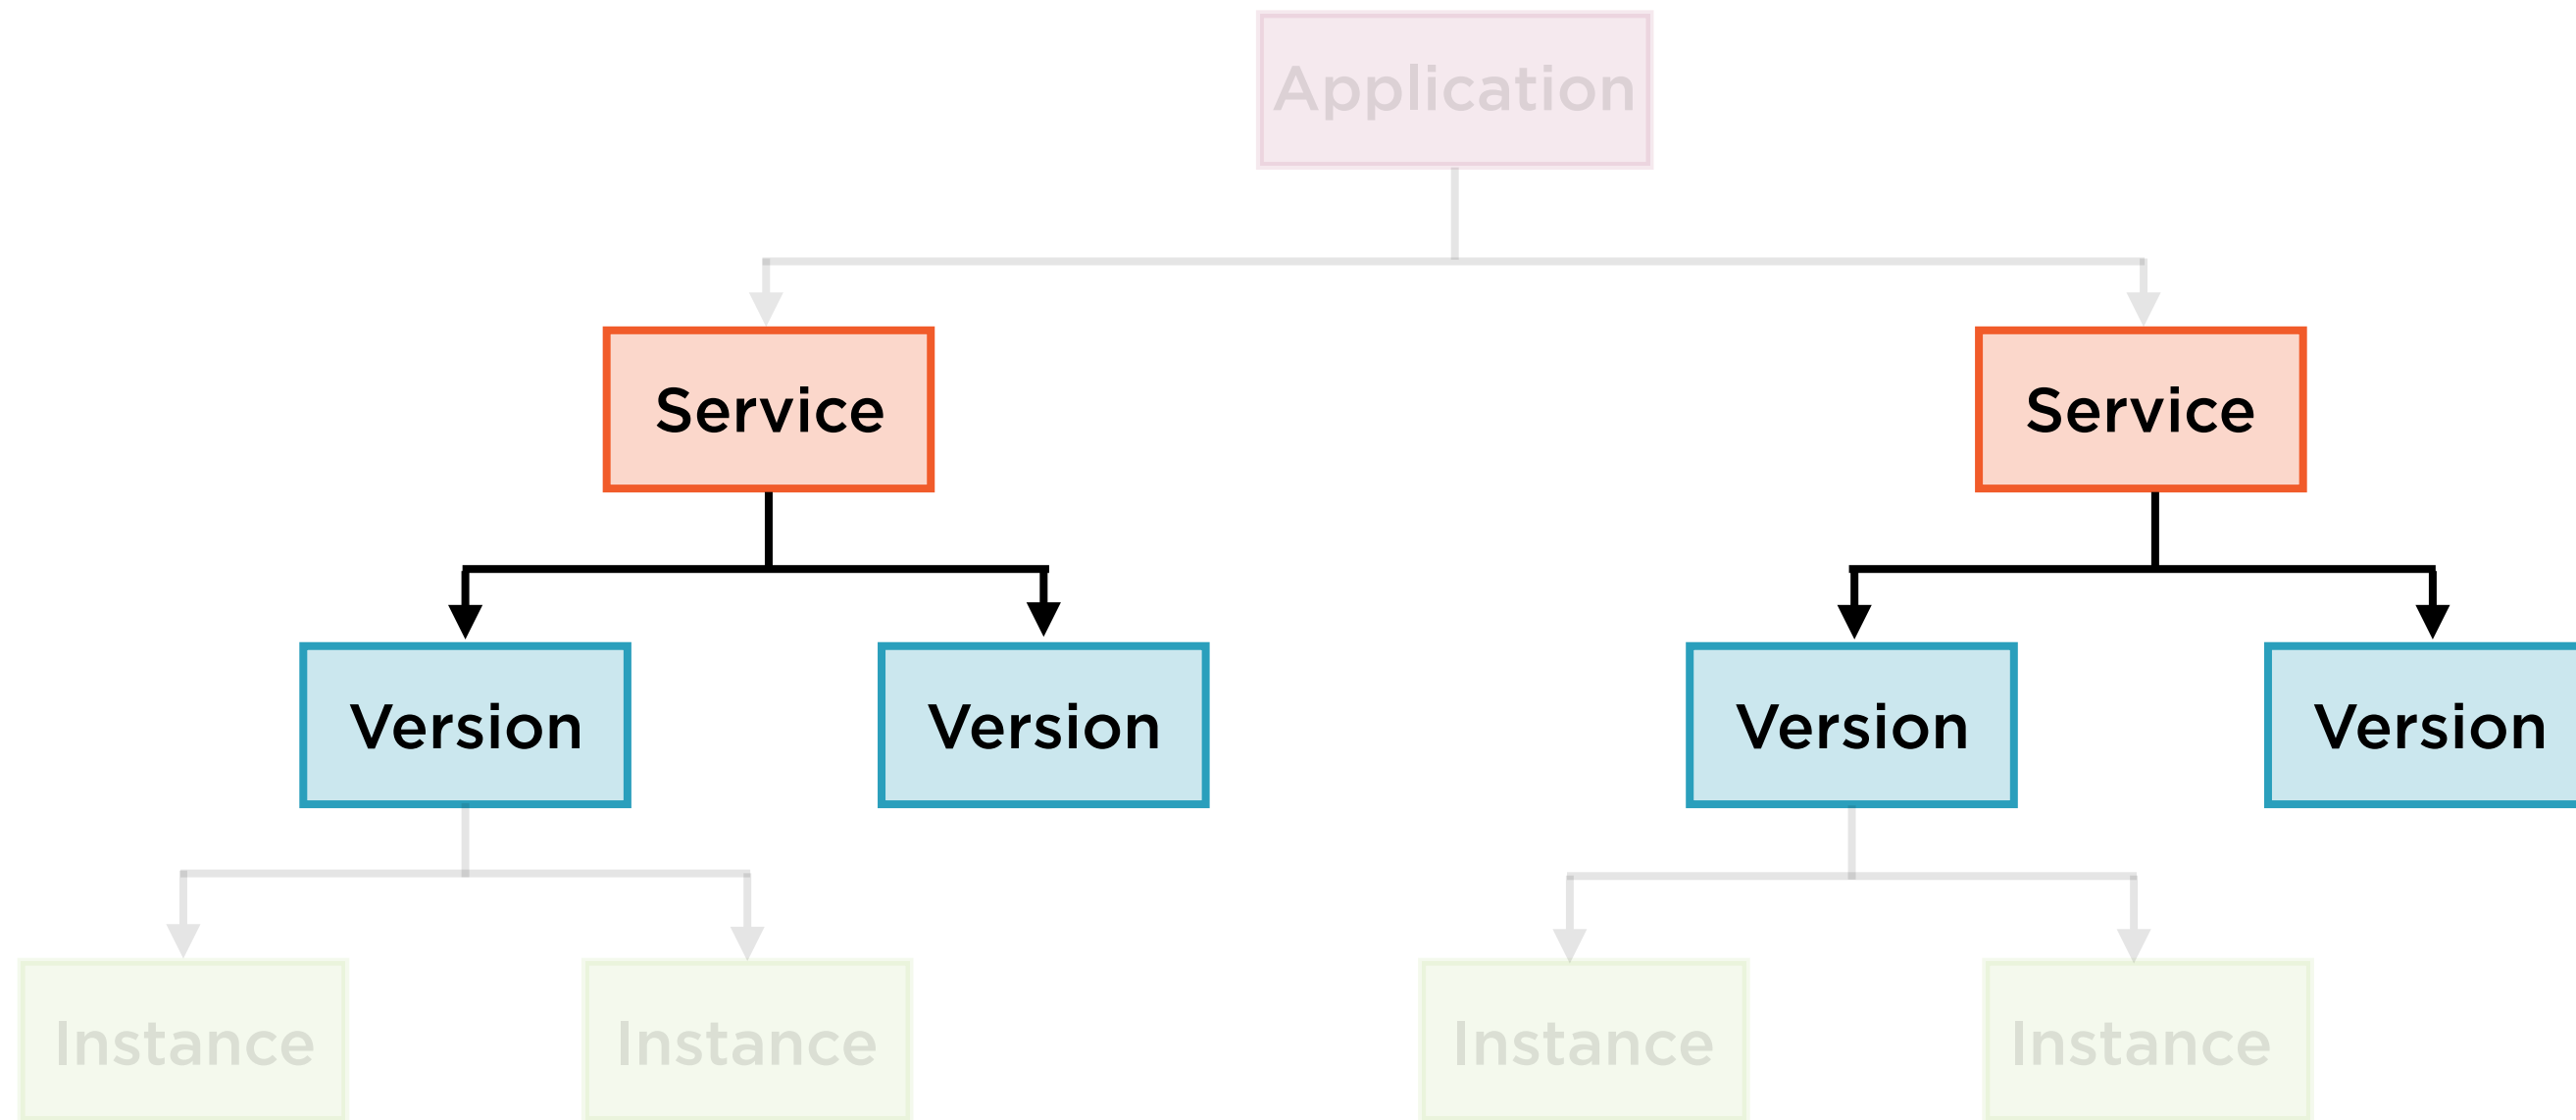
Each service is a logical component which can share features and communicate with one another

Default Service




Every application includes at least one service, the **default service**

Versions



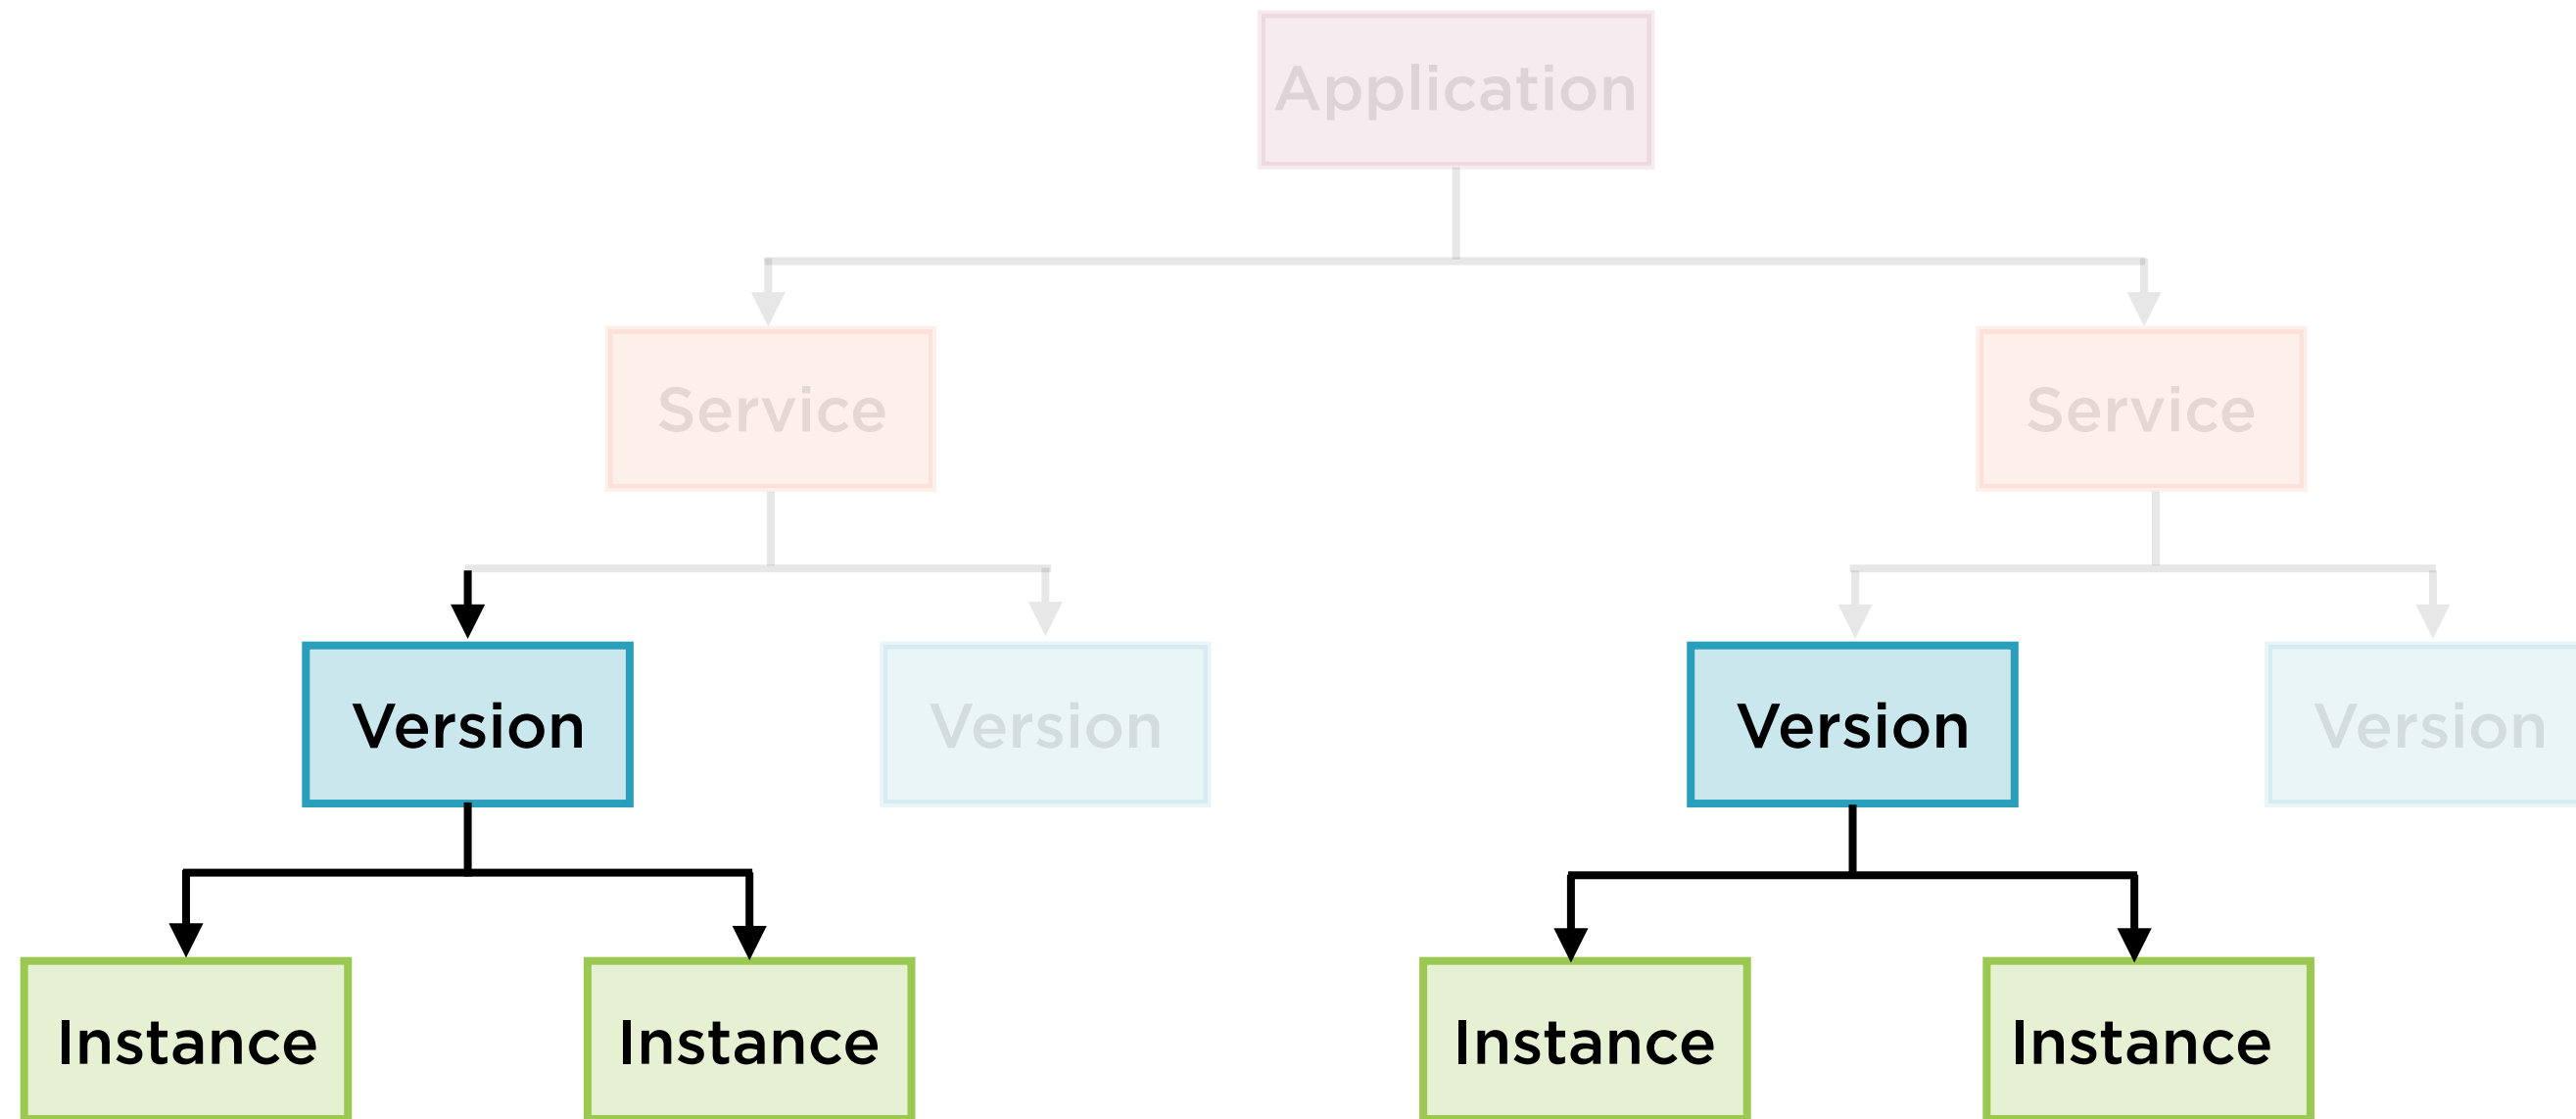
**Multiple versions of every service can be deployed,
traffic is automatically sent to the latest version**

Versions



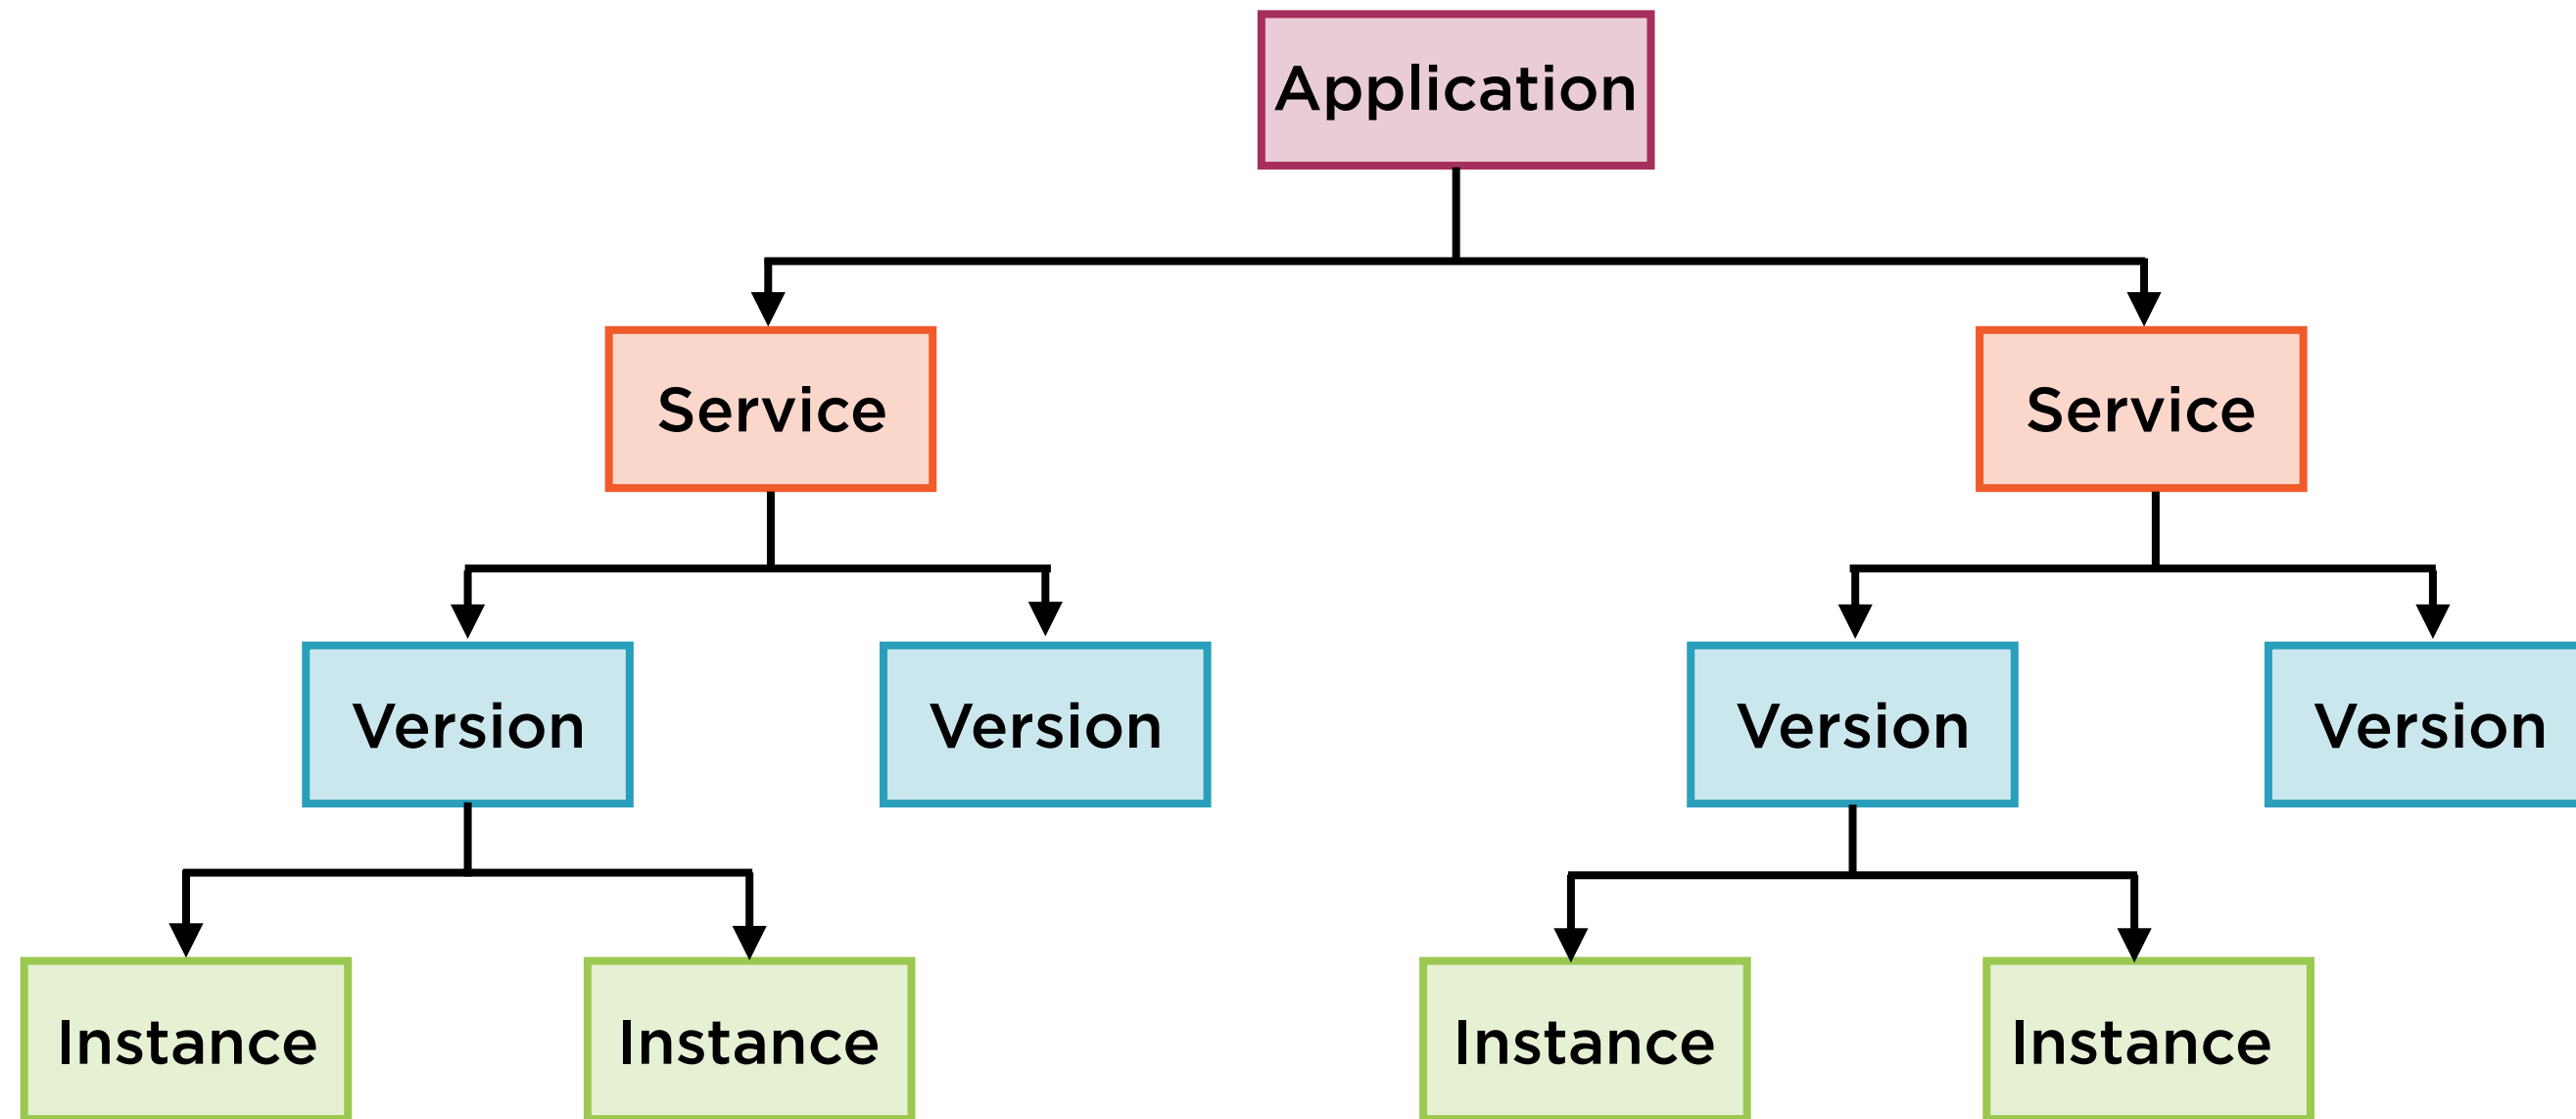
**Every new deployment to a particular service
creates a new version**

Instances



Versions run on one or more instances - can configure App Engine to automatically scale instance based on load

Components of an Application



Requests



Routed to the services or versions that are configured to handle traffic

Target and route requests to specific services and versions

Instance Management

App Engine Environments

Standard

Apps that experience traffic spikes

Instance startup in seconds

Scaling: Manual, Basic or Automatic

No background processes

No SSH debugging

Scale to zero

No installation of third-party binaries

Flexible

Apps that experience consistent traffic

Instance startup in minutes

Scaling: Manual or Automatic

Background processes supported

SSH debugging supported

No scaling to zero (minimum 1 instance)

**Installation of third-party binaries
allowed**

App Engine Environments

Standard

Apps that experience traffic spikes

Instance startup in seconds

Scaling: Manual, Basic or Automatic

No background processes

No SSH debugging

Scale to zero

No installation of third-party binaries

Flexible

Apps that experience consistent traffic

Instance startup in minutes

Scaling: Manual or Automatic

Background processes supported

SSH debugging supported

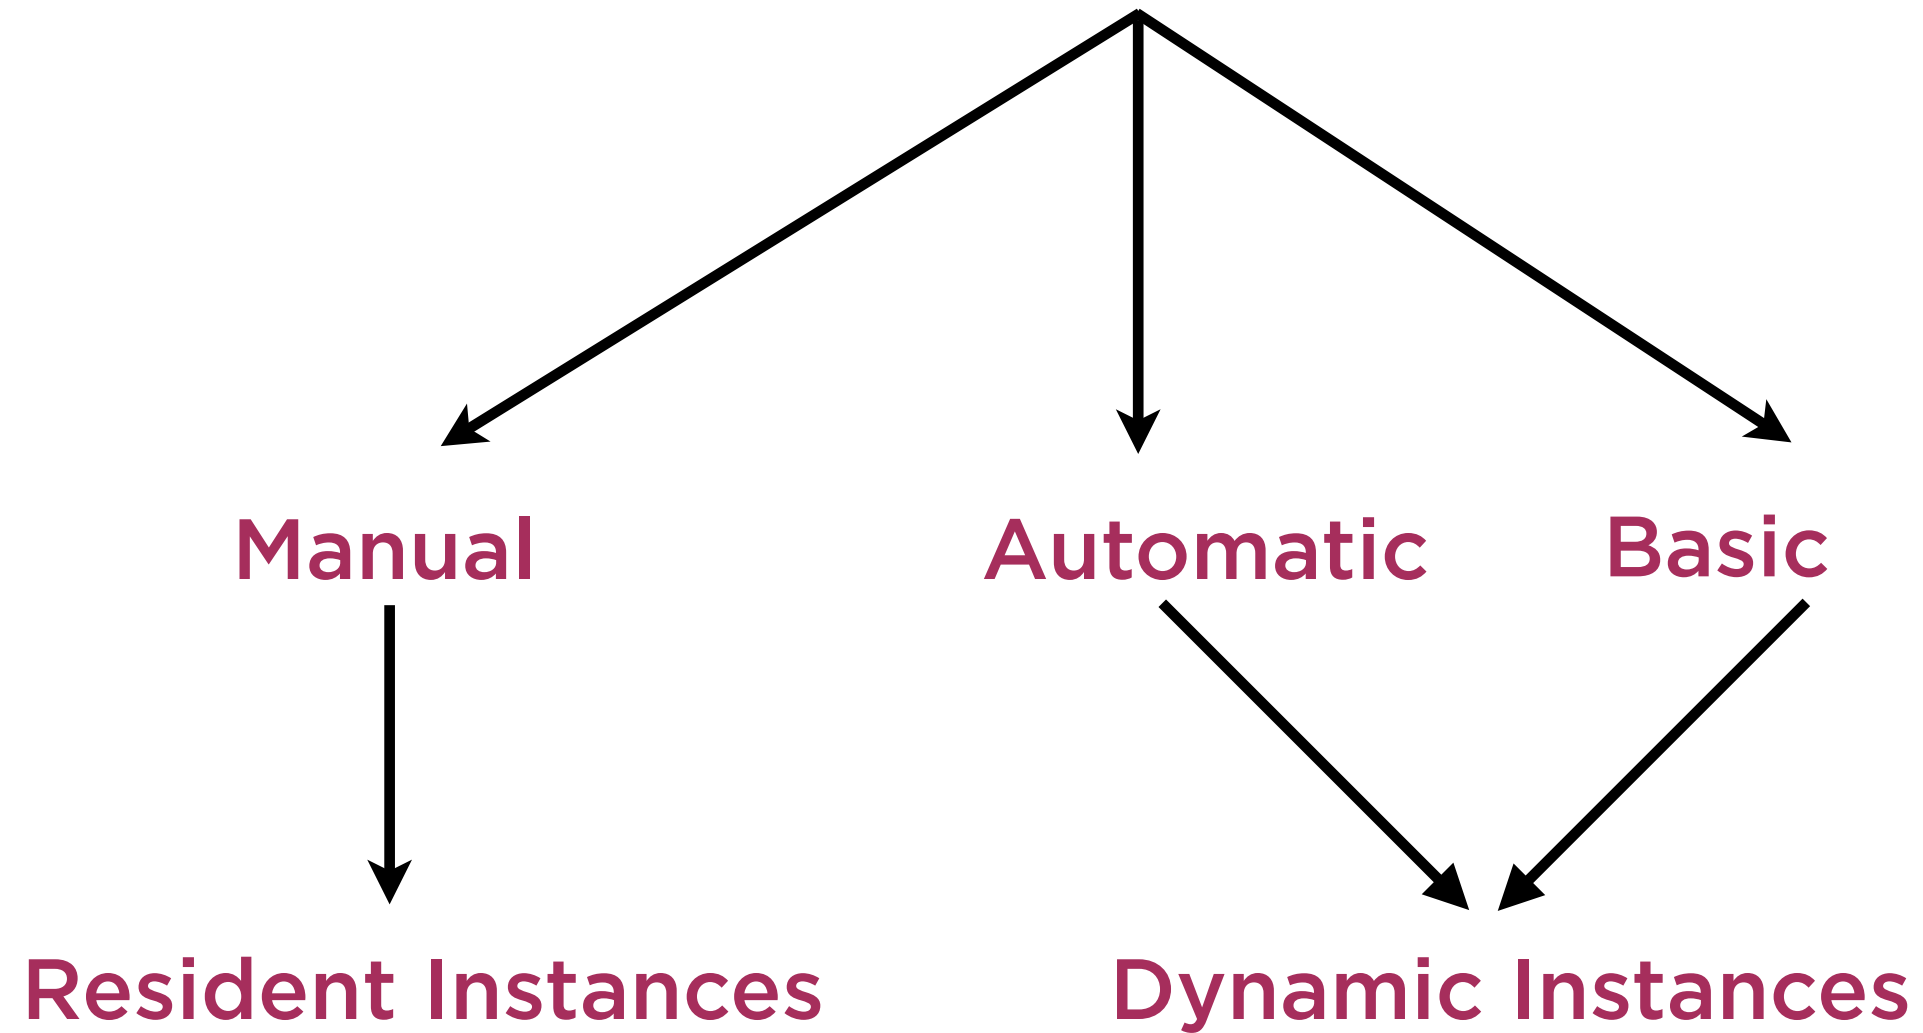
No scaling to zero (minimum 1 instance)

Installation of third-party binaries
allowed

Instances

Computing units on which the application is hosted.
Includes the language runtime, the App Engine APIs, the application code and memory

Scaling and Instance Types




Two Types of Instances


Resident

Dynamic


Three Ways to Scale



Manual



Automatic



Basic

Resident Instances



The diagram consists of two side-by-side rectangular boxes. The left box is dark gray and contains the word 'Resident' in white text. The right box is light gray and contains the word 'Dynamic' in white text.

Resident

Dynamic

Run all the time and are always ready to serve traffic - having resident instances run your app can improve performance

Dynamic Instances




The diagram consists of two side-by-side rectangular boxes. The left box is light gray and contains the word 'Resident' in white text. The right box is dark gray and contains the word 'Dynamic' in white text. Below these boxes is a line of black text.

Resident

Dynamic

Starts up and shuts down automatically based on your application's needs

Three Ways to Scale



Manual




Automatic




Basic

Manual Scaling



Manual



Automatic




Basic

Only uses resident instances to serve traffic, useful for applications that rely on the state of memory over time

Automatic Scaling



Manual



Automatic




Basic

**Use dynamic instance to scale up and down based on the load,
can specify a minimum number of resident instances**

Basic Scaling



Manual



Automatic



Basic

Only uses dynamic instances which are turned down when app is idle - best for intermittent jobs

Instance Classes

Every instance belongs to an instance class which determines its compute resources and pricing

Instance Classes

Instance Class	Memory Limit	CPU Limit	Supported Scaling Types
F1 (default)	128 MB	600 MHz	automatic
F2	256 MB	1.2 GHz	automatic
F4	512 MB	2.4 GHz	automatic
F4_1G	1024 MB	2.4 GHz	automatic
B1	128 MB	600 MHz	manual, basic
B2 (default)	256 MB	1.2 GHz	manual, basic
B4	512 MB	2.4 GHz	manual, basic
B4_1G	1024 MB	2.4 GHz	manual, basic
B8	1024 MB	4.8 GHz	manual, basic

Instance Classes



Instance classes starting with **F** support automatic scaling

Instance classes starting with **B** support manual and basic scaling

Instance Lifecycle

Instance Lifecycle

Running

Stopped

Startup

Warmup

Loading request

Shutdown

All instances of same service and version share same state

Running

Stopped

Auto-scaled instances are always running

Manual/basic scaled instances may be running or stopped

Can stop instances using gcloud, web console or programmatically

Startup

Each service instance created in response to start request

Empty HTTP GET request to `/_ah/start`

Sent by platform, not users

Handled differently in manual, basic and automatic scaling

Warmup

Specific type of loading request

Sent by platform to load app

In advance of live requests

Only for automatic scaling instances

Loading request

Occurs during first request

This is called the loading request

Instance loads libraries and resources

Initialization slows down request processing



Shutdown

Usually shutdown hook can run before instance terminates

Sometimes may not happen

Possible reasons

- Manual stop
- Deployment of new version
- Max memory exceeded
- Instance moved for system reasons

Demo


**Building and deploying a simple App
Engine application**

Demo

Autoscaling in App Engine

Pricing

Pricing




For the standard environment based on instance class

Cost per-instance, per-hour

Flexible environment pay for vCPU, memory and persistent disk

Pricing



**Additional charges for Cloud Datastore,
Memcache, network traffic**

<https://cloud.google.com/appengine/pricing>

Summary

Hosted applications with built-in load balancing and autoscaling

Services and versions

Instances, instance classes and scaling

Choosing between the standard and flexible environment

Deploying and scaling a simple App Engine application