

Deploying Applications on the App Engine Standard Environment



Vitthal Srinivasan

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

**Standard environment for restrictive
lightweight sandbox**

Very few supported runtimes

Very fast scaling

**Integrate with Cloud Datastore for
storage**

**Memcache for in-memory distributed
caching**

App Engine Standard Environment

App Engine Environments

Standard Environment

Flexible Environment

Standard Environment Runtimes

First Generation

Whitelisted extensions and libraries only

External access via URL Fetch API

No file system access

App Engine-specific runtimes

Proprietary sandbox mechanism

Python 2.7, PHP 5.5, Go 1.9

Second Generation

Any extension or library

Full access to external network


Read and write access to /tmp

Unmodified language runtimes

gVisor-based container sandbox

Python 3.7, PHP 7.2, Node.js 8 and 10, Go 1.11, Java 8


Application Execution



App runs in lightweight instance inside sandbox

Very restrictive

App Engine APIs




Standard environment uses built-in App Engine APIs (not client libraries)

Datastore

Memcache

Task queues

Local Development



Run in App Engine SDK

Code need not be portable

Scaling



Scaling from zero

Custom-designed, fast scaling

Zero to thousands very quickly


Health Checks



No health checks used or needed
Instances monitored by platform

Traffic Splitting and Migration

Migrating Traffic




Deploying a new version causes immediate migration of traffic

May cause spike in latency of requests

Migrate traffic gradually

- Add warmup requests to your application
- New instance receive warm up requests before actual traffic

Splitting Traffic



Specify a percentage distribution of traffic

Splitting traffic allows you to conduct A/B testing

Traffic splitting is applied to URLs that do not explicitly target a version

Traffic Splitting




The diagram illustrates two methods for traffic splitting. It consists of a title 'Traffic Splitting' at the top, followed by two side-by-side colored squares. The left square is teal and contains the text 'IP address'. The right square is maroon and contains the text 'Cookies'.

IP address

Cookies

Splitting Traffic Across Multiple Versions



IP address splitting

- Hashes the IP address to a value between 0-999
- Uses that number to route the request

Cookie splitting

- If cookie exists use its value to route
- Otherwise route traffic randomly

Demo

Deploying applications to a specific service


Traffic splitting and migration

Datastore

Google Cloud Datastore

Schemaless NoSQL document database service that provides autoscaling, ACID transactions and SQL-like queries

Choose Datastore For



Highly structured data

- XML, HTML


Highly scaled web serving

- Need very fast reads
- No need for interactivity

Transactional (ACID) support

- With non-relational data

Avoid Datastore For



Relational data data needing full SQL support

Unstructured data e.g. blobs

Interactive querying

Datastore Data Model

Datastore	RDBMS Equivalent
Kind	Table
Entity	Row
Property	Column
Key	Primary key

Datastore Data Model



Kind is the category of the object i.e. Products, Customers



Entity refers to objects of a particular kind




Properties are the attributes of entities



Keys are used to uniquely identify entities

Ancestors



Entities are arranged hierarchically within a namespace

Likes files in a directory system

An entity can have a designated parent

Or can be a root entity



Entity keys are a combination of:
Namespace:Kind:Identifier:Ancestor_Path

Demo

**Implementing a guestbook review
application using Cloud Datastore**

Memcache

Memcache

Distributed in-memory cache in front of, or in place of, persistent storage

Memcache



The diagram consists of two large squares side-by-side. The left square is teal and contains the word 'Shared' in white. The right square is maroon and contains the word 'Dedicated' in white.

Shared

Dedicated

Shared Memcache



Free default for App Engine applications

Provides cache capacity on a best-effort basis

Shared across all App Engine applications

Dedicated Memcache



Provides a fixed cache capacity assigned exclusively to your application

It's billed by the GB-hour of cache size

More predictably and with fewer reads from more costly durable storage

Demo

**Integrating caching in the guestbook
application using Memcache**

Summary

**Standard environment for restrictive
lightweight sandbox**

Very few supported runtimes

Very fast scaling

**Integrate with Cloud Datastore for
storage**

**Memcache for in-memory distributed
caching**